

AI From Zero

Sep 28, 2026 · @Alireza

How to use this book

This book takes you from zero to understanding and building with modern AI, one chapter at a time. It assumes no background in AI, math or programming.

Each chapter has the same four parts:

- **The idea** — a plain-language explanation, usually with an everyday analogy.
- **The picture** — a diagram when the idea has a shape.
- **The code** — short Python you type and run yourself. Reading code teaches little; running and breaking it teaches a lot.
- **Try it** — two or three small exercises. Do them before moving on.

The book has 24 chapters in five parts, plus a math appendix and a concept checklist. Work through them in order; each builds on the last. Ask me whenever you want a chapter expanded, more exercises, or solutions. If a paragraph confuses you, highlight it and leave a comment — I'll rewrite it until it clicks.


What you need: a computer (or just a browser) and about 30–60 minutes per chapter. Chapter 2 sets everything up.

Chapter 1: The big picture

AI, machine learning, deep learning and generative AI are not four separate things. They are nested layers, each one a more specific kind of the one around it.

1.1 Four words, one family

Generative AI sits inside deep learning, inside machine learning, inside AI



The AI family · four nested layers

Read it from the outside in:

- **Artificial Intelligence (AI)** is the whole field: making computers do things we'd call "smart". Early AI was mostly hand-written rules, like a chess program that follows strategies a human typed in.
- **Machine Learning (ML)** is AI that *learns from examples*. You don't write the rules; you show the computer data and it figures out the rules itself.
- **Deep Learning (DL)** is ML using *neural networks* with many layers. It shines on messy data like images, sound and language.
- **Generative AI (GenAI)** is deep learning that *creates* new things — a paragraph, a picture, a song — instead of only labelling or predicting.

1.2 The big shift: from rules to learning

Imagine building a spam filter. The old way, you write rules: "if the email contains *FREE MONEY*, mark it spam". Spammers change one word and your rule fails.

The machine learning way, you collect 10,000 emails already labelled "spam" or "not spam" and let the computer find the patterns. When spam changes, you retrain on new examples.

	Traditional programming	Machine learning
You provide	Rules + data	Data + correct answers
The computer produces	Answers	The rules (<i>a model</i>)
When the world changes	You rewrite the rules	You retrain on new data

1.3 What is a "model"?

A model is a mathematical function with adjustable knobs, called **parameters**. Data goes in, a prediction comes out, and the knobs decide how one turns into the other.

Training means showing the model examples, checking how wrong its predictions are, and nudging the knobs to be a little less wrong. Repeat millions of times and the model gets good.

A simple model might have 2 knobs. A modern large language model has billions. The idea is the same — you'll build a 2-knob model yourself in Chapter 3.

1.4 What is a transformer? (a first look)

A **transformer** is a particular design of neural network, introduced by Google researchers in 2017 in a paper called *Attention Is All You Need*. Almost every modern language AI — ChatGPT, Claude, Gemini — is built on it. The "T" in GPT stands for Transformer (*Generative Pre-trained Transformer*).

Its core trick is called **attention**. When the model reads a sentence, every word gets to *look at every other word* and decide which ones matter for understanding it. Take this sentence:

The animal didn't cross the street because it was too tired.

What does "it" mean? You know instantly it's the animal, not the street. Attention lets the model learn to connect "it" strongly to "animal". Change "tired" to "wide" and a good model shifts "it" towards "street".

Why was this a breakthrough? Older networks read text one word at a time, left to right, and tended to forget the beginning of long passages. Transformers look at all the words at once. That makes them better at long-range meaning and, crucially, fast to train on modern chips (GPUs). Fast training meant researchers could train them on huge amounts of text — which gave us **large language models (LLMs)**.

An LLM's job sounds almost too simple: **predict the next word**. Given "The cat sat on the", it scores every possible next word and picks a likely one ("mat"). Then it adds that word and predicts again. One word at a time, it writes whole essays. Everything impressive it does comes from getting very, very good at that one task.

(Strictly, models work with **tokens** — words or pieces of words — not whole words. We'll ignore that until Part 4, where we build a transformer piece by piece.)

1.5 A taste of code: the world's tiniest text generator

You don't need anything installed yet. Open [Google Colab](#), click *New notebook*, paste this into a cell, and press the play button.

```
from collections import Counter, defaultdict
import random

text = "the cat sat on the mat the cat ate the fish the dog sat on the rug"
words = text.split()

# Count which word follows which
next_words = defaultdict(Counter)
```

```

for current, following in zip(words, words[1:]):
    next_words[current][following] += 1

print(next_words["the"])  # what tends to come after "the"?

# Generate a sentence, one word at a time
word = "the"
sentence = [word]
for _ in range(8):
    options = next_words[word]
    if not options:
        break
    word = random.choices(list(options), weights=list(options.values()))[0]
    sentence.append(word)

print(" ".join(sentence))

```

This program "learned" from data (it counted word pairs) and it "generates" (it predicts the next word, again and again). That is the same basic loop as an LLM.

The difference is enormous, though. This toy only looks at **one** previous word. A transformer looks at **all** previous words through attention, and learns billions of parameters instead of a small table of counts. Run it a few times — the random choice gives a different sentence each time.

1.6 A short history of AI

AI is older than most people think. The ideas came early; what changed recently is that we finally had enough data and computing power to make them work.

Year	Milestone	Why it mattered
1950	Alan Turing asks "Can machines think?" and proposes the <i>imitation game</i> (the Turing test)	Framed the goal of the field
1956	The Dartmouth workshop coins the term "artificial intelligence"	AI becomes a research field
1958	Frank Rosenblatt's <i>perceptron</i> , the first neural network that learns from examples	The seed of deep learning
1960s–1980s	Symbolic AI and expert systems : experts hand-write thousands of if-then rules	Worked in narrow domains; brittle and expensive to maintain
1970s, late 1980s	Two AI winters : results fall short of promises, funding collapses	A lesson in hype that is still relevant
1986	Backpropagation popularized for training multi-layer	The core training algorithm still used today (Chapter 10)

Year	Milestone	Why it mattered
	networks	
1997	IBM's Deep Blue beats chess champion Garry Kasparov	Mostly brute-force search plus hand-tuned rules, not learning
2000s	Classic ML (SVMs, random forests) powers spam filters, search ranking and recommendations	ML becomes everyday technology
2012	AlexNet wins the ImageNet image-recognition contest by a huge margin, trained on GPUs	Starts the deep learning boom (Chapter 12)
2016	DeepMind's AlphaGo beats Go champion Lee Sedol	Deep learning + reinforcement learning (Chapter 20)
2017	The transformer is published	The architecture behind modern LLMs (Chapter 15)
2018–2020	BERT, GPT-2, GPT-3 show that bigger transformers trained on more text keep getting better	The era of large language models
2022	ChatGPT and open image generators reach the public	Generative AI goes mainstream
2023 onwards	Models that see images, use tools, write long reasoning, and act as agents	Where we are now (Chapters 16–18)

Three ingredients drove the modern boom together: **data** (the internet produced enormous amounts of text and images), **compute** (GPUs, originally built for video games, turned out to be perfect for neural network maths), and **algorithms** (better architectures and training tricks). Remove any one and progress stalls.

1.7 Narrow AI vs general AI

Almost every AI system ever built is **narrow AI**: it does one kind of task. A chess engine can't recognize a cat; a spam filter can't translate French.

Artificial general intelligence (AGI) usually means a system that can learn and perform most intellectual tasks a person can, at a human level or beyond. It's a goal, not a product, and experts disagree about how to define it and when, or whether, it will arrive.

Large language models sit somewhere in between: one model can write, translate, summarize, code and answer questions across many subjects, which is far more general than older systems. But they still make basic mistakes a person wouldn't, so treat them as powerful tools rather than minds.

1.8 What today's AI is good and bad at

Every AI model is, at heart, a **pattern-finder**: it predicts based on patterns in its training data. That explains both its strengths and its weaknesses.

Usually good at	Often weak at
Spotting patterns in huge amounts of data	Situations unlike anything in its training data
Language tasks: drafting, summarizing, translating, rewording	Guaranteeing facts are correct (it can hallucinate)
Recognizing images, speech and handwriting	Exact arithmetic and counting, unless given a calculator tool
Writing and explaining code	Knowing what it doesn't know
Doing a repetitive task consistently, at any hour	Events after its training data ends, unless given search
Generating many options quickly	Telling cause from coincidence

Two rules follow. **Garbage in, garbage out**: a model trained on biased or wrong data repeats the bias and the errors (Chapter 24). And **verify what matters**: use AI to draft and assist, and check its output when the stakes are high.

Try it

1. Replace text with a paragraph of your own (a news article works well) and generate again. Does it sound more or less natural?
2. Change the starting word from "the" to another word in your text.
3. Think: why does the output often stop making sense after 4–5 words? (Hint: how much of the sentence does the program remember?)

Key words from this chapter

Term	Plain meaning
Model	A function with adjustable knobs that turns inputs into predictions
Parameters	The knobs; learned during training
Training	Adjusting parameters so predictions become less wrong
Neural network	A model made of layers of simple connected units
Transformer	A neural network design built around attention; the basis of modern LLMs
Attention	The mechanism that lets each word weigh how

Term	Plain meaning
	much every other word matters
LLM	Large language model: a big transformer trained to predict the next token
Token	A word or piece of a word; the unit a language model reads and writes

Chapter 2: Your toolkit

You need exactly one tool to start: Google Colab, a free Python notebook that runs in your browser. Everything in this book runs there.

2.1 Why Python, why notebooks

Python is the language of AI. Nearly every AI library — NumPy, scikit-learn, PyTorch, Hugging Face — is written for it, and its syntax reads almost like English.

A **notebook** mixes code cells and text cells. You run one cell at a time and see the result immediately under it. That makes it perfect for experimenting: change a number, re-run, see what happens.

2.2 Set up in two minutes

1. Go to colab.research.google.com and sign in with a Google account.
2. Click **New notebook**.
3. Type `print("hello, AI")` in the cell and press **Shift + Enter**. The output appears below.
4. Later, for deep learning chapters, you'll switch on a free GPU: *Runtime* → *Change runtime type* → *GPU*.

Prefer working offline? Install Python from python.org and run `pip install jupyter numpy matplotlib scikit-learn`, then start `jupyter notebook`. Colab is simpler, so the book assumes Colab.

2.3 Python in ten minutes

This is only the Python you need for the next few chapters. Run each block in its own cell.

Variables and types — a name that holds a value.

```
hours = 3           # an integer
price = 19.99        # a decimal number (float)
name = "Alireza"     # text (string)
is_ready = True      # True or False (boolean)
print(name, hours * price)
```

Lists — an ordered collection. Data in ML is mostly lists of numbers.

```
scores = [72, 85, 90, 64]
print(scores[0])      # first item: 72
print(len(scores))    # how many: 4
```

```
scores.append(99)          # add one
print(sum(scores) / len(scores))  # average
```

Loops — do something for every item. Training a model is one big loop.

```
for s in scores:
    if s >= 80:
        print(s, "pass with distinction")
    else:
        print(s, "pass")
```

Functions — a reusable recipe. A model is, at heart, a function.

```
def predict_price(size_m2):
    return 3000 * size_m2 + 50000

print(predict_price(80))  # 290000
```

Dictionaries — look things up by name.

```
person = {"name": "Sara", "age": 31}
print(person["age"])
```

2.4 NumPy: math on whole lists at once

NumPy is the library that makes AI math fast. Instead of looping over numbers one by one, you do math on entire arrays in one line.

```
import numpy as np

sizes = np.array([50, 80, 120])      # house sizes in m²
prices = 3000 * sizes + 50000        # applies to every item at once
print(prices)                        # [200000 290000 410000]
print(prices.mean())                 # average price
```

Notice that `predict_price` from 2.3 and this line do the same thing. The NumPy version works on thousands of houses as fast as on one.

2.5 Matplotlib: seeing your data

In AI you always look at your data before modelling it.

```
import matplotlib.pyplot as plt

plt.scatter(sizes, prices)
plt.xlabel("Size (m²)")
plt.ylabel("Price")
plt.show()
```

2.6 More Python you'll see constantly

Slicing — take part of a list. Indexes start at 0; negative numbers count from the end.

```
scores = [72, 85, 90, 64, 99]
print(scores[1:3])  # [85, 90]  items 1 and 2 (the end index is excluded)
print(scores[-1])   # 99        the last item
print(scores[:2])   # [72, 85]  the first two
```

List comprehensions — build a list in one line. You'll see these everywhere in AI code.

```
squares = [x ** 2 for x in range(5)]          # [0, 1, 4, 9, 16]
high = [s for s in scores if s >= 80]         # [85, 90, 99]
words = [w.lower() for w in "The Cat Sat".split()] # ['the', 'cat', 'sat']
f-strings — put values inside text, with formatting.
```

```
accuracy = 0.87654
print(f"Accuracy: {accuracy:.2%}")           # Accuracy: 87.65%
print(f"Loss: {0.123456:.3f}")               # Loss: 0.123
```

2.7 NumPy in depth: shape, indexing, axes, broadcasting

In ML, data is almost always a 2-D array with shape **(number of examples, number of features)**. Getting shapes wrong is the most common bug in AI code, so always print `.shape` when something breaks.

```
import numpy as np

X = np.array([[1, 2, 3],
              [4, 5, 6]])      # 2 examples, 3 features
print(X.shape)                # (2, 3)
print(X[0])                   # first row:    [1 2 3]
print(X[:, 1])                # second column: [2 5]
print(X[X > 2])               # filtering:   [3 4 5 6]
print(X.T.shape)              # transpose:   (3, 2)
print(X.reshape(3, 2))        # same 6 numbers, new shape

Axes: axis=0 means "down the rows" (one result per column); axis=1 means "across the columns"
(one result per row).
```

```
print(X.mean(axis=0))          # mean of each feature: [2.5 3.5 4.5]
print(X.sum(axis=1))           # total of each example: [ 6 15]
```

Broadcasting: NumPy automatically stretches a smaller array to match a bigger one. This one line standardizes every column, which is exactly what `StandardScaler` does in Chapter 4:

```
X_std = (X - X.mean(axis=0)) / X.std(axis=0)   # (2,3) minus (3,) works: the (3,) row
is reused for each row
Useful creators: np.zeros((2, 3)), np.ones(5), np.arange(0, 10, 2), np.linspace(0, 1, 5),
np.random.randn(100, 3) (random numbers from a bell curve).
```

2.8 pandas in five minutes

pandas handles tables with named columns — think of a spreadsheet you control with code. A table is a **DataFrame**; one column is a **Series**.

```
import pandas as pd

df = pd.DataFrame({
    "name": ["Sara", "Omid", "Lena", "Ali"],
    "age": [31, 45, 28, 39],
    "city": ["Toronto", "Ottawa", "Toronto", "Montreal"],
})
print(df["age"].mean())        # 35.75
print(df[df["age"] > 30])      # rows where age > 30
print(df.groupby("city")["age"].mean()) # average age per city
print(df.sort_values("age", ascending=False))
```

```
df["age_next_year"] = df["age"] + 1      # new column
# df = pd.read_csv("file.csv")          # load a real file
```

2.9 Packages and virtual environments

A **package** (or library) is code someone else wrote that you install with `pip`. In Colab, put `!` in front: `!pip install tiktoken`.

On your own computer, use a **virtual environment** for each project: a private folder of packages, so project A's library versions can't break project B.

```
python -m venv .venv                # create it once per project
source .venv/bin/activate           # macOS / Linux   (Windows:
.venv\Scripts\activate)
pip install numpy pandas scikit-learn
pip freeze > requirements.txt        # record exact versions (Chapters 21-22 rely
on this)
```

2.10 Reading error messages

Errors are normal. Read the **last line first**: it names the error type and the cause. The lines above show where it happened.

Error	Usual cause
<code>NameError: name 'x' is not defined</code>	Typo, or you didn't run the cell that creates <code>x</code>
<code>TypeError</code>	Wrong kind of value, e.g. adding text to a number
<code>IndexError / KeyError</code>	Asking for a list position or dictionary key that doesn't exist
<code>ModuleNotFoundError</code>	The package isn't installed: <code>pip install</code> it
<code>ValueError: shapes ... not aligned</code>	Array shapes don't match: print <code>.shape</code> of each

When stuck, paste the full error into a search engine or an AI assistant along with the code that caused it.

Try it

1. Make a list of five temperatures in °C and convert them all to °F with NumPy ($F = C \times 9/5 + 32$).
2. Write a function `is_hot(temp_c)` that returns True above 25.
3. Plot your five temperatures with `plt.plot()` instead of `plt.scatter()`. What's the difference?

Chapter 3: How a machine learns

Every model in this book — from a 2-knob line to a billion-knob transformer — learns with the same loop: guess, measure the error, nudge the knobs, repeat. In this chapter you build that loop from scratch.

3.1 The problem

You have 50 houses with their size and sale price. You want to predict the price of a house you haven't seen. Plot the data and it roughly forms a line: bigger houses cost more.

So our model is a line with two knobs:

$$\text{predicted price} = w \cdot \text{size} + b$$

w (the **weight**) is how much each extra m² adds. b (the **bias**) is the starting price. Training means finding the w and b that fit the data best.

3.2 Measuring "wrong": the loss

For each house, compare the prediction with the real price. Square the difference (so negatives don't cancel positives, and big misses count extra), then average over all houses. This number is the **loss** — here, the *mean squared error*:

$$L = \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

$\hat{y}$ is the prediction, y is the real price, n is the number of houses. Lower loss = better model. Training is simply "make L as small as possible".

3.3 Walking downhill: gradient descent

Imagine standing on a foggy hillside, trying to reach the valley. You can't see far, but you can feel which way the ground slopes under your feet. So you take a small step downhill, feel again, step again.

That's **gradient descent**. The "slope" is the **gradient**: it tells us, for each knob, whether turning it up would raise or lower the loss. We turn each knob a little in the direction that lowers it:

$$w \leftarrow w - \eta \frac{\partial L}{\partial w} \quad b \leftarrow b - \eta \frac{\partial L}{\partial b}$$

η (eta) is the **learning rate**: the size of each step. Too small and training crawls. Too big and you overshoot the valley and bounce around.

You don't need to calculate gradients by hand — libraries do it automatically. For this simple model the formulas are short, so we write them out once to see there's no magic.

3.4 The code: a model that learns

```
import numpy as np
import matplotlib.pyplot as plt

# 1. Make some data: 50 houses, true rule = 3000 per m² + 50,000, plus noise
np.random.seed(0)
sizes = np.random.uniform(40, 150, 50)
prices = 3000 * sizes + 50000 + np.random.normal(0, 20000, 50)

# Scale numbers down so training is stable
x = sizes / 100
```

```

y = prices / 100000

# 2. Start with knobs at zero
w, b = 0.0, 0.0
learning_rate = 0.1

# 3. The training loop
for step in range(1000):
    pred = w * x + b                # guess
    error = pred - y
    loss = (error ** 2).mean()      # measure
    grad_w = 2 * (error * x).mean() # which way is downhill?
    grad_b = 2 * error.mean()
    w -= learning_rate * grad_w     # nudge the knobs
    b -= learning_rate * grad_b
    if step % 200 == 0:
        print(f"step {step}: loss {loss:.4f}")

print(f"Learned: price ≈ {w * 1000:.0f} × size + {b * 100000:.0f}")

# 4. See the result
plt.scatter(sizes, prices, label="real houses")
plt.plot(sizes, (w * x + b) * 100000, color="red", label="model")
plt.xlabel("Size (m²)"); plt.ylabel("Price"); plt.legend(); plt.show()

```

You should see the loss drop fast (about 12.7 to 0.03) and the model learn roughly **2,944 per m² + 52,000**. That's close to the hidden true rule of 3,000 + 50,000; the gap comes from the random noise in the data.

Nobody told the program the rule. It found it from examples. That is machine learning.

3.5 The same thing in three lines

In practice you use a library. **scikit-learn** is the standard one for classic ML:

```

from sklearn.linear_model import LinearRegression

model = LinearRegression().fit(sizes.reshape(-1, 1), prices)
print(model.coef_[0], model.intercept_) # about 2944 and 52090
print(model.predict([[100]]))           # price of a 100 m² house

```

Same answer. Now you know what's happening inside.

3.6 How this connects to transformers

A transformer like the ones behind ChatGPT or Claude trains with *exactly this loop*. The differences are scale, not kind:

	Our house model	A large language model
Knobs (parameters)	2	Billions
Input	One number (size)	A long sequence of tokens
Prediction	A price	The next token

	Our house model	A large language model
Loss	Mean squared error	Cross-entropy (how surprised it was by the real next token)
Learning method	Gradient descent	Gradient descent (a smarter variant called Adam)

3.7 One gradient step, by hand

Seeing the numbers once makes the loop concrete. Take three points that follow $y = 2x$ exactly: $x = [1, 2, 3]$, $y = [2, 4, 6]$. Start with $w = 0$, $b = 0$ and learning rate 0.1.

1. **Guess:** every prediction is $0 \times x + 0 = 0$, so predictions are $[0, 0, 0]$.
2. **Errors** (prediction – truth): $[-2, -4, -6]$.
3. **Loss:** the mean of the squared errors = $(4 + 16 + 36) / 3 \approx \mathbf{18.67}$.
4. **Gradients:** $\text{grad_w} = 2 \times \text{mean}(\text{error} \times x) = 2 \times \text{mean}(-2, -8, -18) \approx -18.67$; $\text{grad_b} = 2 \times \text{mean}(\text{error}) = 2 \times (-4) = -8$. Both are negative, meaning "increase me to reduce the loss".
5. **Update:** $w = 0 - 0.1 \times (-18.67) \approx 1.87$; $b = 0 - 0.1 \times (-8) = 0.8$.

New predictions with $w \approx 1.87$, $b = 0.8$ are $[2.67, 4.53, 6.40]$. The errors are now $[0.67, 0.53, 0.40]$ and the loss is about **0.30** — a 60-fold drop in one step. The predictions now overshoot slightly, so the next step's gradients will be positive and pull w and b back. After a few hundred steps they settle at $w = 2$, $b = 0$.

3.8 Where the gradient formulas come from

The loss for our line is:

$$L(w, b) = \frac{1}{n} \sum_{i=1}^n (wx_i + b - y_i)^2$$

To find how L changes when w changes, use the **chain rule** (Appendix A.2): the derivative of (something)² is $2 \times$ (something), multiplied by the derivative of the something itself. The something is $(wx_i + b - y_i)$, the error. Its derivative with respect to w is x_i , and with respect to b is 1. So:

$$\frac{\partial L}{\partial w} = \frac{2}{n} \sum_{i=1}^n (\hat{y}_i - y_i) x_i \quad \frac{\partial L}{\partial b} = \frac{2}{n} \sum_{i=1}^n (\hat{y}_i - y_i)$$

These are exactly the two lines `grad_w = 2 * (error * x).mean()` and `grad_b = 2 * error.mean()` in the code. Every gradient in deep learning is found the same way, just through many more layers — and frameworks do it for you.

3.9 Batch, stochastic and mini-batch gradient descent

How many examples should you look at before each step?

Variant	Examples per step	Pros	Cons
---------	-------------------	------	------

Variant	Examples per step	Pros	Cons
Batch gradient descent	All of them	Accurate, smooth steps	Very slow and memory-hungry on big data
Stochastic gradient descent (SGD)	1	Many cheap steps; noise can escape bad spots	Very noisy path
Mini-batch gradient descent	Typically 32–512	Best of both; ideal for GPUs	One more setting to choose

Almost all modern training uses mini-batches. Confusingly, people still call it "SGD".

3.10 The shape of the loss: bowls, valleys and local minima

Plot the loss for every combination of w and b and you get a surface. For linear regression with squared error it's a smooth **bowl** with one lowest point (the problem is **convex**), so gradient descent always finds the best answer.

Neural networks have bumpy landscapes with many dips (**local minima**) and flat passes (**saddle points**). In practice this matters less than you'd fear: with millions of dimensions, most dips are nearly as good as each other, and tricks like momentum (Chapter 11) carry training past flat spots.

The learning rate controls the step size on that surface. Too small and you crawl down the slope for hours. Too large and each step overshoots the bottom and lands higher up the other side; the loss grows until it becomes `inf` or `nan`. Try it in the exercises below.

3.11 Why we scaled the inputs

In the code we divided sizes by 100 and prices by 100,000. Without scaling, size (40–150) and price (hundreds of thousands) make the gradient for w enormous compared with the one for b . The bowl becomes a long, narrow valley, and gradient descent zigzags between its walls instead of rolling down the middle. Scaling every feature to similar ranges makes the bowl round, so one learning rate works for all knobs. This is why scaling (Chapter 4.7) matters so much for any model trained by gradient descent.

3.12 Couldn't we just solve it directly?

For linear regression, yes. There's an exact formula, the **normal equation**, which gives the best weights in one calculation, and scikit-learn's `LinearRegression` uses a closely related method:

$$\mathbf{w} = (X^T X)^{-1} X^T \mathbf{y}$$

So why learn gradient descent? Because for almost every other model — logistic regression, neural networks, transformers — no such formula exists, and for huge datasets even the formula is too slow. Gradient descent works for all of them.

Try it

1. Change `learning_rate` to `0.01`, then to `1.5`. What happens to the loss in each case? Why?

2. Change the true rule in step 1 (say 5000 per m²) and check the model finds it.
3. Increase the noise from 20000 to 80000. How close does the model get now?

Part 2 — Classic machine learning

Chapter 4: Data and features

A model can only be as good as its data. Practitioners often say 80% of ML work is preparing data, and only 20% is modelling.

4.1 The vocabulary of a dataset

A dataset is a table. Each **row** is one example (a passenger, an email, a house). Each **column** is a **feature** — a measurable property. The column you want to predict is the **label** or **target**.

Term	Meaning	Titanic example
Example / sample	One row	One passenger
Feature (X)	An input column	Age, ticket class, fare
Label / target (y)	What we predict	Survived (1) or not (0)
Numerical feature	A number with meaning in its size	Age, fare
Categorical feature	One of a fixed set of values	Sex, port of embarkation

Data also comes as text, images, audio and time series. Classic ML mostly works on tables; deep learning (Part 3) handles the rest.

4.2 Load and explore (EDA)

Exploratory data analysis (EDA) means looking at your data before modelling. We'll use the Titanic passenger list and the **pandas** library, Python's spreadsheet.

```
import pandas as pd
import seaborn as sns

df = sns.load_dataset("titanic")
print(df.shape)          # (891, 15): 891 passengers, 15 columns
print(df.head())         # first five rows
df.info()                # column types and non-missing counts
print(df.describe())      # min, max, mean of numeric columns
print(df.isna().sum())    # missing values per column
print(df["survived"].value_counts(normalize=True))  # 62% died, 38% survived
```

Things to look for: missing values, impossible values (a negative age), extreme **outliers**, and how balanced the label is.

4.3 Clean: missing values and outliers

About 20% of ages are missing. You have three options: drop those rows, fill them with a typical value (**imputation**), or predict them with another model. Filling with the median is a simple, robust start.

```
df = df[["survived", "pclass", "sex", "age", "sibsp", "fare", "embarked"]].copy()
df["age"] = df["age"].fillna(df["age"].median())
df = df.dropna(subset=["embarked"]) # only 2 rows missing: just drop them
```

Outliers (a fare of 512 when most are under 50) aren't always errors. Check whether they're real before removing them.

4.4 Feature engineering

Feature engineering means creating new, more useful columns from existing ones. Good features often beat fancier models.

```
df["family_size"] = df["sibsp"] + 1 # siblings/spouses aboard + the passenger
```

Other common tricks: extracting the day of week from a date, the length of a text, or a ratio of two columns.

4.5 Encode categories as numbers

Models only understand numbers. **One-hot encoding** turns a category into 0/1 columns, one per value:

```
df = pd.get_dummies(df, columns=["sex", "embarked"], drop_first=True)
# "sex" becomes sex_male (1 or 0); "embarked" becomes embarked_Q, embarked_S
```

For ordered categories (small < medium < large) use **ordinal encoding**: 0, 1, 2. Don't do this for unordered ones — the model would think "Southampton > Cherbourg".

4.6 Split: train, validation, test

This is the most important habit in ML. You must test your model on data it has **never seen**, or you'll fool yourself.

- **Training set** (~70–80%): the model learns from it.
- **Validation set** (~10–15%): you use it to compare models and choose settings.
- **Test set** (~10–20%): touched **once**, at the very end, for an honest final score.

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop(columns="survived")
y = df["survived"]
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y)
print(X_train.shape, X_test.shape) # (711, 8) (178, 8)
```

stratify=y keeps the same survival ratio in both sets. random_state makes the split repeatable.

(Chapter 8 replaces the single validation set with cross-validation.)

4.7 Scale numeric features

Age runs 0–80 while fare runs 0–512. Many models (kNN, SVM, neural networks, anything using gradient descent) work badly when features have very different ranges.

- **Standardization:** shift and scale to mean 0, standard deviation 1. The usual default.
- **Normalization (min-max):** squeeze into 0–1. Common for images and neural nets.

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler().fit(X_train)      # learn mean/std from TRAIN only
X_train_s = scaler.transform(X_train)
X_test_s = scaler.transform(X_test)        # apply the same numbers to test
```

4.8 The silent killer: data leakage

Data leakage is when information from the test set, or from the future, sneaks into training. Your score looks great; the real-world model fails.

Common causes: fitting the scaler on all data before splitting, filling missing values with a mean computed over all rows, or including a feature that wouldn't exist at prediction time ("was refunded" when predicting fraud). The rule: **split first, then fit every preprocessing step on the training set only**. Chapter 8's pipelines enforce this automatically.

4.9 Features from text

Models need numbers, so text must be converted. The classic methods count words.

- **Bag of words:** one column per word in the vocabulary; each cell counts how often that word appears in the document. Word order is lost.
- **n-grams:** also count pairs ("not good") or triples of words, to keep some order.
- **TF-IDF** (term frequency × inverse document frequency): weight each word by how often it appears in *this* document, discounted by how many documents contain it. Words like "the" appear everywhere and get low weight; distinctive words get high weight.

$$\text{tf-idf}(t, d) = \text{tf}(t, d) \times \log \frac{N}{\text{df}(t)}$$

N is the number of documents and $\text{df}(t)$ the number containing term t . (scikit-learn uses a slightly smoothed version of this.)

```
from sklearn.feature_extraction.text import TfidfVectorizer

docs = ["the driver braked hard",
        "the driver was speeding",
        "heavy rain and hard braking"]
vec = TfidfVectorizer()
M = vec.fit_transform(docs)          # 3 documents x vocabulary size
print(vec.get_feature_names_out())
print(M.toarray().round(2))          # "speeding" scores higher than "the" in
document 2
```

TF-IDF plus logistic regression is still a strong, fast baseline for text classification. Modern systems use **embeddings** instead (Chapter 14), which capture meaning rather than exact words: "braked" and "braking" look unrelated to TF-IDF but similar to an embedding model.

4.10 Features from dates and times

A raw timestamp is rarely useful. Extract the parts that drive behaviour:

```
import numpy as np, pandas as pd

df = pd.DataFrame({"timestamp": pd.to_datetime(["2026-03-02 08:15", "2026-03-07 23:40"])})
df["day_of_week"] = df["timestamp"].dt.dayofweek      # Monday = 0 ... Sunday = 6
df["hour"] = df["timestamp"].dt.hour
df["is_weekend"] = df["day_of_week"] >= 5
# cyclical encoding: hour 23 should be close to hour 0
df["hour_sin"] = np.sin(2 * np.pi * df["hour"] / 24)
df["hour_cos"] = np.cos(2 * np.pi * df["hour"] / 24)
print(df)
```

Why sine and cosine? As plain numbers, 23 and 0 look far apart, but 11 pm and midnight are neighbours. Placing hours around a circle fixes that. Other useful time features: month, holiday flags, and "time since" features (days since last purchase, hours since last break).

4.11 Categories with thousands of values

One-hot encoding a column like "city" or "product ID" with 10,000 values creates 10,000 columns. Alternatives:

- **Frequency encoding:** replace each category with how often it appears.
- **Target encoding:** replace each category with the average label for that category (e.g. the average survival rate per ticket type). Powerful, but it leaks badly if computed on the same rows you train on. scikit-learn's `TargetEncoder` avoids this by computing it with internal cross-validation.
- **Grouping rare values:** merge categories seen fewer than, say, 20 times into "other".
- **Learned embeddings:** in neural networks, give each category its own small vector learned during training (the same idea as word embeddings in Chapter 14).

4.12 Imbalanced labels

In many real problems the interesting class is rare: fraud, equipment failure, a serious crash. Three habits from the start: split with `stratify=y` so the rare class appears in every split; never judge the model by accuracy alone (Chapter 9); and consider class weights or resampling during training (Chapter 8.7).

4.13 Label quality

Labels usually come from people, and people disagree and make mistakes. If two experts label the same 100 examples and agree only 80% of the time, no model can reliably exceed that level of agreement with either of them.

- Write clear **labelling guidelines** with examples of hard cases.
- Have a sample labelled twice and measure **inter-annotator agreement** (Cohen's kappa, Chapter 17).
- Review the examples your model gets wrong: many will turn out to be labelling errors.

The idea that improving data often beats improving the model is called **data-centric AI**.

4.14 How much data do I need?

There's no fixed number; it depends on how complex the pattern is and how noisy the data is. Rough guidance: simple tabular models can work with hundreds to thousands of rows; deep learning from scratch usually needs tens of thousands or more; transfer learning and pre-trained models (Chapters 12 and 16) cut that dramatically. The reliable way to find out is a **learning curve** (Chapter 8.7): if the validation score is still rising as you add data, more data will help.

Try it

1. Instead of the median, fill missing ages with the median age *per ticket class* (`df.groupby("pclass")["age"].transform("median")`).
2. Create a feature `is_alone` that is 1 when `family_size == 1`.
3. Find a leakage risk: if you were predicting hospital readmission, which of these features would leak — age, diagnosis, "number of days readmitted"?

Chapter 5: Supervised learning

Supervised learning means learning from labelled examples: every training row comes with the right answer. It covers two task types — **regression** (predict a number) and **classification** (predict a category).

5.1 The four ways machines learn

Type	Learns from	Example
Supervised	Inputs + correct answers	Spam filter, price prediction
Unsupervised	Inputs only; finds structure	Customer segments (Chapter 7)
Self-supervised	Inputs that label themselves	An LLM predicting the next word of real text (Part 4)
Reinforcement	Rewards from trial and error	Game-playing agents (Chapter 20)

5.2 Regression: predicting numbers

You met **linear regression** in Chapter 3. With many features it becomes one weight per feature:

$$\hat{y} = w_1x_1 + w_2x_2 + \dots + w_nx_n + b$$

Ridge and **Lasso** are linear regression plus a penalty that keeps weights small, which fights overfitting (Chapter 8). Lasso can push useless weights all the way to zero.

```
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.metrics import mean_absolute_error

X, y = load_diabetes(return_X_y=True)  # predict disease progression from 10
```

```

measurements
X_tr, X_te, y_tr, y_te = train_test_split(X, y, test_size=0.2, random_state=42)

for model in [LinearRegression(), Ridge(alpha=1.0), Lasso(alpha=0.1)]:
    model.fit(X_tr, y_tr)
    print(type(model).__name__, mean_absolute_error(y_te, model.predict(X_te)))

```

5.3 Classification: predicting categories

Logistic regression, despite its name, is a classifier. It computes a linear score, then squashes it into a probability between 0 and 1 with the **sigmoid** function:

$$P(y = 1) = \sigma(z) = \frac{1}{1 + e^{-z}}, \quad z = w \cdot x + b$$

If the probability is above 0.5 (the **decision threshold**), predict class 1. Its loss is **cross-entropy** (log loss): it punishes being confidently wrong very hard. The same loss trains LLMs.

For more than two classes, **softmax** generalizes sigmoid: it turns a score per class into probabilities that sum to 1.

5.4 Four classic classifiers

- **k-nearest neighbours (kNN)** — no real training; to classify a new point, look at the k most similar training examples and take a vote. Simple, but slow on big data and needs scaled features.
- **Support vector machine (SVM)** — finds the boundary with the widest possible gap (the *margin*) between classes. With the **kernel trick** it can draw curved boundaries.
- **Naive Bayes** — uses probability (Bayes' theorem) and assumes features are independent. Fast, a classic for text classification.
- **Logistic regression** — the reliable, interpretable baseline. Always try it first.

5.5 Compare them on the Titanic

Use `X_train_s`, `X_test_s`, `y_train`, `y_test` from Chapter 4.

```

from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB

models = {
    "Logistic regression": LogisticRegression(),
    "k-nearest neighbours": KNeighborsClassifier(n_neighbors=5),
    "SVM": SVC(),
    "Naive Bayes": GaussianNB(),
}
for name, model in models.items():
    model.fit(X_train_s, y_train)
    print(f"{name:22s} accuracy: {model.score(X_test_s, y_test):.3f}")

```

Expect roughly 0.81, 0.80, 0.83 and 0.73. Notice every scikit-learn model has the same interface: `fit()` to train, `predict()` to use, `score()` to evaluate. Learn it once, use it everywhere.

5.6 Which model when?

Situation	Good first choice
Need to explain every decision	Logistic / linear regression
Tabular data, want best accuracy	Gradient boosting (Chapter 6)
Short texts, little data	Naive Bayes, logistic regression
Small dataset, complex boundary	SVM
Images, audio, long text	Deep learning (Part 3)

5.7 Logistic regression, worked through

Suppose a trained model for Titanic survival has learned this score (the numbers are illustrative):

$$z = 1.0 + 2.5 \times \text{is_female} - 1.0 \times \text{ticket_class}$$

- A woman in 1st class: $z = 1.0 + 2.5 - 1.0 = 2.5$, and $\sigma(2.5) \approx \mathbf{0.92}$. Predict survived.
- A man in 3rd class: $z = 1.0 + 0 - 3.0 = -2.0$, and $\sigma(-2.0) \approx \mathbf{0.12}$. Predict died.

Reading the weights: each weight is a change in the *log-odds*. Raising e to the weight gives an **odds ratio**: $e^{2.5} \approx 12$, so being female multiplies the odds of survival by about 12, other things equal. This interpretability is why logistic regression is popular in medicine, finance and insurance.

Why cross-entropy? For one example with true label y (0 or 1) and predicted probability p :

$$\text{loss} = -[y \log p + (1 - y) \log(1 - p)]$$

If the woman really survived, the loss is $-\log(0.92) \approx 0.08$: small. If the man had actually survived, the loss would be $-\log(0.12) \approx 2.1$: about 25 times larger. Being confidently wrong is punished much harder than being unsure, which pushes the model towards honest probabilities.

5.8 Decision boundaries

A classifier divides the feature space into regions, one per class. The line between them is the **decision boundary**, and each model family draws a characteristic shape:

Model	Boundary shape
Logistic regression	A straight line (a flat plane in more dimensions)
kNN	Jagged, follows individual training points
SVM with RBF kernel	Smooth curves
Decision tree	Boxes made of horizontal and vertical cuts
Neural network	Any smooth shape, given enough neurons

```
import matplotlib.pyplot as plt
from sklearn.datasets import make_moons
```

```

from sklearn.inspection import DecisionBoundaryDisplay
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC

X, y = make_moons(n_samples=300, noise=0.25, random_state=0) # two interleaved
half-moons
fig, axes = plt.subplots(1, 2, figsize=(10, 4))
for ax, model in zip(axes, [LogisticRegression(), SVC(kernel="rbf")]):
    model.fit(X, y)
    DecisionBoundaryDisplay.from_estimator(model, X, ax=ax, alpha=0.3)
    ax.scatter(X[:, 0], X[:, 1], c=y, s=10)
    ax.set_title(type(model).__name__)
plt.show()

```

Logistic regression can only draw a straight line through the moons; the SVM curves around them.

5.9 kNN: distance and the choice of k

kNN depends entirely on how "similar" is measured:

- **Euclidean distance:** straight-line distance, the square root of the sum of squared differences. The default.
- **Manhattan distance:** the sum of absolute differences, like walking city blocks.
- **Cosine distance:** compares direction, not size; common for text and embeddings (Chapter 14).

Because distances add up differences across features, an unscaled feature with large numbers (fare) drowns out one with small numbers (class). Always scale before kNN.

Small k (1–3) follows every quirk of the training data and overfits; large k smooths too much and underfits. With many features, kNN also suffers from the **curse of dimensionality**: in high dimensions, all points end up roughly equally far apart, so "nearest" stops meaning much.

5.10 SVMs: margins and the kernel trick

An SVM looks for the boundary with the widest empty corridor (the **margin**) between the classes. Only the training points touching the corridor edges matter; they're called **support vectors**, hence the name.

- **C** controls how strict the margin is. Large C: few training mistakes allowed, narrow margin, risk of overfitting. Small C: wider margin, some mistakes tolerated, smoother boundary.
- The **kernel trick** lets an SVM behave as if the data had been mapped into a much higher-dimensional space where a straight cut works, without ever computing that space. The **RBF kernel** is the usual choice; its **gamma** setting controls how far each point's influence reaches (large gamma = wiggly boundary).

SVMs shine on small and medium datasets with clear structure; they become slow beyond roughly a hundred thousand rows.

5.11 Naive Bayes, by hand

Say 40% of emails are spam. The word "free" appears in 50% of spam emails but only 5% of normal ones. An email arrives containing "free". How likely is it spam? Bayes' theorem (Appendix A.3):

$$P(\text{spam} \mid \text{free}) = \frac{0.5 \times 0.4}{0.5 \times 0.4 + 0.05 \times 0.6} = \frac{0.20}{0.23} \approx 0.87$$

For a whole email, naive Bayes multiplies the evidence from every word, assuming the words are independent of each other given the class. That assumption is clearly false ("free" and "offer" go together), which is why it's called *naive* — but it works surprisingly well, trains instantly, and needs little data.

5.12 More than two classes

- **Multi-class** (one label from several, e.g. car / truck / bus): either train one yes/no model per class and pick the most confident (**one-vs-rest**), or use a model with a **softmax** output that produces one probability per class directly. scikit-learn chooses automatically.
- **Multi-label** (several labels can be true at once, e.g. tags on a news article): train one independent yes/no output per label, each with its own sigmoid.

Try it

1. Change `n_neighbors` in kNN to 1, 15 and 50. Which is best? Why is 1 risky?
2. Train kNN on *unscaled* `X_train` and compare. How much does scaling matter?
3. Print `LogisticRegression().fit(X_train_s, y_train).coef_`. Which feature pushes survival up most?

Chapter 6: Trees and ensembles

For tabular data (spreadsheets, databases), tree ensembles — random forests and gradient boosting — are usually the most accurate models, often beating deep learning.

6.1 Decision trees: twenty questions

A **decision tree** predicts by asking a series of yes/no questions: "Is the passenger male? Is their age under 4? Are they in 3rd class?" Each question splits the data; each final **leaf** gives a prediction.

Training picks, at each step, the question that best separates the classes, measured by **Gini impurity** or **entropy** (how mixed the classes are in each group).

```
from sklearn.tree import DecisionTreeClassifier, export_text

tree = DecisionTreeClassifier(max_depth=3, random_state=42)
tree.fit(X_train, y_train)          # trees don't need scaled features
print(tree.score(X_test, y_test))  # about 0.81
print(export_text(tree, feature_names=list(X_train.columns)))
```

The printout shows the learned questions. The first split is `sex_male` — the tree rediscovered "women and children first". Trees are very **interpretable**, but a deep tree memorizes the training data (overfits). `max_depth` limits that.

6.2 Ensembles: the wisdom of crowds

An **ensemble** combines many models so their errors cancel out. Two main strategies:

- **Bagging** (bootstrap aggregating): train many models in parallel, each on a random sample of the data, then average their votes. Reduces variance.
- **Boosting**: train models one after another, each focusing on the mistakes of the ones before. Reduces bias.

6.3 Random forests

A **random forest** is bagging with decision trees, plus a twist: each split may only consider a random subset of features. That makes the trees different from each other, and the average far more robust than any one tree.

```
from sklearn.ensemble import RandomForestClassifier
import pandas as pd
```

```
forest = RandomForestClassifier(n_estimators=300, random_state=42)
forest.fit(X_train, y_train)
print(forest.score(X_test, y_test))    # about 0.82
```

```
importance = pd.Series(forest.feature_importances_, index=X_train.columns)
print(importance.sort_values(ascending=False))    # fare, age, sex at the top
```

Feature importance tells you which inputs the model relies on most — useful for explaining results and for spotting leakage (a suspiciously dominant feature).

6.4 Gradient boosting

Gradient boosting builds shallow trees one at a time. Each new tree predicts the *remaining errors* of the ensemble so far, and is added with a small weight (the learning rate). It's gradient descent (Chapter 3) done with trees.

```
from sklearn.ensemble import HistGradientBoostingClassifier
```

```
gb = HistGradientBoostingClassifier(random_state=42)
gb.fit(X_train, y_train)
print(gb.score(X_test, y_test))    # about 0.83
```

In industry and Kaggle competitions, the go-to libraries are **XGBoost**, **LightGBM** and **CatBoost** (pip install xgboost). They're faster and have more options, but work the same way.

6.5 Key settings (hyperparameters)

Setting	Effect of increasing it
n_estimators (number of trees)	More stable; slower. For boosting, too many can overfit
max_depth	Trees can learn more complex patterns; more overfitting risk
min_samples_leaf	Each leaf must cover more examples; smoother, less overfitting
learning_rate (boosting)	Faster learning; less precise. Small rate + more

Setting	Effect of increasing it
	trees usually wins

Chapter 8 shows how to choose these automatically.

6.6 Gini impurity, worked example

Gini impurity measures how mixed a group is: the chance that two randomly picked members belong to different classes. For class proportions p_k :

$$\text{Gini} = 1 - \sum_k p_k^2$$

A pure group (all one class) has Gini 0; a 50/50 group has the maximum, 0.5. Take 10 passengers: 6 died, 4 survived.

- **Before splitting:** $\text{Gini} = 1 - (0.6^2 + 0.4^2) = 1 - 0.52 = \mathbf{0.48}$.
- **Split by sex.** Women (4 people: 1 died, 3 survived): $1 - (0.25^2 + 0.75^2) = 0.375$. Men (6 people: 5 died, 1 survived): $1 - ((5/6)^2 + (1/6)^2) \approx 0.278$.
- **Weighted average after the split:** $0.4 \times 0.375 + 0.6 \times 0.278 \approx \mathbf{0.317}$.

The split reduces impurity by $0.48 - 0.317 \approx 0.16$. The tree computes this reduction for every possible question and picks the largest. **Entropy**, the alternative, is $-\sum p_k \log_2 p_k$ (Appendix A.5); it usually picks the same splits.

6.7 How a tree grows

1. For every feature, sort its values and try a threshold between each pair of neighbours ("age < 9.5?", "fare < 26.3?"). For categories, try groupings.
2. Compute the weighted impurity of the two resulting groups. Keep the best question overall.
3. Repeat inside each group.
4. Stop when a group is pure, too small (`min_samples_leaf`) or the tree is deep enough (`max_depth`).

A leaf predicts the majority class (and the class proportions as probabilities). **Regression trees** work the same way but measure impurity by variance, and each leaf predicts the average target of its examples.

Trees need no scaling, handle mixed feature types, and capture interactions automatically ("young AND first class"). Their weaknesses: a single tree changes a lot if the data changes a little, and a tree **can't extrapolate** — asked about a house twice as large as any it has seen, it predicts the value of the largest houses it knows.

6.8 Bootstrap sampling and out-of-bag scores

Each tree in a random forest trains on a **bootstrap sample**: n rows drawn *with replacement* from the n training rows. Some rows appear several times and, on average, about 37% don't appear at all.

Those left-out rows are **out-of-bag (OOB)** for that tree, and each row can be scored by the trees that never saw it. The result is a free validation estimate, without a separate split:

```
forest = RandomForestClassifier(n_estimators=300, oob_score=True, random_state=42)
forest.fit(X_train, y_train)
print(forest.oob_score_) # close to the cross-validation score
```

6.9 Gradient boosting, step by step

Predict three house prices: 200, 300 and 400 (thousand).

1. **Start** with the simplest prediction: the mean, 300 for every house. Errors (residuals) are -100, 0, +100.
2. **Tree 1** is trained to predict those residuals. Suppose it learns them perfectly: -100, 0, +100.
3. **Add a fraction** of it (learning rate 0.1): predictions become 290, 300, 310. The new residuals are -90, 0, +90.
4. **Tree 2** fits the new residuals, and so on. Each tree fixes part of what is still wrong.

After a few hundred trees the predictions approach the truth. Why only add a fraction each time? Small steps are more cautious; together with many trees they generalize better than a few big steps. For classification, the "residuals" are the gradients of the log loss, which is where the name comes from.

Use early stopping so you don't have to guess the number of trees:

`HistGradientBoostingClassifier(early_stopping=True)` holds out some data and stops once the validation score stops improving.

6.10 XGBoost, LightGBM and CatBoost

Library	Known for
XGBoost	The pioneer: regularized boosting, robust, widely used in industry
LightGBM	Very fast on large data; grows trees leaf by leaf
CatBoost	Handles categorical features natively and well, with little tuning
scikit-learn HistGradientBoosting	Built in, fast, handles missing values; inspired by LightGBM

On most tabular problems, all four land within a point or two of each other after tuning. Pick one and spend the effort on features and evaluation instead.

Try it

1. Train a tree with no `max_depth`. Compare its training score (`tree.score(X_train, y_train)`) with its test score. What do you see?
2. Plot the tree visually with `from sklearn.tree import plot_tree; plot_tree(tree, feature_names=list(X_train.columns), filled=True)`.
3. Install XGBoost and train `XGBClassifier()` on the same data.

Chapter 7: Unsupervised learning

Unsupervised learning finds structure in data that has no labels. The two main jobs are **clustering** (grouping similar things) and **dimensionality reduction** (compressing many features into a few).

7.1 Clustering with k-means

k-means splits data into k groups. It starts with k random centre points, assigns every point to its nearest centre, moves each centre to the middle of its points, and repeats until nothing moves.

```
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X, _ = make_blobs(n_samples=300, centers=4, random_state=42) # we pretend we don't
know there are 4
km = KMeans(n_clusters=4, n_init=10, random_state=42).fit(X)

plt.scatter(X[:, 0], X[:, 1], c=km.labels_)
plt.scatter(*km.cluster_centers_.T, marker="x", s=200, color="black")
plt.show()
```

Real uses: customer segmentation, grouping similar documents, compressing image colours.

7.2 Choosing k

You usually don't know how many groups exist. Two tools help:

- **Elbow method** — plot `km.inertia_` (total distance from points to their centre) for $k = 1, 2, 3 \dots$. It always falls; pick the k where it stops falling steeply.
- **Silhouette score** — from -1 to 1 : how much closer each point is to its own cluster than to the next one. Higher is better.

```
from sklearn.metrics import silhouette_score

for k in range(2, 7):
    km = KMeans(n_clusters=k, n_init=10, random_state=42).fit(X)
    print(k, round(km.inertia_), round(silhouette_score(X, km.labels_), 3))
# k=4 gives the best silhouette (about 0.79)
```

7.3 Other clustering methods

Method	Idea	Good when
k-means	Nearest centre	Round, similar-sized groups
Hierarchical	Merge closest groups step by step into a tree (dendrogram)	You want to see groups at every level
DBSCAN	Dense regions are clusters; sparse points are noise	Odd shapes, outliers, unknown k
Gaussian mixture	Each cluster is a bell curve; soft membership	Overlapping groups

7.4 Dimensionality reduction with PCA

A handwritten-digit image of 8×8 pixels has 64 features. Humans can't picture 64 dimensions. **Principal Component Analysis (PCA)** finds the few directions along which the data varies most, and projects onto them.

```
from sklearn.datasets import load_digits
from sklearn.decomposition import PCA

digits = load_digits()                # 1,797 images, 64 features each
pca = PCA(n_components=2)
X2 = pca.fit_transform(digits.data)
print(pca.explained_variance_ratio_.sum()) # about 0.29 of the information kept in
just 2 numbers

plt.scatter(X2[:, 0], X2[:, 1], c=digits.target, cmap="tab10", s=8)
plt.colorbar(); plt.show()
```

Even in 2D, the digits form visible groups. Uses of PCA: visualization, removing noise, and speeding up other models.

For visualization only, **t-SNE** and **UMAP** usually give clearer pictures than PCA because they preserve neighbourhoods rather than straight-line directions.

7.5 Anomaly detection

A close cousin: finding rows that don't fit any pattern — fraud, faulty sensors, network attacks.

`sklearn.ensemble.IsolationForest` is the standard starting point.

7.6 k-means, step by step

Six points on a line: 1, 2, 3, 10, 11, 12. Find $k = 2$ clusters, starting (badly) with centres at 1 and 2.

1. **Assign** each point to the nearest centre: {1} goes to centre 1; {2, 3, 10, 11, 12} go to centre 2.
2. **Move** each centre to the mean of its points: centre 1 = 1; centre 2 = $(2 + 3 + 10 + 11 + 12) / 5 = 7.6$.
3. **Assign again**: 1, 2 and 3 are now closer to centre 1; 10, 11 and 12 to centre 2.
4. **Move again**: centre 1 = 2, centre 2 = 11.
5. **Assign again**: nothing changes. Done — the clusters are {1, 2, 3} and {10, 11, 12}.

What k-means minimizes is the **inertia**: the sum of squared distances from each point to its centre.

Its limits: you must choose k ; it assumes roughly round clusters of similar size; it's thrown off by outliers and unscaled features; and different random starts can give different answers. scikit-learn softens the last problem with **k-means++** (spreading the initial centres out) and `n_init=10` (run ten times, keep the best).

7.7 Hierarchical clustering and linkage

Agglomerative clustering starts with every point as its own cluster and repeatedly merges the two closest clusters until one remains. The full merge history is drawn as a tree, a **dendrogram**; cutting it at a chosen height gives the clusters.

"Closest clusters" can be defined in several ways, called **linkage**:

Linkage	Distance between two clusters	Tends to produce
Single	Their two closest points	Long chains; sensitive to noise
Complete	Their two farthest points	Compact clusters
Average	Average of all point pairs	A compromise
Ward	Increase in total variance if merged	Even, round clusters (like k-means); the usual default

```
from scipy.cluster.hierarchy import linkage, dendrogram
import matplotlib.pyplot as plt

Z = linkage(X[:50], method="ward")  # first 50 points of the blobs from 7.1
dendrogram(Z); plt.show()
```

7.8 DBSCAN's two settings

DBSCAN defines clusters as dense regions. It has two settings:

- **eps**: the radius of a neighbourhood.
- **min_samples**: how many points must be within eps for a point to count as a dense **core point**.

Core points that are within eps of each other join the same cluster; points near a core point but not dense themselves are **border points**; everything else is **noise**, labelled -1. DBSCAN finds the number of clusters by itself and handles odd shapes that defeat k-means:

```
from sklearn.datasets import make_moons
from sklearn.cluster import DBSCAN, KMeans

Xm, _ = make_moons(n_samples=300, noise=0.05, random_state=0)
print(KMeans(n_clusters=2, n_init=10).fit_predict(Xm)[:10])  # cuts each moon in half
print(DBSCAN(eps=0.2, min_samples=5).fit_predict(Xm)[:10])  # follows each moon
To pick eps, plot every point's distance to its k-th nearest neighbour in sorted order and look for the bend. DBSCAN struggles when clusters have very different densities.
```

7.9 PCA, intuitively

Picture a cloud of points in 2-D stretched along a diagonal, like a cigar. The **first principal component** is the direction along the cigar, where the data varies most. The second is perpendicular to it, where the data varies least. Projecting every point onto the first direction keeps most of the information in one number instead of two.

In more dimensions it's the same: each next component is the direction of the most remaining variation, perpendicular to all previous ones. Mathematically, the components are the **eigenvectors** of the data's covariance matrix, and each one's **eigenvalue** is the variance it captures.

Practical points: standardize features first (otherwise the feature with the biggest numbers dominates); plot cumulative `explained_variance_ratio_` to choose how many components to keep; and remember each component is a mix of the original features, so it's harder to interpret.

7.10 Reading t-SNE and UMAP plots carefully

These plots are great for spotting that groups exist, and easy to over-read:

- The **size** of a cluster in the plot means little; dense and sparse groups can look the same.
- The **distance between clusters** often means little; two far-apart blobs aren't necessarily very different.
- Settings (t-SNE's **perplexity**, UMAP's **n_neighbors**) and the random seed can change the picture a lot. Try several.
- Use them for exploration, not as features for a model or as proof of anything.

Try it

1. Run k-means on the digits with `k = 10` and compare `km.labels_` with `digits.target`. Did it find the digits without being told?
2. Try `PCA(n_components=0.95)`: how many components keep 95% of the variance?
3. Replace PCA with `from sklearn.manifold import TSNE; TSNE(n_components=2).fit_transform(digits.data)` and compare the plot.

Chapter 8: Training classical models well

Training a model is one line (`model.fit`). Training it *well* — so it works on data it has never seen — is the real skill. This chapter is the full workflow.

8.1 The training workflow, end to end

1. Define the problem and the metric that matters (Chapter 9).
2. Collect data, then split off a test set and lock it away.
3. Explore and clean the training data (Chapter 4).
4. Build a **baseline**: the simplest reasonable model (or even "always predict the most common class").
5. Try a few model families with cross-validation.
6. Tune the best one's hyperparameters.
7. Evaluate once on the test set.
8. Retrain on all the data, save it, deploy it (Part 5), monitor it.

8.2 Overfitting and underfitting

- **Underfitting**: the model is too simple to capture the pattern. Bad on training data, bad on test data.
- **Overfitting**: the model memorizes the training data, including its noise. Great on training data, bad on test data.
- The goal is the sweet spot between: it **generalizes**.

This is the **bias–variance trade-off**. *Bias* is error from wrong assumptions (too simple). *Variance* is error from sensitivity to the particular training sample (too complex).

See it happen by fitting curves of different complexity to 30 noisy points:

```
import numpy as np
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
from sklearn.pipeline import make_pipeline
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

rng = np.random.default_rng(0)
x = np.sort(rng.uniform(0, 1, 30)).reshape(-1, 1)
y = np.sin(2 * np.pi * x).ravel() + rng.normal(0, 0.2, 30)
x_tr, x_te, y_tr, y_te = train_test_split(x, y, test_size=0.3, random_state=0)

for degree in [1, 4, 15]:
    m = make_pipeline(PolynomialFeatures(degree), LinearRegression()).fit(x_tr, y_tr)
    print(degree, "train:", round(mean_squared_error(y_tr, m.predict(x_tr)), 3),
          "test:", round(mean_squared_error(y_te, m.predict(x_te)), 3))
```

Degree 1 (a straight line) underfits. Degree 4 fits well on both. Degree 15 gets almost zero training error but a test error hundreds of times worse — textbook overfitting.

Symptom	Diagnosis	Fixes
Train error high, test error high	Underfitting	More complex model, better features, less regularization
Train error low, test error high	Overfitting	More data, simpler model, regularization, early stopping
Both low, close together	Good fit	Ship it (after the test set confirms)

8.3 Regularization

Regularization adds a penalty for complexity to the loss, so the model prefers simpler explanations.

$$L_{\text{total}} = L_{\text{data}} + \lambda \sum_j w_j^2 \quad (\text{L2 / Ridge}) \quad L_{\text{total}} = L_{\text{data}} + \lambda \sum_j |w_j| \quad (\text{L1 / Lasso})$$

λ controls the strength (alpha in scikit-learn, and $C = 1/\lambda$ in logistic regression and SVM). For trees, regularization means limits like `max_depth` and `min_samples_leaf`. In neural networks you'll meet **dropout** and **weight decay** (Chapter 11).

8.4 Cross-validation

A single validation split is noisy: you might just get lucky. **k-fold cross-validation** splits the training data into k parts, trains k times (each time holding out a different part), and averages the scores.

```

from sklearn.model_selection import cross_val_score

scores = cross_val_score(pipe, X_train, y_train, cv=5)  # pipe is built in 8.5
print(scores.round(3), "mean:", scores.mean().round(3), "±", scores.std().round(3))
# e.g. [0.79 0.72 0.824 0.838 0.803] mean: 0.795 ± 0.041

```

The spread ($\pm$) tells you how much to trust the difference between two models. Variants: **stratified k-fold** for classification (the default), **group k-fold** when rows from one customer must stay together, and **time-series split** when the future must never train the past.

8.5 Pipelines: preprocessing without leakage

A **pipeline** chains preprocessing and the model into one object. Every step is fitted on the training folds only, so cross-validation can't leak. This version starts from the *raw* Titanic data:

```

import seaborn as sns
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split

raw = sns.load_dataset("titanic")[["survived", "pclass", "sex", "age", "sibsp",
                                   "fare", "embarked"]]
X, y = raw.drop(columns="survived"), raw["survived"]
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
                                                    random_state=42, stratify=y)

numeric = ["pclass", "age", "sibsp", "fare"]
categorical = ["sex", "embarked"]
prep = ColumnTransformer([
    ("num", make_pipeline(SimpleImputer(strategy="median"), StandardScaler()),
     numeric),
    ("cat", make_pipeline(SimpleImputer(strategy="most_frequent"),
                         OneHotEncoder(handle_unknown="ignore")), categorical),
])
pipe = Pipeline([("prep", prep), ("model", RandomForestClassifier(random_state=42))])

```

8.6 Hyperparameter tuning

Parameters are learned by training (weights, tree splits). **Hyperparameters** are settings you choose before training (tree depth, learning rate, k). Tuning means searching for the best hyperparameters using cross-validation.

```

from sklearn.model_selection import GridSearchCV

grid = GridSearchCV(pipe, {
    "model__n_estimators": [100, 300],
    "model__max_depth": [3, 5, None],
    "model__min_samples_leaf": [1, 5],
}, cv=5, n_jobs=-1)
grid.fit(X_train, y_train)
print(grid.best_params_, grid.best_score_)  # best CV score about 0.82
print(grid.score(X_test, y_test))          # final, honest test score about 0.80

```

- **Grid search** tries every combination — thorough, slow.
- **Random search** (RandomizedSearchCV) samples combinations — usually as good, much faster.
- **Bayesian optimization** (the **Optuna** library) learns which regions look promising — best for expensive models.

Notice the test score is slightly below the CV score. That's normal: tuning picks the setting that happened to do best on those folds, which flatters it a little.

8.7 Other training issues

- **Class imbalance** (e.g. 1% fraud): use `class_weight="balanced"`, resample the data (oversample the rare class or undersample the common one), and judge with precision/recall, not accuracy.
- **Learning curves**: plot score vs training-set size (`sklearn.model_selection.learning_curve`). If test score still rises at the end, more data will help.
- **Reproducibility**: fix `random_state`, record library versions and the exact data used.
- **Saving a model**: `import joblib; joblib.dump(grid.best_estimator_, "model.joblib")`, then `joblib.load("model.joblib")` to use it later. Chapter 22 serves this file from a Docker container.

8.8 L1 vs L2: why one creates zeros

Both penalties discourage large weights, but differently:

- **L2 (Ridge)** penalizes the *square* of each weight. The push towards zero is strong for big weights and fades as a weight gets small. Result: all weights shrink, but few become exactly zero.
- **L1 (Lasso)** penalizes the *absolute value*. The push towards zero is the same size however small the weight is, so weak features get pushed all the way to 0. Result: automatic **feature selection**.
- **Elastic Net** mixes both.

```
from sklearn.datasets import load_diabetes
from sklearn.linear_model import Lasso, Ridge

X, y = load_diabetes(return_X_y=True)
for alpha in [0.01, 0.1, 1.0]:
    lasso = Lasso(alpha=alpha).fit(X, y)
    ridge = Ridge(alpha=alpha).fit(X, y)
    print(alpha, "Lasso zero weights:", (lasso.coef_ == 0).sum(), "| Ridge zero
weights:", (ridge.coef_ == 0).sum())
```

As alpha grows, Lasso switches off more and more of the 10 features; Ridge keeps them all.

8.9 Early stopping

Models trained step by step (gradient boosting, neural networks) start by learning real patterns and then begin memorizing noise. **Early stopping** watches the validation score after each round and stops once it hasn't improved for a set number of rounds (the **patience**), then keeps the best version. It works as regularization because it limits how far the weights can move from their simple starting point, and it saves training time.

8.10 Nested cross-validation

If you tune hyperparameters with cross-validation and then report that same best score, it's optimistic: you picked the setting that happened to look best on those folds. **Nested cross-validation** separates the two jobs. An inner loop tunes; an outer loop measures how well the whole "tune and train" procedure performs on data it never saw.

```
from sklearn.model_selection import GridSearchCV, cross_val_score

inner = GridSearchCV(pipe, {"model__max_depth": [3, 5, None]}, cv=3) # pipe from 8.5
outer_scores = cross_val_score(inner, X_train, y_train, cv=5)        # tuning
repeated inside each outer fold
print(outer_scores.mean())    # an honest estimate for the tuned model
```

Here it trains the model about 50 times instead of once, so it's mostly used when data is small and an honest estimate matters, as in research or medical work.

8.11 Resampling and SMOTE

For rare classes, you can rebalance the *training* data:

- **Random undersampling:** drop examples of the common class. Fast; throws data away.
- **Random oversampling:** duplicate examples of the rare class. Can overfit to those exact rows.
- **SMOTE:** creates new synthetic rare-class examples by picking a rare example and one of its rare-class neighbours, and placing a point somewhere on the line between them.

Resample only inside the training folds, never the validation or test data, which must keep their real proportions. The imbalanced-learn library provides a pipeline that does this correctly:

```
# pip install imbalanced-learn
from imblearn.pipeline import Pipeline as ImbPipeline
from imblearn.over_sampling import SMOTE

imb_pipe = ImbPipeline([("prep", prep), ("smote", SMOTE(random_state=42)),
                        ("model", RandomForestClassifier(random_state=42))])
# use it exactly like pipe: cross_val_score(imb_pipe, X_train, y_train, cv=5,
scoring="f1")
```

Often `class_weight="balanced"` plus a well-chosen threshold (Chapter 9.4) works just as well, with less complexity.

8.12 When your model is bad: a debugging checklist

1. **Check the data.** Print shapes, look at rows, check for leakage (a suspiciously perfect score is a warning sign) and mislabelled examples.
2. **Compare with a baseline.** Is the model actually beating "predict the most common class" or "predict the average"?
3. **Compare training and validation scores.** Big gap: overfitting. Both poor: underfitting or uninformative features.
4. **Read the errors.** Look at 20–50 wrong predictions. Are they hard cases, labelling mistakes, or a missing feature?

5. **Check feature importance.** Do the top features make sense to someone who knows the domain?
6. **Draw a learning curve.** Would more data help, or is the model capacity the limit?
7. **Simplify.** Remove features and steps until it works, then add things back one at a time.
8. **Fix the randomness.** Set `random_state` everywhere so a change in score comes from your change, not from luck.

Try it

1. Replace `GridSearchCV` with `RandomizedSearchCV(pipe, ..., n_iter=10)`. Is it as good? How much faster?
2. Swap the model in the pipeline for `LogisticRegression(max_iter=1000)` and tune `model__C` over `[0.01, 0.1, 1, 10]`.
3. Plot a learning curve for the pipeline. Would more passengers help?

Chapter 9: Evaluating classical models

A single accuracy number can hide a useless model. Evaluation means choosing the metric that matches the real cost of mistakes, comparing against a baseline, and testing on truly unseen data.

9.1 Always start with a baseline

```
from sklearn.dummy import DummyClassifier

dummy = DummyClassifier(strategy="most_frequent").fit(X_train, y_train)
print(dummy.score(X_test, y_test))  # about 0.62: "everyone dies" is 62% accurate!
```

Our tuned forest scores about 0.80. The honest improvement is $0.62 \rightarrow 0.80$, not $0 \rightarrow 0.80$. If 99% of transactions are legitimate, a model that never flags fraud is 99% accurate and worthless.

9.2 The confusion matrix

For a yes/no classifier, every prediction is one of four outcomes:

	Predicted: no	Predicted: yes
Actually no	True negative (TN)	False positive (FP) — a false alarm
Actually yes	False negative (FN) — a miss	True positive (TP)

```
from sklearn.metrics import confusion_matrix, classification_report

best = grid.best_estimator_  # from Chapter 8
pred = best.predict(X_test)
print(confusion_matrix(y_test, pred))  # [[104  6]
                                       # [ 29 40]]
print(classification_report(y_test, pred, digits=3))
```

9.3 Classification metrics

Metric	Question it answers	Formula
--------	---------------------	---------

Metric	Question it answers	Formula
Accuracy	What fraction of all predictions were right?	$(TP + TN) / \text{all}$
Precision	When it says "yes", how often is it right?	$TP / (TP + FP)$
Recall (sensitivity)	Of the real "yes" cases, how many did it catch?	$TP / (TP + FN)$
Specificity	Of the real "no" cases, how many did it clear?	$TN / (TN + FP)$
F1 score	A single balance of precision and recall	$2 \times P \times R / (P + R)$
ROC-AUC	How well does it rank positives above negatives, across all thresholds?	Area under the ROC curve; 0.5 = random, 1 = perfect
PR-AUC	Same idea, focused on the rare positive class	Area under the precision–recall curve
Log loss	Are its probabilities trustworthy?	Cross-entropy; lower is better

Our model: precision 0.87 but recall only 0.58 for survivors — when it predicts survival it's usually right, but it misses 42% of actual survivors.

Which to prioritize depends on the cost of errors:

- **Recall** when misses are expensive: cancer screening, fraud, a driver falling asleep.
- **Precision** when false alarms are expensive: spam filters (losing a real email), recommending a costly intervention.
- **F1** or **PR-AUC** when classes are imbalanced and both matter.
- For **multi-class** problems, report per-class scores and a *macro* average (every class equal) or *weighted* average (by class size).

9.4 The threshold is a choice

The classifier outputs a probability; 0.5 is just the default cut-off. Moving it trades precision against recall:

```
from sklearn.metrics import precision_score, recall_score, roc_auc_score

proba = best.predict_proba(X_test)[: , 1]
print("ROC-AUC:", roc_auc_score(y_test, proba))    # about 0.84
for t in [0.3, 0.5, 0.7]:
    p = (proba >= t).astype(int)
    print(t, "precision", round(precision_score(y_test, p), 2), "recall",
          round(recall_score(y_test, p), 2))
```

```
# 0.3 -> precision 0.66, recall 0.81
# 0.5 -> precision 0.87, recall 0.58
```

Pick the threshold on validation data according to the business cost, never on the test set.

9.5 Regression metrics

Metric	Meaning	Note
MAE (mean absolute error)	Average size of the error, in the target's units	Easy to explain: "off by \$12,000 on average"
RMSE (root mean squared error)	Like MAE but punishes big errors more	Use when large misses are especially bad
MAPE	Average error as a percentage	Breaks when true values are near zero
R ²	Fraction of variance explained; 1 = perfect, 0 = no better than predicting the mean	Can be negative for a terrible model

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
# mae = mean_absolute_error(y_true, y_pred); rmse = mean_squared_error(y_true,
y_pred) ** 0.5; r2 = r2_score(y_true, y_pred)
```

9.6 Evaluating clustering

Without labels, use internal scores: **silhouette** (higher is better), **Davies–Bouldin** (lower is better). With labels, **adjusted Rand index** compares clusters with the truth. Most importantly: do the clusters make sense to a domain expert?

9.7 Beyond the metric

- **Calibration**: when the model says 70%, does it happen 70% of the time? Check with `sklearn.calibration.CalibrationDisplay`.
- **Slice analysis**: compute metrics per group (age band, region, device). A good average can hide a group where the model fails (Chapter 24).
- **Error analysis**: read 20–50 wrong predictions by hand. It's the fastest way to find bad labels, leakage or missing features.
- **Statistical confidence**: small test sets give noisy scores. A 1-point difference on 178 passengers is probably noise; use cross-validation spreads or bootstrap confidence intervals.
- **Online evaluation**: after deployment, measure the real outcome with **A/B tests** and monitor for **drift** (Chapter 23).

9.8 Every metric from one confusion matrix

Our tuned forest's test results were TN = 104, FP = 6, FN = 29, TP = 40 (179 passengers). Every classification metric comes from these four numbers:

Metric	Calculation	Result
--------	-------------	--------

Metric	Calculation	Result
Accuracy	$(104 + 40) / 179$	0.80
Precision	$40 / (40 + 6)$	0.87
Recall	$40 / (40 + 29)$	0.58
Specificity	$104 / (104 + 6)$	0.95
False positive rate	$6 / (104 + 6)$	0.05
F1	$2 \times 0.87 \times 0.58 / (0.87 + 0.58)$	0.70

The story they tell together: the model is cautious. It rarely raises a false alarm (specificity 0.95), but it misses many real survivors (recall 0.58). Whether that's acceptable depends on the use.

9.9 How ROC and precision–recall curves are built

A **ROC curve** shows every possible threshold at once:

1. Sort the test examples by predicted probability, highest first.
2. Start with the threshold above every score: nothing is predicted positive (true-positive rate 0, false-positive rate 0).
3. Lower the threshold past each example in turn. Each positive you pass raises the true-positive rate (recall); each negative raises the false-positive rate.
4. Plot recall (y) against false-positive rate (x). A perfect model shoots straight up, then across; a random one follows the diagonal.

ROC-AUC has a neat meaning: the probability that a randomly chosen positive gets a higher score than a randomly chosen negative. With two positives scored 0.9 and 0.6 and two negatives scored 0.7 and 0.2, there are four positive–negative pairs, and the positive wins in three of them ($0.9 > 0.7$, $0.9 > 0.2$, $0.6 > 0.2$). $AUC = 3/4 = 0.75$.

A **precision–recall curve** plots precision against recall over the same thresholds. A random model's precision equals the share of positives, so with 1% fraud a PR-AUC of 0.30 can be excellent while ROC-AUC looks deceptively high for weak models. For rare classes, prefer PR curves.

9.10 Micro, macro and weighted averages

With several classes you get one score per class, and must combine them. Suppose recall is 0.95 for class A (80 examples), 0.50 for B (16 examples) and 0.00 for C (4 examples):

- **Macro average:** plain mean of the class scores: $(0.95 + 0.50 + 0.00) / 3 \approx \mathbf{0.48}$. Every class counts equally, so failing on rare C hurts a lot.
- **Weighted average:** mean weighted by class size: $(80 \times 0.95 + 16 \times 0.50 + 4 \times 0) / 100 = \mathbf{0.84}$. Dominated by big classes.
- **Micro average:** pool every individual prediction first. For recall in single-label problems it equals accuracy, here 0.84.

The gap between 0.48 and 0.84 is the warning: the model ignores class C. Report macro averages when every class matters.

9.11 Log loss, worked

Three predictions:

True label	Predicted P(positive)	Loss = $-\log(\text{probability given to the true label})$
1	0.9	$-\log(0.9) = 0.11$
0	0.2	$-\log(0.8) = 0.22$
1	0.4	$-\log(0.4) = 0.92$

Mean log loss ≈ 0.42 . A model that always says 0.5 scores $\log(2) \approx 0.69$, so anything above 0.69 is worse than a coin. The **Brier score**, the mean squared difference between the probability and the 0/1 outcome, is a gentler alternative.

9.12 Choosing the threshold by cost

When errors have known costs, pick the threshold that minimizes total cost on validation data. Suppose a missed fraud (FN) costs \$500 and reviewing a false alarm (FP) costs \$10:

```
import numpy as np
from sklearn.metrics import confusion_matrix

def total_cost(y_true, proba, threshold, cost_fn=500, cost_fp=10):
    tn, fp, fn, tp = confusion_matrix(y_true, proba >= threshold).ravel()
    return cost_fn * fn + cost_fp * fp

thresholds = np.linspace(0.01, 0.99, 99)
costs = [total_cost(y_val, proba_val, t) for t in thresholds] # use validation data
print("best threshold:", thresholds[int(np.argmin(costs))])
```

If the probabilities are well calibrated, theory gives the answer directly: predict positive whenever $p > \text{cost_FP} / (\text{cost_FP} + \text{cost_FN}) = 10 / 510 \approx 0.02$. With costs this lopsided, flagging even a 2% risk is worth it.

Try it

1. Plot the ROC curve with `from sklearn.metrics import RocCurveDisplay; RocCurveDisplay.from_predictions(y_test, proba)`.
2. Find the threshold that gives recall of at least 0.75. What precision do you get?
3. Compute the forest's accuracy separately for men and women. Is it equally good for both?

Part 3 — Deep learning

Chapter 10: Neural networks and backpropagation

A neural network is many simple units stacked in layers. Each unit does logistic regression's job — a weighted sum and a squashing function — and stacking them lets the network learn patterns no single line can.

10.1 The artificial neuron

A **neuron** takes inputs, multiplies each by a weight, adds a bias, and passes the result through an **activation function**:

$$\text{output} = f(w_1x_1 + w_2x_2 + \dots + w_nx_n + b)$$

The activation adds **non-linearity**. Without it, a stack of layers collapses into one big linear equation and can only draw straight lines.

Activation	Formula (plain)	Where it's used
ReLU	$\max(0, z)$	Default for hidden layers: simple and fast
Sigmoid	$1 / (1 + e^{-z})$, output 0–1	Output for yes/no problems
Tanh	Like sigmoid but –1 to 1	Older networks, RNNs
Softmax	Turns scores into probabilities summing to 1	Output for multi-class problems (and LLMs)
GELU	A smooth ReLU	Transformers

10.2 Layers

A network has an **input layer** (your features), one or more **hidden layers**, and an **output layer** (the prediction). "Deep" learning simply means many hidden layers. Each layer transforms the data into a more useful representation: in an image network, early layers detect edges, middle layers shapes, late layers whole objects.

Passing an input through all layers to get a prediction is the **forward pass**.

10.3 Backpropagation

Training works like Chapter 3: compute the loss, find the gradient, step downhill. The problem is that a deep network has millions of weights spread across layers. How much did a weight in layer 1 contribute to the final error?

Backpropagation answers this with the **chain rule** from calculus. It starts at the loss and works backwards layer by layer, passing along "how much the loss changes if this value changes". One

backward pass gives the gradient for every weight at once. It's the engine behind all deep learning, including LLMs.

10.4 Build one from scratch: solving XOR

XOR outputs 1 when exactly one input is 1. No straight line separates the classes, so logistic regression fails. A network with one hidden layer solves it.

```
import numpy as np

X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]], dtype=float)
y = np.array([[0], [1], [1], [0]], dtype=float)

rng = np.random.default_rng(42)
W1 = rng.normal(0, 1, (2, 4)); b1 = np.zeros((1, 4)) # 2 inputs -> 4 hidden neurons
W2 = rng.normal(0, 1, (4, 1)); b2 = np.zeros((1, 1)) # 4 hidden -> 1 output
sigmoid = lambda z: 1 / (1 + np.exp(-z))
lr = 1.0

for step in range(5000):
    # forward pass
    h = np.tanh(X @ W1 + b1)
    p = sigmoid(h @ W2 + b2)
    loss = -np.mean(y * np.log(p) + (1 - y) * np.log(1 - p)) # cross-entropy

    # backward pass (chain rule, layer by layer)
    dz2 = (p - y) / len(X)
    dW2 = h.T @ dz2; db2 = dz2.sum(0, keepdims=True)
    dz1 = (dz2 @ W2.T) * (1 - h ** 2) # derivative of tanh
    dW1 = X.T @ dz1; db1 = dz1.sum(0, keepdims=True)

    # gradient descent step
    W1 -= lr * dW1; b1 -= lr * db1
    W2 -= lr * dW2; b2 -= lr * db2
    if step % 1000 == 0:
        print(step, round(loss, 4))

print(p.round(3).ravel()) # [0. 0.999 1. 0.001] -> XOR learned
```

That's a complete neural network: forward pass, loss, backward pass, update. Every framework automates exactly these steps.

10.5 Key training concepts

- **Epoch:** one full pass over the training data.
- **Batch:** a small group of examples processed together (e.g. 32). **Mini-batch gradient descent** updates the weights after each batch — faster and less memory than using all data at once.
- **Weight initialization:** weights start small and random. All zeros would make every neuron learn the same thing.
- **Vanishing / exploding gradients:** in very deep networks, gradients can shrink to nothing or blow up as they pass back through many layers. ReLU, careful initialization, normalization layers and **residual connections** (shortcuts that skip layers) fix this — all used in transformers.

- **Universal approximation:** a big enough network can approximate almost any function. The hard part is training it to generalize.

10.6 A forward pass with real numbers

A tiny network: 2 inputs, 2 hidden ReLU neurons, 1 sigmoid output. Input $x = (1, 2)$.

- **Hidden neuron 1**, weights (0.5, 1.0), bias 0: $z = 0.5 \times 1 + 1.0 \times 2 + 0 = 2.5$. $\text{ReLU}(2.5) = \mathbf{2.5}$.
- **Hidden neuron 2**, weights (-1.0, 0.5), bias 0.5: $z = -1.0 + 1.0 + 0.5 = 0.5$. $\text{ReLU}(0.5) = \mathbf{0.5}$.
- **Output neuron**, weights (1.0, -2.0), bias 0.2: $z = 2.5 - 1.0 + 0.2 = 1.7$. $\text{Sigmoid}(1.7) \approx \mathbf{0.85}$.

The network predicts an 85% chance of class 1. In matrix form, the hidden layer is one line, $h = \text{ReLU}(xW_1 + b_1)$, exactly as in the XOR code.

10.7 Backpropagation on one neuron, by hand

Take a single sigmoid neuron with cross-entropy loss. A useful fact: for this pair, the gradient of the loss with respect to z is simply $\mathbf{p - y}$ (prediction minus truth).

Input $x = 2$, true label $y = 1$, starting weight $w = 0.5$, bias $b = 0$, learning rate 1.

1. Forward: $z = 0.5 \times 2 + 0 = 1$; $p = \sigma(1) \approx 0.731$; loss $= -\log(0.731) \approx 0.313$.
2. Backward: $dL/dz = p - y = -0.269$. By the chain rule, $dL/dw = dL/dz \times x = -0.538$ and $dL/db = dL/dz \times 1 = -0.269$.
3. Update: $w = 0.5 + 0.538 = 1.038$; $b = 0 + 0.269 = 0.269$.
4. Check: $z = 1.038 \times 2 + 0.269 \approx 2.35$; $p \approx 0.91$; loss ≈ 0.09 . It dropped from 0.31.

In a deep network the chain just gets longer. For a weight in the first layer:

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial z_3} \cdot \frac{\partial z_3}{\partial h_2} \cdot \frac{\partial h_2}{\partial h_1} \cdot \frac{\partial h_1}{\partial w_1}$$

Backprop computes these products from the output backwards, reusing each layer's result for the layer before it. That reuse is why one backward pass costs only about twice a forward pass, even with billions of weights. It also explains vanishing gradients: multiply many numbers smaller than 1 and the product heads towards 0.

10.8 Why non-linearity is essential

Without activation functions, two layers are $W_2(W_1x) = (W_2W_1)x$ — a single matrix, so a single linear layer. Stacking a hundred of them still gives one straight-line model. The activation between layers is what lets each layer bend the space, so the network can carve out curved boundaries like XOR's. ReLU is popular because it's cheap and its gradient is exactly 1 for positive inputs, so gradients pass through without shrinking.

10.9 Weight initialization

- If weights start too large, activations grow layer after layer and explode; too small, and they shrink to nothing.
- **Xavier (Glorot) initialization** scales random starting weights by roughly $1 / \sqrt{\text{number of inputs}}$, keeping the signal's size steady for tanh and sigmoid.

- **He (Kaiming) initialization** uses roughly $\sqrt{2 / \text{number of inputs}}$, compensating for ReLU switching off half its inputs.
- Weights must be *random*: if all start equal, every neuron in a layer receives the same gradient and stays identical forever (the **symmetry** problem).

PyTorch's layers use sensible defaults, so you rarely set this by hand, but it explains why very deep networks once failed to train.

10.10 How big should a network be?

A fully connected layer with n inputs and m outputs has $n \times m$ weights plus m biases. Chapter 11's first layer, $784 \rightarrow 256$, has $784 \times 256 + 256 = 200,960$ parameters.

- **Width** (neurons per layer) adds capacity within a layer; **depth** (number of layers) lets the network build features out of features, which is usually more efficient.
- A practical recipe: start small, grow the network until it overfits the training data, then add regularization (Chapter 11.5) and more data until validation performance stops improving.

Try it

1. Change the hidden layer from 4 neurons to 2. Does it still learn XOR? Try a few random seeds.
2. Set `lr = 0.01`. How many more steps does it need?
3. Replace `np.tanh` with ReLU (`np.maximum(0, z)`). What must change in the backward pass? (Hint: ReLU's derivative is 1 where $z > 0$, else 0.)

Chapter 11: Training deep networks in PyTorch

PyTorch automates what you did by hand in Chapter 10: it computes gradients for you (**autograd**), runs on GPUs, and provides ready-made layers. It's the framework behind most modern AI research, including most LLMs.

In Colab, switch on a GPU first: *Runtime* → *Change runtime type* → *GPU*. PyTorch is pre-installed.

11.1 Tensors and autograd

A **tensor** is PyTorch's array — like a NumPy array, but it can live on a GPU and remember how it was computed.

```
import torch

w = torch.tensor(3.0, requires_grad=True)  # track gradients for w
x = torch.tensor(2.0)
loss = (w * x - 10) ** 2                  # prediction 6, target 10
loss.backward()                           # backpropagation, automatically
print(w.grad)                             # tensor(-16.): increase w to reduce loss
```

11.2 A complete training script

We'll classify **Fashion-MNIST**: 70,000 small greyscale images of 10 clothing types. For simplicity we use the test set for validation here; in real work, split off a separate validation set.

```

import torch, torch.nn as nn
from torch.utils.data import DataLoader
from torchvision import datasets, transforms

device = "cuda" if torch.cuda.is_available() else "cpu"

# 1. Data: Dataset + DataLoader (batches and shuffling)
tfm = transforms.ToTensor()
train_ds = datasets.FashionMNIST("data", train=True, download=True, transform=tfm)
test_ds = datasets.FashionMNIST("data", train=False, download=True, transform=tfm)
train_dl = DataLoader(train_ds, batch_size=64, shuffle=True)
test_dl = DataLoader(test_ds, batch_size=256)

# 2. Model
model = nn.Sequential(
    nn.Flatten(), # 28x28 image -> 784 numbers
    nn.Linear(784, 256), nn.ReLU(), nn.Dropout(0.2),
    nn.Linear(256, 128), nn.ReLU(), nn.Dropout(0.2),
    nn.Linear(128, 10), # one score per class
).to(device)

# 3. Loss and optimizer
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.AdamW(model.parameters(), lr=1e-3, weight_decay=1e-4)

def evaluate(loader):
    model.eval() # turns dropout off
    correct, total_loss = 0, 0.0
    with torch.no_grad(): # no gradients needed
        for xb, yb in loader:
            xb, yb = xb.to(device), yb.to(device)
            out = model(xb)
            total_loss += loss_fn(out, yb).item() * len(xb)
            correct += (out.argmax(1) == yb).sum().item()
    n = len(loader.dataset)
    return total_loss / n, correct / n

# 4. The training loop
for epoch in range(5):
    model.train()
    for xb, yb in train_dl:
        xb, yb = xb.to(device), yb.to(device)
        loss = loss_fn(model(xb), yb) # forward pass + loss
        optimizer.zero_grad() # clear old gradients
        loss.backward() # backward pass
        optimizer.step() # update weights
    val_loss, val_acc = evaluate(test_dl)
    print(f"epoch {epoch + 1}: val loss {val_loss:.3f}, val accuracy {val_acc:.3f}")

torch.save(model.state_dict(), "fashion_mlp.pt") # save the trained weights

```

Expect roughly 87–89% accuracy after 5 epochs. The five lines inside the loop — forward, loss, zero_grad, backward, step — are the same whether you train this tiny model or a 70-billion-parameter LLM.

11.3 Optimizers

Optimizer	Idea	When
SGD	Plain gradient descent on mini-batches	Simple; still used for some vision models
SGD + momentum	Keeps a running direction, like a rolling ball	Smoother, faster than plain SGD
Adam	Momentum + a separate adaptive step size per weight	The safe default
AdamW	Adam with weight decay done correctly	Standard for transformers and LLMs

11.4 The learning rate: the most important setting

If you tune only one thing, tune the learning rate. Common tools:

- **Learning-rate schedules:** start higher, decay over time (`torch.optim.lr_scheduler.CosineAnnealingLR`).
- **Warmup:** start tiny and ramp up over the first few hundred steps. Essential for transformers.
- **Gradient clipping:** cap gradient size to avoid sudden blow-ups (`torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)`).

11.5 Regularization and stability tools

- **Dropout:** randomly switch off a fraction of neurons during training, so the network can't rely on any one. Off at evaluation time (`model.eval()`).
- **Weight decay:** L2-style penalty on weights (Chapter 8).
- **Early stopping:** stop when validation loss stops improving; keep the best checkpoint.
- **Data augmentation:** create varied copies of training data (flipped, cropped, rotated images).
- **Batch normalization / layer normalization:** rescale activations inside the network to keep training stable. Transformers use layer norm.
- **Mixed precision:** compute in 16-bit numbers for about 2× speed and half the memory (`torch.autocast`).

11.6 Reading loss curves

Always plot training loss and validation loss per epoch.

What you see	Meaning	Action
Both falling	Still learning	Train longer
Train falling, validation rising	Overfitting	Stop early, add regularization or data
Both flat and high	Underfitting or learning rate too	Bigger model, higher learning

What you see	Meaning	Action
	low	rate
Loss jumps to NaN	Learning rate too high, or bad data	Lower the rate, clip gradients, check inputs

11.7 The computation graph behind autograd

Every operation on a tensor with `requires_grad=True` is recorded in a **computation graph**: which values were combined, and how. `loss.backward()` walks that graph from the loss back to every weight, applying the chain rule at each node.

```
import torch

x = torch.tensor([1.0, 2.0, 3.0], requires_grad=True)
y = (x ** 2).sum()          # y = x12 + x22 + x32
y.backward()
print(x.grad)               # tensor([2., 4., 6.]) the derivative 2x

y = (x ** 2).sum()
y.backward()
print(x.grad)               # tensor([4., 8., 12.]) gradients ADD UP across calls
```

That accumulation is why the training loop calls `optimizer.zero_grad()` every step. Two related tools: with `torch.no_grad()`: switches recording off (for evaluation, saving memory and time), and `tensor.detach()` gives a copy cut off from the graph.

11.8 Your own Dataset

For your own data, write a small `Dataset` class with three methods. The `DataLoader` then handles batching and shuffling.

```
import torch, seaborn as sns
from torch.utils.data import Dataset, DataLoader

class TitanicDataset(Dataset):
    def __init__(self):
        df = sns.load_dataset("titanic")[["survived", "pclass", "age",
"fare"]].dropna()
        self.X = torch.tensor(df[["pclass", "age", "fare"]].values,
dtype=torch.float32)
        self.y = torch.tensor(df["survived"].values, dtype=torch.long)

    def __len__(self):
        # how many examples
        return len(self.y)

    def __getitem__(self, i):
        # return example number i
        return self.X[i], self.y[i]

loader = DataLoader(TitanicDataset(), batch_size=32, shuffle=True)
xb, yb = next(iter(loader))
print(xb.shape, yb.shape)          # torch.Size([32, 3]) torch.Size([32])
```

For large datasets, `__getitem__` loads one file at a time from disk (an image, an audio clip), so the full dataset never needs to fit in memory. `DataLoader(..., num_workers=2)` loads batches in parallel while the GPU trains.

11.9 How Adam works

Adam keeps two running averages for every weight: the average gradient (m , the **momentum**) and the average squared gradient (v , how large the gradients typically are):

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t \quad v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2 \quad w \leftarrow w - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

g is the current gradient; the defaults are $\beta_1 = 0.9$ and $\beta_2 = 0.999$; the hats mean a small correction for the first few steps, when the averages start at zero.

In plain words: momentum smooths out noisy mini-batch gradients so the weights keep moving in a consistent direction, and dividing by $\sqrt{v}$ gives every weight its own step size — weights with large, jumpy gradients take smaller steps, rarely updated weights take bigger ones. **AdamW** applies weight decay directly to the weights rather than through the gradient, which works better and is the standard for transformers.

11.10 Choosing the batch size

Larger batches	Smaller batches
Smoother, more accurate gradient per step	Noisier gradients, which can help generalization
Better use of the GPU, faster epochs	Fewer examples processed in parallel
More memory: the usual cause of "CUDA out of memory"	Fits on smaller GPUs
Often needs a larger learning rate	Works with smaller learning rates

Start with 32–128 and use the largest that fits in memory. If you need a bigger effective batch than fits, use **gradient accumulation**: run several small batches, add up their gradients, then take one optimizer step.

11.11 Debugging deep networks

1. **Check the starting loss.** With 10 balanced classes and random weights, cross-entropy should start near $\log(10) \approx 2.30$. Much higher means something's wrong with the inputs or the output layer.
2. **Overfit one batch.** Train repeatedly on a single batch; the loss should fall close to 0. If it can't, there's a bug (wrong labels, a missing `zero_grad`, a frozen layer, an extra softmax before `CrossEntropyLoss`, which applies its own).
3. **Look at your inputs.** Display a few images with their labels after all transforms.
4. **Switch modes correctly.** `model.train()` while training, `model.eval()` when evaluating; dropout and batch norm behave differently in each.

5. **Read the common errors.** "Expected all tensors to be on the same device": move the model and every batch with `.to(device)`. "mat1 and mat2 shapes cannot be multiplied": a layer's input size doesn't match the previous layer's output; print shapes.
6. **Watch for NaN.** Lower the learning rate, clip gradients, and check for division by zero or log of zero in custom code.

```
# overfit-one-batch test
xb, yb = next(iter(train_dl))
xb, yb = xb.to(device), yb.to(device)
for step in range(200):
    loss = loss_fn(model(xb), yb)
    optimizer.zero_grad(); loss.backward(); optimizer.step()
print(loss.item()) # should be close to 0
```

11.12 Saving and resuming training

Save the optimizer state and the epoch along with the weights, so a long run can resume exactly where it stopped:

```
torch.save({"model": model.state_dict(), "optimizer": optimizer.state_dict(),
"epoch": epoch}, "checkpoint.pt")

ckpt = torch.load("checkpoint.pt")
model.load_state_dict(ckpt["model"])
optimizer.load_state_dict(ckpt["optimizer"])
start_epoch = ckpt["epoch"] + 1
```

Try it

1. Record `val_loss` each epoch in a list, train for 20 epochs, and plot it. When does overfitting start?
2. Swap AdamW for `torch.optim.SGD(model.parameters(), lr=0.01)`. How does it compare?
3. Remove the dropout layers and compare training vs validation accuracy.

Chapter 12: Convolutional networks and transfer learning

A **convolutional neural network (CNN)** is built for images. Instead of treating every pixel as unrelated, it slides small pattern detectors across the image, so a feature learned in one corner is recognized everywhere.

12.1 Why flattening images is wasteful

Chapter 11's network flattened each 28×28 image into 784 numbers and lost the idea that neighbouring pixels belong together. A 1000×1000 colour photo would need 3 million inputs per neuron — far too many weights.

12.2 The building blocks

- **Convolution:** a small grid of weights (a **filter** or **kernel**, often 3×3) slides over the image. At each position it computes a weighted sum, producing a **feature map** that lights up where its pattern appears (a vertical edge, a curve, later an eye). The same filter weights are reused at every position — that's what keeps CNNs small.

- **Channels:** a colour image has 3 (red, green, blue). Each conv layer outputs many channels, one per filter.
- **Stride and padding:** stride is how far the filter jumps each step; padding adds a border so edges aren't lost.
- **Pooling:** shrink each feature map, e.g. **max pooling** keeps the largest value in every 2×2 block. Makes the network cheaper and less sensitive to small shifts.
- **Fully connected head:** at the end, flatten and use ordinary layers to classify.

Stacked, these learn a hierarchy: edges → textures → parts → objects.

12.3 A CNN in PyTorch

Replace the model in Chapter 11's script with this; everything else stays the same.

```
import torch.nn as nn

model = nn.Sequential(
    nn.Conv2d(1, 32, kernel_size=3, padding=1), nn.ReLU(), # 1 input channel -> 32
    feature maps, 28x28
    nn.MaxPool2d(2), # -> 14x14
    nn.Conv2d(32, 64, kernel_size=3, padding=1), nn.ReLU(),
    nn.MaxPool2d(2), # -> 7x7
    nn.Flatten(), # 64 * 7 * 7 = 3136
    nn.Linear(64 * 7 * 7, 128), nn.ReLU(), nn.Dropout(0.3),
    nn.Linear(128, 10),
).to(device)
```

Expect about 91–92% after 5 epochs — better than the plain network, with fewer weights in the early layers.

12.4 Famous architectures

Model	Year	Contribution
LeNet	1998	Early CNN for reading handwritten digits
AlexNet	2012	Won ImageNet by a huge margin on GPUs; started the deep learning boom
VGG	2014	Showed deeper stacks of small 3×3 filters work well
ResNet	2015	Residual connections made 100+ layer networks trainable
Vision Transformer (ViT)	2020	Applied transformers to images by cutting them into patches

12.5 Transfer learning: don't start from zero

Training a big image model needs millions of images. **Transfer learning** takes a model already trained on a huge dataset (ImageNet, 1.2 million images) and adapts it to your task with far less data. The early layers already know edges and textures; you only retrain the last part.

```
import torch, torch.nn as nn
from torchvision import models

model = models.resnet18(weights=models.ResNet18_Weights.DEFAULT)  # pre-trained on
ImageNet
for p in model.parameters():
    p.requires_grad = False  # freeze
everything
model.fc = nn.Linear(model.fc.in_features, 3)  # new head: e.g. 3
of your classes
optimizer = torch.optim.AdamW(model.fc.parameters(), lr=1e-3)  # train only the
head

Feed it your own images with torchvision.datasets.ImageFolder("my_images",
transform=models.ResNet18_Weights.DEFAULT.transforms()), where my_images has one sub-
folder per class. With a few hundred images per class you can often get strong accuracy.
```

Fine-tuning goes one step further: after training the head, unfreeze some or all layers and continue with a small learning rate. The same pattern — pre-train big once, adapt cheaply many times — is exactly how LLMs are used (Chapter 16).

12.6 Other vision tasks

- **Object detection:** find and box every object (YOLO, Faster R-CNN).
- **Segmentation:** label every pixel (U-Net, Mask R-CNN, Segment Anything).
- **Image generation:** create new images (Chapter 19).

12.7 A convolution by hand

Take a tiny 3×5 greyscale image: dark on the left (0), bright on the right (9).

	col 1	col 2	col 3	col 4	col 5
row 1	0	0	0	9	9
row 2	0	0	0	9	9
row 3	0	0	0	9	9

And a 3×3 **vertical-edge filter** whose three rows are all $(-1, 0, +1)$: it subtracts the left column from the right column.

Slide the filter across. At each position, multiply the 9 overlapping numbers pairwise and add them up:

- **Columns 1–3** (all 0): every product is 0, so the output is **0**. No edge.
- **Columns 2–4**: each row contributes $-1 \times 0 + 0 \times 0 + 1 \times 9 = 9$; three rows give **27**. Edge detected.
- **Columns 3–5**: each row contributes $-1 \times 0 + 0 \times 9 + 1 \times 9 = 9$; output **27**.

The feature map is [0, 27, 27]: high exactly where dark meets bright. A CNN isn't given this filter; it learns the nine numbers by backpropagation, along with dozens of others that detect other patterns.

```
import torch, torch.nn.functional as F

image = torch.tensor([[0, 0, 0, 9, 9]] * 3, dtype=torch.float32).view(1, 1, 3, 5)
edge = torch.tensor([[ -1, 0, 1]] * 3, dtype=torch.float32).view(1, 1, 3, 3)
print(F.conv2d(image, edge))  # tensor([[[[ 0., 27., 27.]]]])
```

(Strictly, deep learning's "convolution" is what mathematicians call cross-correlation — the filter isn't flipped. It makes no difference once the filter is learned.)

12.8 The output-size formula

For an input of width n , filter size k , padding p and stride s :

$$\text{output size} = \left\lfloor \frac{n + 2p - k}{s} \right\rfloor + 1$$

- 28×28 input, 3×3 filter, padding 1, stride 1: $(28 + 2 - 3) / 1 + 1 = \mathbf{28}$. Size kept.
- The same without padding: $(28 - 3) / 1 + 1 = \mathbf{26}$.
- 2×2 max pooling with stride 2 on 28: $(28 - 2) / 2 + 1 = \mathbf{14}$. Size halved.

This is how we knew the CNN in 12.3 ends with $64 \times 7 \times 7$ values before the linear layer.

12.9 Why CNNs need so few parameters

A conv layer's parameters are $(\text{input channels} \times k \times k + 1) \times \text{output channels}$:

Layer in the 12.3 model	Parameters
Conv2d($1 \rightarrow 32, 3 \times 3$)	$32 \times (1 \times 9) + 32 = 320$
Conv2d($32 \rightarrow 64, 3 \times 3$)	$64 \times (32 \times 9) + 64 = 18,496$
Linear($3136 \rightarrow 128$)	$3136 \times 128 + 128 = 401,536$
Linear($128 \rightarrow 10$)	1,290

The first layer produces 32 feature maps of 28×28 with just 320 parameters, because the same small filter is reused at every position (**weight sharing**). A fully connected layer producing the same 25,088 outputs from 784 pixels would need about 19.7 million. Notice most of this model's parameters sit in the classifier head, which is why modern CNNs replace it with **global average pooling** (average each feature map to one number).

12.10 Receptive field

The **receptive field** of a neuron is the patch of the original image that can influence it. After one 3×3 conv it's 3×3 ; after two stacked 3×3 convs, 5×5 ; each pooling layer roughly doubles it. Deep layers therefore "see" large parts of the image, which is how they recognize whole objects while early layers only see edges.

12.11 A residual block in code

ResNet's key idea: instead of each block computing a new output, it computes a *correction* that is added to its input.

```
import torch.nn as nn, torch.nn.functional as F

class ResidualBlock(nn.Module):
    def __init__(self, channels):
        super().__init__()
        self.conv1 = nn.Conv2d(channels, channels, 3, padding=1)
        self.bn1 = nn.BatchNorm2d(channels)
        self.conv2 = nn.Conv2d(channels, channels, 3, padding=1)
        self.bn2 = nn.BatchNorm2d(channels)

    def forward(self, x):
        out = F.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        return F.relu(x + out) # the shortcut: add the input back
```

Two reasons it works. If a block isn't useful, it can learn to output nearly zero and simply pass its input through, so adding layers can't make the network worse. And during backpropagation, gradients flow straight back along the shortcuts, which prevents them from vanishing over 100+ layers. Transformers use exactly the same trick (Chapter 15).

Try it

1. Add data augmentation: `transforms.Compose([transforms.RandomHorizontalFlip(), transforms.ToTensor()])` for the training set only.
2. Count parameters of both models: `sum(p.numel() for p in model.parameters())`.
3. Make a folder with photos of two kinds of objects from your phone and fine-tune ResNet-18 on them.

Chapter 13: Sequences — RNNs, LSTMs and their limits

Text, speech, stock prices and sensor readings are **sequences**: order matters. Before transformers, **recurrent neural networks (RNNs)** were the main tool for them. Understanding their limits explains why transformers won.

13.1 The recurrent idea

An RNN reads a sequence one step at a time and carries a **hidden state** — a memory vector — from each step to the next:

$$h_t = \tanh(W_x x_t + W_h h_{t-1} + b)$$

At word t , the new memory h_t mixes the current word x_t with the previous memory h_{t-1} . The same weights are used at every step, so it handles sequences of any length.

13.2 The problem: forgetting

In a long sentence, information from early words must pass through dozens of steps. During training, gradients shrink at every step on the way back (the **vanishing gradient** problem from Chapter 10). The network struggles to connect "The keys that I left on the kitchen table yesterday **are...**" back to "keys".

13.3 LSTMs and GRUs

The **Long Short-Term Memory (LSTM)** network (1997) adds a separate **cell state** and three **gates** — small learned switches:

- **Forget gate:** what to erase from memory.
- **Input gate:** what new information to store.
- **Output gate:** what to reveal at this step.

The **GRU** is a simpler version with two gates. Both remember far longer than plain RNNs, and they powered machine translation, speech recognition and phone keyboards through the 2010s.

13.4 An LSTM in PyTorch

A sentiment classifier skeleton: token ids go in, positive/negative comes out.

```
import torch, torch.nn as nn

class SentimentLSTM(nn.Module):
    def __init__(self, vocab_size, embed_dim=64, hidden_dim=128):
        super().__init__()
        self.embed = nn.Embedding(vocab_size, embed_dim)  # word id -> vector
        self.lstm = nn.LSTM(embed_dim, hidden_dim, batch_first=True)
        self.head = nn.Linear(hidden_dim, 2)

    def forward(self, token_ids):  # shape: (batch, sequence_length)
        x = self.embed(token_ids)
        _, (h_last, _) = self.lstm(x)  # final hidden state summarizes
        # the sentence
        return self.head(h_last[-1])

model = SentimentLSTM(vocab_size=10_000)
fake_batch = torch.randint(0, 10_000, (8, 50))  # 8 sentences of 50 tokens
print(model(fake_batch).shape)  # torch.Size([8, 2])
Train it with the same loop as Chapter 11.
```

13.5 Sequence-to-sequence and the birth of attention

For translation, an **encoder** RNN read the whole source sentence into one vector and a **decoder** RNN wrote the translation from it. Squeezing a whole sentence into one vector was a bottleneck.

In 2014, researchers added **attention**: at each output word, the decoder could look back at *all* encoder states and focus on the relevant ones. Translation quality jumped.

Then came the key question: if attention is doing the important work, do we need the slow, step-by-step recurrence at all? The 2017 answer — no — is the transformer (Chapter 15).

13.6 Why transformers replaced RNNs

	RNN / LSTM	Transformer
Reads a sequence	One step at a time	All positions at once
Training speed on GPUs	Slow: steps can't run in parallel	Fast: fully parallel
Long-range connections	Pass through many steps; fade	Direct: any word attends to any word
Cost as sequences grow	Grows linearly	Grows with the square of the length

RNNs are still useful for small, streaming or on-device problems, and newer **state-space models** (such as Mamba) revive the recurrent idea for very long sequences.

13.7 Time series

Forecasting (sales, energy demand) is a sequence task too. Start with simple baselines ("same as last week"), then classical methods (ARIMA, exponential smoothing) or gradient boosting on lag features, before trying LSTMs or transformers. Always split time series by date, never randomly.

13.8 Unrolling and backpropagation through time

To train an RNN, picture it **unrolled**: a copy of the network for every time step, chained together, all sharing the same weights. A 50-word sentence becomes a 50-layer-deep network. Backpropagation through this chain is called **backpropagation through time (BPTT)**.

That depth causes the trouble from Chapter 10. The gradient flowing back from word 50 to word 1 is multiplied by the recurrent weights at every step. If those factors are mostly below 1, the gradient vanishes and early words are forgotten; if above 1, it explodes. Two standard fixes: **gradient clipping** (cap the gradient's size) for explosions, and **truncated BPTT** (only backpropagate through, say, the last 100 steps) to keep long sequences affordable.

13.9 The LSTM equations

At each step the LSTM computes three gates (numbers between 0 and 1 from sigmoids) and a candidate memory, from the previous hidden state h and the current input x :

$$\begin{array}{ll}
 f_t &= \sigma(W_f[h_{t-1}, x_t] + b_f) && \text{forget gate} \\
 i_t &= \sigma(W_i[h_{t-1}, x_t] + b_i) && \text{input gate} \\
 \tilde{c}_t &= \tanh(W_c[h_{t-1}, x_t] + b_c) && \text{candidate memory} \\
 c_t &= f_t \odot c_{t-1} + i_t \odot \tilde{c}_t && \text{new cell state} \\
 o_t &= \sigma(W_o[h_{t-1}, x_t] + b_o) && \text{output gate} \\
 h_t &= o_t \odot \tanh(c_t) && \text{new hidden state}
 \end{array}$$

$\odot$ means element-by-element multiplication. The crucial line is the cell state update: the old memory is kept (scaled by the forget gate) and new information is *added*. Because it's addition rather than repeated multiplication, gradients can travel back through many steps without vanishing — the same principle as residual connections in ResNets and transformers.

13.10 Bidirectional RNNs

For tasks where the whole sequence is available up front (classifying a review, tagging each word as a name or place), it helps to read both ways. A **bidirectional** RNN runs one RNN left-to-right and another right-to-left and joins their states, so each word's representation reflects what came before *and* after it. In PyTorch: `nn.LSTM(embed_dim, hidden_dim, batch_first=True, bidirectional=True)`. It can't be used for generating text, where the future doesn't exist yet. Encoder transformers like BERT get the same benefit by attending in both directions.

13.11 Teacher forcing

When training a decoder to produce a sequence, what should it receive as the previous word: its own (possibly wrong) guess, or the correct one? **Teacher forcing** feeds the correct previous word during training. It trains much faster and more stably, because one early mistake doesn't derail the rest of the sentence.

The cost is **exposure bias**: at generation time the model only ever sees its own outputs, a situation it never practised. Transformers train with teacher forcing too — that's exactly what the shifted targets in Chapter 15's `get_batch` are: at every position the model sees the real text so far and predicts the next real token.

13.12 Greedy decoding vs beam search

Greedy decoding picks the single most likely word at each step. It can miss better sentences overall. Suppose the first word is "The" (0.5) or "A" (0.4). After "The", the best next word has probability 0.4, so the pair scores $0.5 \times 0.4 = 0.20$. After "A", the best next word has probability 0.9: $0.4 \times 0.9 = 0.36$. Greedy picks "The" and never finds the better sentence.

Beam search keeps the B most probable partial sentences at every step (the **beam width**, often 4–10), extends each, and keeps the best B again. With B = 2 it holds both "The" and "A" and finds the 0.36 path. Beam search suits tasks with one right answer, like translation and speech recognition. Chat models usually use sampling instead (Chapter 15.6), because beam search tends to produce bland, repetitive text.

Try it

1. Replace `nn.LSTM` with `nn.GRU`. What else in `forward` must change? (Hint: a GRU returns only `h`, not `(h, c)`.)
2. Count the parameters of the model. Which layer has the most?
3. Write down in your own words why a transformer can train faster than an LSTM.

Part 4 — Generative AI and transformers

Chapter 14: Tokens and embeddings

Neural networks only understand numbers. Before an LLM can read text, it chops it into **tokens** and turns each token into a list of numbers, an **embedding**, whose position in space captures meaning.

14.1 Tokenization

A **tokenizer** splits text into pieces from a fixed **vocabulary** (typically 30,000–200,000 pieces) and maps each piece to an integer id.

- Split by **characters**: tiny vocabulary, but very long sequences.
- Split by **words**: short sequences, but a huge vocabulary and no way to handle new words.
- Split into **subwords** (the modern choice): common words stay whole, rare words break into pieces. "unbelievably" → "un", "believ", "ably".

The most common algorithm is **byte-pair encoding (BPE)**: start from single characters, then repeatedly merge the most frequent adjacent pair into a new token, until the vocabulary reaches the target size.

```
# pip install tiktoken
import tiktoken

enc = tiktoken.get_encoding("cl100k_base")      # a GPT-4-era tokenizer
ids = enc.encode("Transformers are unbelievably useful!")
print(ids)                                     # a list of integers
print([enc.decode([i]) for i in ids])          # the pieces each integer stands for
print(len(ids), "tokens")
```

Why it matters in practice: models have a **context window** measured in tokens, API pricing is per token, and a rough rule for English is 1 token ≈ ¾ of a word. Languages other than English, and code, often need more tokens per word.

14.2 From ids to vectors: embeddings

An integer id carries no meaning: token 1045 isn't "closer" to 1046. So the model has an **embedding table** — a learned matrix with one row per token. Looking up a token gives its vector, for example 768 or 4,096 numbers.

```
import torch.nn as nn
embedding = nn.Embedding(num_embeddings=50_000, embedding_dim=768)  # learned during training
```

During training, tokens used in similar contexts drift to similar vectors. Distance in this space becomes a measure of meaning.

14.3 The famous example: word2vec

In 2013, **word2vec** learned word vectors by predicting a word from its neighbours. The vectors captured relationships as directions:

$$\overrightarrow{\text{king}} - \overrightarrow{\text{man}} + \overrightarrow{\text{woman}} \approx \overrightarrow{\text{queen}}$$

The "male → female" direction is roughly the same everywhere in the space, as are "country → capital" and "present → past tense".

14.4 Contextual embeddings

word2vec gives "bank" one vector whether it's a river bank or a savings bank. Transformers fix this: the token's embedding is only the starting point, and each attention layer updates it using the surrounding words. By the last layer, "bank" near "river" and "bank" near "loan" have very different vectors.

14.5 Sentence embeddings and semantic search

An **embedding model** turns a whole sentence or document into one vector. Texts with similar meaning land close together, even with no words in common. We measure closeness with **cosine similarity** (1 = same direction, 0 = unrelated).

```
# pip install sentence-transformers
from sentence_transformers import SentenceTransformer, util

model = SentenceTransformer("all-MiniLM-L6-v2")
sentences = [
    "How do I reset my password?",
    "I forgot my login credentials.",
    "What's the weather in Ottawa today?",
]
vectors = model.encode(sentences)
print(vectors.shape)                # (3, 384): 384 numbers per sentence
print(util.cos_sim(vectors, vectors)) # sentences 1 and 2 score far higher
                                     # than 1 and 3
```

This is the core of **semantic search**, recommendation, clustering documents, and **RAG** (Chapter 18). Collections of embeddings are stored in **vector databases** (FAISS, Chroma, pgvector, Pinecone) that can find the nearest vectors among millions in milliseconds.

14.6 Position

Attention on its own sees a bag of tokens with no order: "dog bites man" and "man bites dog" look the same. So transformers add **positional information** to each embedding. The original transformer added fixed sine-wave patterns; most modern LLMs use **rotary position embeddings (RoPE)**, which rotate vectors by an angle that depends on position.

14.7 BPE merges, by hand

A tiny corpus with word counts: low (5 times), lower (2), newest (6), widest (3). Start with single characters.

1. **Count adjacent pairs** across the corpus, weighted by word frequency. "e s" appears in newest and widest: $6 + 3 = 9$ times. "s t" also 9. "l o" and "o w" 7 each.
2. **Merge the most frequent pair** into a new token: "e" + "s" → "es". Words become n e w e s t, w i d e s t.
3. **Recount and repeat.** "e s t" now appears 9 times → merge into "est". Then "l o" → "lo", then "lo w" → "low".

After four merges the vocabulary contains "est" and "low". The word "lowest" never appeared, but it can be written as "low" + "est" — which is how subword tokenizers handle new words. Real tokenizers run tens of thousands of merges over billions of words, and modern ones (like tiktoken's) work on bytes, so any text in any language can be encoded.

14.8 How word2vec learns

word2vec rests on an old linguistic idea: a word's meaning is shown by the words that tend to appear around it. "Coffee" and "tea" both appear near "cup", "drink" and "hot", so they should get similar vectors.

The **skip-gram** version trains on a simple task: given a word, predict which words appear within a few positions of it. With **negative sampling**, it becomes a yes/no problem: is this (word, context) pair real, or a random fake? The score for a pair is the dot product of the two words' vectors passed through a sigmoid — logistic regression on embeddings. Training pushes real pairs' vectors together and fake pairs apart. After millions of updates, words used in similar ways end up close in the space. Nobody ever tells it what a word means.

14.9 Cosine similarity, worked

$$\cos(\mathbf{a}, \mathbf{b}) = \frac{\mathbf{a} \cdot \mathbf{b}}{\|\mathbf{a}\| \|\mathbf{b}\|}$$

- $\mathbf{a} = (1, 2, 2)$, $\mathbf{b} = (2, 4, 4)$: $\mathbf{b}$ is exactly twice $\mathbf{a}$, same direction. Cosine = **1**.
- $\mathbf{a} = (1, 0)$, $\mathbf{b} = (0, 1)$: at right angles, nothing in common. Cosine = **0**.
- $\mathbf{a} = (1, 2, 3)$, $\mathbf{b} = (4, 5, 6)$: dot product $4 + 10 + 18 = 32$; lengths $\sqrt{14} \approx 3.74$ and $\sqrt{77} \approx 8.77$; cosine = $32 / (3.74 \times 8.77) \approx \mathbf{0.97}$. Very similar.

Cosine ignores length and compares direction only. That's useful for text: a long and a short document about the same topic point the same way. If vectors are normalized to length 1 first, cosine similarity is just the dot product, which is why the RAG code in Chapter 18 uses `normalize_embeddings=True`.

14.10 How big are embeddings?

Model	Vector size
word2vec (classic)	300
all-MiniLM-L6-v2 (Chapter 14.5)	384
BERT-base	768
Typical modern embedding models for search	768–3,072
Internal token vectors in large LLMs	4,096 to over 16,000

Larger vectors can capture finer distinctions but cost more memory and slower search: a million documents with 1,024 numbers each, stored as 4-byte floats, take about 4 GB. Pick the smallest model that does well on your own retrieval test (Chapter 17.7).

14.11 Tokenizer quirks you'll notice

- **Spaces are part of tokens.** " cat" (with a leading space) and "cat" are often different tokens.
- **Numbers split unpredictably.** "12345" may become "123" + "45", one reason LLMs are shaky at arithmetic without a calculator tool.

- **Letters are hidden.** The model sees "straw" + "berry", not ten letters, which is why counting letters in a word can trip it up.
- **Other languages cost more.** Text in many non-English languages takes more tokens for the same meaning, so it's more expensive and fills the context window faster.
- **Special tokens** mark structure: end of text, start of a turn, system vs user vs assistant (Chapter 16.3).
- **Counts differ between models.** Each model family has its own tokenizer, so "1,000 tokens" isn't the same amount of text everywhere.

Try it


1. Tokenize a sentence in English and the same sentence in Persian or French. Which uses more tokens?
2. Add five sentences of your own to the embedding example, including two paraphrases. Do the paraphrases score highest?
3. Build a tiny search engine: embed 10 FAQ answers, embed a question, and print the answer with the highest cosine similarity.

Chapter 15: Attention and the transformer

A transformer is a stack of identical blocks, each doing two things: **attention** (tokens share information with each other) and a **feed-forward network** (each token thinks on its own). This chapter explains both, then you build a working mini-GPT.

15.1 The whole architecture at a glance

A GPT model predicts one token, then feeds it back in



Decoder-only transformer (GPT) · one prediction step

Text becomes tokens, tokens become vectors, the vectors pass through N transformer blocks, and the final layer scores every token in the vocabulary. The model samples one, appends it, and runs again.

15.2 Attention: queries, keys and values

Think of a library. You arrive with a **query** ("books about cats"). Every book has a **key** on its spine (its topic). You compare your query with every key, and take the **values** (contents) of the books that match best.

In self-attention, every token produces all three, by multiplying its vector by three learned weight matrices:

- **Query (Q)**: what am I looking for?
- **Key (K)**: what do I contain?
- **Value (V)**: what information will I pass on if chosen?

Each token's query is compared (dot product) with every token's key. High scores mean "relevant". Softmax turns the scores into weights summing to 1, and the token's new vector is the weighted mix of all the values:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Dividing by $\sqrt{d_k}$ (the key length) keeps the scores from getting too large, which would make softmax too spiky to train. In "The animal didn't cross the street because it was too tired", the query from "it" matches the key from "animal" strongly, so "it" absorbs information about the animal.

15.3 Masking, heads and the rest of the block

- **Causal mask:** a GPT must predict the next token without seeing it, so each token may only attend to itself and earlier tokens. Future positions get a score of $-\infty$ before softmax, which becomes a weight of 0.
- **Multi-head attention:** run several attentions in parallel (e.g. 12 heads), each with its own Q, K, V. One head might track grammar, another who "it" refers to. Their outputs are joined and mixed.
- **Feed-forward network:** two linear layers with a GELU in between, applied to each token separately. Much of the model's factual knowledge seems to be stored here.
- **Residual connections:** each part's output is *added* to its input ($x + \text{attention}(x)$), so information and gradients can flow straight through 100+ layers.
- **Layer normalization:** rescales each vector to keep numbers stable.

15.4 Three families of transformers

Family	Attention	Trained to	Examples	Best at
Encoder-only	Sees the whole text both ways	Fill in masked words	BERT, RoBERTa	Classification, search embeddings
Decoder-only	Causal: only earlier tokens	Predict the next token	GPT, Claude, Llama, Mistral	Generating text, chat, code
Encoder-decoder	Encoder reads, decoder writes	Map input text to output text	Original Transformer, T5	Translation, summarization

Nearly all modern chat LLMs are decoder-only.

15.5 Build a tiny GPT

This is a complete, character-level GPT trained on Shakespeare's plays, about 1.8 million parameters. It's a simplified version of Andrej Karpathy's well-known *nanoGPT*. Use a GPU runtime in Colab; training takes a few minutes.

```
import torch, torch.nn as nn, torch.nn.functional as F
import urllib.request
```

```

# ---- data: characters are our tokens ----
url = "https://raw.githubusercontent.com/karpathy/char-
rnn/master/data/tinyshakespeare/input.txt"
text = urllib.request.urlopen(url).read().decode("utf-8")
chars = sorted(set(text)); vocab_size = len(chars)
stoi = {c: i for i, c in enumerate(chars)}; itos = {i: c for c, i in stoi.items()}
encode = lambda s: [stoi[c] for c in s]
decode = lambda ids: "".join(itos[i] for i in ids)
data = torch.tensor(encode(text), dtype=torch.long)
n = int(0.9 * len(data)); train_data, val_data = data[:n], data[n:]

device = "cuda" if torch.cuda.is_available() else "cpu"
block_size, batch_size = 128, 64 # context length, sequences per batch
n_embd, n_head, n_layer, dropout = 192, 6, 4, 0.1

def get_batch(split):
    d = train_data if split == "train" else val_data
    ix = torch.randint(len(d) - block_size - 1, (batch_size,))
    x = torch.stack([d[i:i + block_size] for i in ix])
    y = torch.stack([d[i + 1:i + block_size + 1] for i in ix]) # targets: shifted
    by one
    return x.to(device), y.to(device)

# ---- the model ----
class CausalSelfAttention(nn.Module):
    def __init__(self):
        super().__init__()
        self.qkv = nn.Linear(n_embd, 3 * n_embd) # makes Q, K, V in one go
        self.proj = nn.Linear(n_embd, n_embd)
        self.drop = nn.Dropout(dropout)
        self.register_buffer("mask", torch.tril(torch.ones(block_size, block_size)))

    def forward(self, x):
        B, T, C = x.shape
        q, k, v = self.qkv(x).split(n_embd, dim=2)
        # split into heads: (B, heads, T, head_size)
        q = q.view(B, T, n_head, C // n_head).transpose(1, 2)
        k = k.view(B, T, n_head, C // n_head).transpose(1, 2)
        v = v.view(B, T, n_head, C // n_head).transpose(1, 2)
        att = (q @ k.transpose(-2, -1)) / (k.size(-1) ** 0.5) # scores
        att = att.masked_fill(self.mask[:T, :T] == 0, float("-inf")) # no peeking
        ahead
        att = self.drop(F.softmax(att, dim=-1)) # weights
        out = (att @ v).transpose(1, 2).contiguous().view(B, T, C) # mix values,
        join heads
        return self.drop(self.proj(out))

class Block(nn.Module):
    def __init__(self):
        super().__init__()
        self.ln1, self.attn = nn.LayerNorm(n_embd), CausalSelfAttention()
        self.ln2 = nn.LayerNorm(n_embd)
        self.ffn = nn.Sequential(nn.Linear(n_embd, 4 * n_embd), nn.GELU(),
                                  nn.Linear(4 * n_embd, n_embd), nn.Dropout(dropout))

```

```

    def forward(self, x):
        x = x + self.attn(self.ln1(x))    # residual connection
        x = x + self.ffn(self.ln2(x))    # residual connection
        return x

class TinyGPT(nn.Module):
    def __init__(self):
        super().__init__()
        self.tok = nn.Embedding(vocab_size, n_embd)
        self.pos = nn.Embedding(block_size, n_embd)
        self.blocks = nn.Sequential(*[Block() for _ in range(n_layer)])
        self.ln_f = nn.LayerNorm(n_embd)
        self.head = nn.Linear(n_embd, vocab_size)

    def forward(self, idx, targets=None):
        B, T = idx.shape
        x = self.tok(idx) + self.pos(torch.arange(T, device=idx.device))
        logits = self.head(self.ln_f(self.blocks(x)))
        loss = None if targets is None else F.cross_entropy(logits.view(-1,
vocab_size), targets.view(-1))
        return logits, loss

    @torch.no_grad()
    def generate(self, idx, max_new_tokens, temperature=1.0):
        for _ in range(max_new_tokens):
            logits, _ = self(idx[:, -block_size:])
            probs = F.softmax(logits[:, -1, :] / temperature, dim=-1)
            idx = torch.cat([idx, torch.multinomial(probs, 1)], dim=1)    # sample,
append, repeat
        return idx

# ---- training: the same loop as Chapter 11 ----
model = TinyGPT().to(device)
print(sum(p.numel() for p in model.parameters()) / 1e6, "million parameters")
opt = torch.optim.AdamW(model.parameters(), lr=1e-3)
for step in range(3000):
    x, y = get_batch("train")
    _, loss = model(x, y)
    opt.zero_grad(); loss.backward(); opt.step()
    if step % 500 == 0:
        model.eval()
        with torch.no_grad():
            _, val_loss = model(*get_batch("val"))
        model.train()
        print(f"step {step}: train loss {loss.item():.2f}, val loss
{val_loss.item():.2f}")

start = torch.zeros((1, 1), dtype=torch.long, device=device)    # start from a newline
print(decode(model.generate(start, 500)[0].tolist()))

```

The loss starts around 4.2 (random guessing among ~65 characters) and falls to roughly 1.5. The output won't make sense, but it will look like a play: character names, line breaks, old-English-sounding words. You built a language model from nothing. Real LLMs use the same architecture with subword tokens, thousands of dimensions, dozens of layers and trillions of training tokens.

15.6 Sampling: how the next token is chosen

- **Greedy**: always take the most likely token. Repetitive.
- **Temperature**: divide scores before softmax. Below 1 is safer and more focused; above 1 is more creative and more random.
- **Top-k**: sample only from the k most likely tokens.
- **Top-p (nucleus)**: sample from the smallest set of tokens whose probabilities add up to p (e.g. 0.9).

15.7 Scaling and efficiency

Researchers found **scaling laws**: loss falls predictably as you increase parameters, data and compute together. That's why models grew from millions to hundreds of billions of parameters. Tricks that make big models practical include **FlashAttention** (faster attention on GPUs), the **KV cache** (reuse past keys and values when generating), **mixture of experts** (only part of the network runs for each token), and **quantization** (storing weights in 8 or 4 bits).

15.8 Attention by hand, with tiny numbers

Three tokens, "the cat sat", with 2-number vectors (real models use thousands). Compute the new vector for "sat".

Token	Key (K)	Value (V)
the	(0, 1)	(1, 0)
cat	(1, 1)	(0, 2)
sat	(1, 0)	(2, 1)

The query for "sat" is (1, 0).

1. **Scores** = query · each key: the → $1 \times 0 + 0 \times 1 = 0$; cat → $1 \times 1 + 0 \times 1 = 1$; sat → $1 \times 1 + 0 \times 0 = 1$.
2. **Scale** by $\sqrt{d_k} = \sqrt{2} \approx 1.41$: 0, 0.71, 0.71.
3. **Softmax**: $e^0 = 1$, $e^{0.71} \approx 2.03$, $e^{0.71} \approx 2.03$; divided by their total 5.06, the weights are **0.20, 0.40, 0.40**.
4. **Mix the values**: $0.20 \times (1, 0) + 0.40 \times (0, 2) + 0.40 \times (2, 1) = \mathbf{(1.00, 1.20)}$.

"sat" now carries 40% of its information from "cat" (who sat?) and 20% from "the". With the causal mask, "cat" would only be allowed to look at "the" and itself. Check it in code:

```
import numpy as np

q = np.array([1.0, 0.0])
K = np.array([[0, 1], [1, 1], [1, 0]], dtype=float)
V = np.array([[1, 0], [0, 2], [2, 1]], dtype=float)
scores = K @ q / np.sqrt(2)
weights = np.exp(scores) / np.exp(scores).sum()
print(weights.round(2), (weights @ V).round(2))  # [0.2 0.4 0.4] [1. 1.2]
```

Where do Q, K and V come from? Each is the token's vector multiplied by its own learned matrix (W_Q , W_K , W_V). Training adjusts those matrices so that useful queries match useful keys.

15.9 Why divide by $\sqrt{d}$?

A dot product adds up d products. For random vectors, the result's spread grows with $\sqrt{d}$: with $d = 4,096$, raw scores would typically be around ± 64 . Softmax on numbers that large becomes almost one-hot (one weight near 1, the rest near 0), and its gradients become tiny, so learning stalls. Dividing by $\sqrt{d}$ brings the scores back to a spread of about 1 whatever the size of the model.

15.10 Counting a transformer's parameters

With model width d (the vector size) and L layers, each block has:

- **Attention:** four $d \times d$ matrices (Q, K, V and the output projection) = $4d^2$.
- **Feed-forward:** $d \rightarrow 4d \rightarrow d = 8d^2$.

So about $12d^2$ per layer, plus the embedding table (vocabulary size $\times d$):

$$\text{parameters} \approx 12 L d^2 + Vd$$

Check against GPT-2 small ($d = 768$, $L = 12$, vocabulary 50,257): $12 \times 12 \times 768^2 \approx 85$ million, plus $50,257 \times 768 \approx 39$ million, total $\approx$ **124 million** — its published size. Our tiny GPT: $12 \times 4 \times 192^2 \approx 1.8$ million. Notice the d^2 : doubling the width quadruples the parameters per layer.

15.11 The KV cache

When generating, each new token needs the keys and values of every earlier token. Recomputing them at every step would repeat almost all the work. The **KV cache** stores each layer's keys and values as they're produced, so each new step only computes the newest token. Generation becomes much faster.

The price is memory: 2 (keys and values) $\times$ layers $\times d \times$ tokens $\times$ bytes per number. For a model with 32 layers, $d = 4,096$, an 8,000-token conversation and 2-byte numbers, that's about **4.3 GB** for a single conversation. Modern models shrink this with **grouped-query attention**, where several query heads share one set of keys and values.

15.12 Why long context is expensive

Every token attends to every earlier token, so the attention score table has $T \times T$ entries for T tokens. Doubling the context length quadruples that work; at 128,000 tokens it's about 16 billion scores per head per layer. Engineering helps — **FlashAttention** computes attention in small tiles without storing the whole table, and some models use **sliding-window** or **sparse** attention — but long prompts still cost more and run slower. Models can also use information in the middle of very long inputs less reliably than information near the start or end, so put the most important instructions and documents where they're easy to find, and test it.

Try it

1. Generate with `temperature=0.5` and `temperature=1.5`. Describe the difference.
2. Double `n_layer` to 8. Does validation loss improve? How much slower is training?
3. Remove the causal mask line and train again. Why does the training loss drop suspiciously fast, and why is the model useless for generation?

Chapter 16: Training LLMs

An assistant like ChatGPT or Claude is made in two big phases. **Pre-training** teaches a model language and knowledge by predicting the next token on trillions of tokens. **Post-training** turns that raw text predictor into a helpful, honest, safe assistant.

16.1 The pipeline

Stage	Data	What the model learns	Rough cost
1. Pre-training	Trillions of tokens of web pages, books, code, papers	Grammar, facts, reasoning patterns; becomes a base model	Weeks to months on thousands of GPUs
2. Supervised fine-tuning (SFT)	Thousands to millions of example conversations written or checked by people	To follow instructions and answer as an assistant	Hours to days
3. Preference tuning (RLHF, DPO, RLAIF)	Pairs of answers ranked better/worse	Which answers people prefer: helpful, harmless, honest	Days
4. RL on verifiable tasks	Maths, code and puzzles with checkable answers	Longer step-by-step reasoning ("reasoning models")	Days to weeks
5. Evaluation and red-teaming	Benchmarks, human raters, adversarial testers	Nothing — it checks the model (Chapter 17)	Ongoing

16.2 Pre-training

Pre-training is exactly what you did with the tiny GPT in Chapter 15, at vast scale. The loss is cross-entropy on the next token: how surprised the model was by the real continuation.

$$L = -\frac{1}{T} \sum_{t=1}^T \log P_{\theta}(x_t \mid x_1, \dots, x_{t-1})$$

It's called **self-supervised**: the text labels itself, so no human annotation is needed. What matters most:

- **Data quality**: deduplication, removing spam and toxic content, filtering by quality, and a careful mix of domains (web, code, maths, languages). Data work is often the biggest differentiator between labs.
- **Scale**: a common rule of thumb from the *Chinchilla* paper is about 20 training tokens per parameter for compute-efficient training; modern models are often trained on far more. Training compute is roughly $6 \times \text{parameters} \times \text{tokens floating-point operations}$.

- **Distributed training:** no single GPU can hold a large model. **Data parallelism** gives each GPU a different batch; **tensor** and **pipeline parallelism** split the model itself across GPUs; **FSDP/ZeRO** shards weights and optimizer state to save memory.
- **Stability:** warmup, learning-rate decay, gradient clipping, mixed precision and frequent checkpoints (Chapter 11), because a crash three weeks in is very expensive.

The result is a **base model**: it continues text brilliantly, but if you ask it a question it may just write more questions.

16.3 Supervised fine-tuning (SFT)

SFT continues training on example conversations: a prompt and an ideal response. The loss is the same next-token cross-entropy, usually computed only on the response tokens. Conversations are formatted with a **chat template**, special tokens marking who is speaking (system, user, assistant). Quality beats quantity: a few thousand excellent examples can transform a base model's behaviour.

16.4 Preference tuning: RLHF, DPO and RLAIF

Some qualities are easier to judge than to write: which of two answers is clearer, kinder, or more accurate?

- **RLHF** (reinforcement learning from human feedback): people rank pairs of answers; a **reward model** is trained to predict their preferences; then the LLM is optimized with reinforcement learning (classically **PPO**) to score highly, with a penalty for drifting too far from the SFT model.
- **DPO** (direct preference optimization): skips the separate reward model and RL loop, and adjusts the LLM directly from preferred/rejected pairs. Simpler and widely used.
- **RLAIF / Constitutional AI**: an AI model, guided by a written set of principles, produces some of the feedback instead of humans. Anthropic uses this approach for Claude.

16.5 Reasoning models

Newer models are further trained with reinforcement learning on problems whose answers can be checked automatically (a maths result, code that passes tests). Rewarding correct final answers teaches the model to "think" longer, write out steps, and check its own work before answering.

16.6 Fine-tuning an LLM yourself

You won't pre-train an LLM, but you can fine-tune an open model on your own data. First ask whether you need to:

Approach	Use when	Cost
Prompting (Chapter 18)	Always try first	Free
RAG (Chapter 18)	The model needs your facts or documents	Low
Fine-tuning	You need a consistent style, format or narrow skill, or a smaller, cheaper model that	Medium

Approach	Use when	Cost
	matches a big one	
Pre-training from scratch	A new language or domain with no suitable model; big budget	Very high

Full fine-tuning updates every weight — memory-hungry. **Parameter-efficient fine-tuning (PEFT)** freezes the model and trains a tiny add-on. The most popular is **LoRA** (low-rank adaptation): it adds small matrices beside the existing weights and trains only those, often under 1% of the parameters. **QLoRA** also loads the frozen model in 4-bit, so a 7-billion-parameter model can be fine-tuned on one consumer GPU.

16.7 Code: LoRA fine-tuning with Hugging Face

This fine-tunes a small open model (0.5 billion parameters, fits on Colab's free GPU) on 2,000 example conversations. Hugging Face libraries change quickly; if an argument errors, check the current [TRL documentation](#).

```
# pip install -U transformers datasets peft trl accelerate
from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForCausalLM, pipeline
from peft import LoraConfig
from trl import SFTTrainer, SFTConfig

model_name = "Qwen/Qwen2.5-0.5B-Instruct"
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(model_name)

# Conversations stored as {"messages": [{"role": "user", ...}, {"role": "assistant", ...}]}
dataset = load_dataset("trl-lib/Capybara", split="train[:2000]")

lora = LoraConfig(r=16, lora_alpha=32, lora_dropout=0.05,
                  target_modules="all-linear", task_type="CAUSAL_LM")
args = SFTConfig(output_dir="qwen-lora", num_train_epochs=1,
                 per_device_train_batch_size=4, gradient_accumulation_steps=4,
                 learning_rate=2e-4, logging_steps=20)

trainer = SFTTrainer(model=model, args=args, train_dataset=dataset,
                    peft_config=lora, processing_class=tokenizer)
trainer.train()          # watch the training loss fall in the logs
trainer.save_model()     # saves only the small LoRA adapter

chat = pipeline("text-generation", model=trainer.model, tokenizer=tokenizer)
print(chat([{"role": "user", "content": "Explain overfitting in two sentences."}],
          max_new_tokens=100)[0]["generated_text"])
```

To fine-tune on your own data, replace the dataset with a list of your conversations in the same messages format: `Dataset.from_list(my_conversations)`. For preference tuning, TRL's `DPOTrainer` takes prompt, chosen and rejected columns in the same style.

16.8 Practical tips

- **Data first:** 500 clean, consistent examples beat 50,000 messy ones. Read your data.
- **Hold out an evaluation set** and compare against the untuned model (Chapter 17). Fine-tuning can make a model worse at things you didn't train on (**catastrophic forgetting**).
- **Typical LoRA settings:** rank 8–64, learning rate 1e-4 to 3e-4, 1–3 epochs. More epochs on small data overfits fast.
- **Track experiments** with a tool like Weights & Biases or MLflow (`report_to="wandb"`).

16.9 The reward model's loss

A reward model reads a prompt x and a response y and outputs one number, $r(x, y)$. It's trained on pairs where people preferred response y_w (winner) over y_l (loser). The **Bradley–Terry** model says the probability that the winner is preferred is the sigmoid of the difference in rewards, so the loss is:

$$L_{RM} = -\log \sigma(r(x, y_w) - r(x, y_l))$$

The model learns to give preferred answers higher scores. RLHF then trains the LLM (the **policy** π) to maximize reward while staying close to the SFT model π_{ref} :

$$\max_{\pi} \mathbb{E}[r(x, y)] - \beta \text{KL}(\pi \parallel \pi_{ref})$$

The KL penalty matters: without it, the model finds strange outputs that fool the reward model into high scores (**reward hacking**, Chapter 20.10), and its language degrades.

16.10 DPO's loss

DPO showed that the same goal can be reached without a reward model or RL, directly from preference pairs:

$$L_{DPO} = -\log \sigma \left(\beta \left[\log \frac{\pi(y_w | x)}{\pi_{ref}(y_w | x)} - \log \frac{\pi(y_l | x)}{\pi_{ref}(y_l | x)} \right] \right)$$

Read it as: increase how much more likely the model makes the preferred answer *compared with the reference model*, and decrease the same for the rejected answer. β controls how far it may drift from the reference. In code it's just another training loss over log-probabilities, which is why DPO became so popular.

16.11 LoRA's maths

A layer's weight matrix W has $d \times k$ numbers. LoRA freezes W and learns a low-rank update made of two thin matrices, B ($d \times r$) and A ($r \times k$), with a small rank r :

$$W' = W + \frac{\alpha}{r} BA$$

For $d = k = 4,096$ and $r = 16$, full fine-tuning updates 16.8 million numbers in that layer; LoRA trains $16 \times (4,096 + 4,096) \approx 131,000$ — under 1%. B starts at zero, so training begins exactly at the original model. The bet behind LoRA is that the *change* needed for a new task is simple (low-rank) even when the model itself is not. After training, BA can be added into W permanently, so the fine-tuned model runs at the same speed as the original. You can also keep several small adapters and swap them on one base model.

16.12 Quantization

Quantization stores weights with fewer bits:

Precision	Bytes per parameter	Size of a 7-billion-parameter model
float32	4	~28 GB
bfloat16 / float16	2	~14 GB
int8	1	~7 GB
4-bit	0.5	~3.5 GB

8-bit usually loses very little quality; 4-bit loses a bit more but lets large models run on laptops and single GPUs. Common formats: **GGUF** (for llama.cpp and Ollama on ordinary computers), **GPTQ** and **AWQ** (for GPU serving), and **NF4** from bitsandbytes (used by QLoRA).

16.13 Chat templates in practice

Each model family expects conversations formatted with its own special tokens. The tokenizer knows the right template:

```
from transformers import AutoTokenizer

tok = AutoTokenizer.from_pretrained("Qwen/Qwen2.5-0.5B-Instruct")
messages = [{"role": "system", "content": "You are concise."},
             {"role": "user", "content": "What is LoRA?"}]
print(tok.apply_chat_template(messages, tokenize=False, add_generation_prompt=True))
# <|im_start|>system\nYou are concise.<|im_end|>\n<|im_start|>user\nWhat is
LoRA?<|im_end|>\n<|im_start|>assistant\n
```

Feeding a model text in the wrong format, or fine-tuning with one template and running with another, quietly makes it worse. Always use `apply_chat_template`; TRL's `SFTTrainer` does this for you when data is in the `messages` format.

16.14 How much GPU memory does training need?

Task	Rough memory per parameter	7B model
Inference in bfloat16	2 bytes	~14 GB + KV cache
Full fine-tuning with AdamW, mixed precision	~16 bytes (weights, gradients, a float32 master copy, two Adam averages)	~110 GB + activations: several GPUs
LoRA	~2 bytes (frozen weights) + small adapters	~16–20 GB
QLoRA (4-bit base)	~0.5 bytes + small adapters	~6–10 GB: one consumer GPU

Activations (values saved during the forward pass for the backward pass) add more, growing with batch size and sequence length. **Gradient checkpointing** trades extra computation for less activation memory by recomputing them during the backward pass.

Try it

1. Before training, ask the base model a question; ask again after. What changed?
2. Create 50 conversations in a particular style (for example, answers that always end with a one-line summary) and fine-tune on them.
3. Change the LoRA rank r from 16 to 4. Compare training loss and the adapter's file size.

Chapter 17: Evaluating LLMs

Evaluating an LLM is harder than evaluating a classifier, because there is rarely one correct answer. Good LLM evaluation combines several layers: automatic metrics, public benchmarks, your own task-specific test set, an LLM judge calibrated against humans, and monitoring in production.

17.1 The layers of LLM evaluation

Layer	Question	Tools	Weakness
Intrinsic	How well does it model language?	Perplexity, validation loss	Says little about usefulness
Public benchmarks	How does it compare with other models?	MMLU-style knowledge tests, GSM8K-style maths, HumanEval-style coding, leaderboards	Contamination, saturation, not your task
Task evals	Does it do <i>my</i> job well?	Your own test set + reference answers + metrics	Needs effort to build
LLM-as-judge	Is this open-ended answer good?	A strong model grading with a rubric	Judge bias; must be validated
Human evaluation	Do people actually find it good?	Expert raters, side-by-side comparisons	Slow, expensive, raters disagree
Safety evals	Can it be misused or cause harm?	Red-teaming, jailbreak suites, bias tests	Never complete
Online	Does it work for real users?	A/B tests, thumbs up/down, task completion, monitoring	Only after launch

17.2 Perplexity

Perplexity is e raised to the average next-token loss. Intuitively: how many tokens the model is "choosing between" at each step. Lower means the text is less surprising to the model.

```
import math, torch
from transformers import AutoTokenizer, AutoModelForCausalLM

tok = AutoTokenizer.from_pretrained("gpt2")
lm = AutoModelForCausalLM.from_pretrained("gpt2")

def perplexity(text):
    ids = tok(text, return_tensors="pt").input_ids
    with torch.no_grad():
        loss = lm(ids, labels=ids).loss      # average next-token cross-entropy
    return math.exp(loss.item())

print(perplexity("The cat sat on the mat. ")) # low
print(perplexity("Mat the on sat cat the. ")) # much higher
Use it to compare checkpoints of the same model during training, not models with different tokenizers.
```

17.3 Reference-based metrics

When you have a correct answer to compare with:

Metric	Measures	Good for
Exact match / accuracy	Output equals the answer (after normalizing)	Multiple choice, short factual answers, classification
Token F1	Word overlap with the answer	Extractive question answering
BLEU / ROUGE	n-gram overlap with a reference	Translation / summarization; weak for modern LLMs
BERTScore / embedding similarity	Similarity of meaning, via embeddings	Paraphrased answers
pass@k	Share of problems where at least one of k generated programs passes the unit tests	Code generation
JSON / schema validity	Output parses and matches the required format	Structured outputs, tool calls

17.4 Benchmarks and their traps

Public benchmarks cover knowledge, maths, coding, reasoning, long context and tool use. Use them to shortlist models, then test on your own task. Watch out for:

- **Contamination**: benchmark questions leaked into training data, inflating scores.
- **Saturation**: when top models all score above 90%, the benchmark stops separating them.

- **Prompt sensitivity:** small wording or formatting changes can move scores several points. Keep prompts fixed when comparing.

17.5 Build your own evaluation set

This is the most valuable evaluation you'll ever do.

1. Collect 50–500 realistic inputs, ideally from real users or logs.
2. Cover the common cases, known edge cases, and inputs that should be refused or escalated.
3. Write a reference answer or a clear grading rubric for each.
4. Split into a **development set** (to iterate prompts on) and a **held-out set** (for final comparisons), just like Chapter 4.
5. Version it with your code (Chapter 21) and run it on every prompt or model change, like unit tests. A drop is a **regression**.
6. Add every real failure you find in production as a new test case.

17.6 LLM-as-judge

For open-ended outputs, a strong LLM can grade answers against a rubric. Scores are cheap and fast, but only trustworthy once you've checked them against human ratings.

```
# pip install anthropic (set the ANTHROPIC_API_KEY environment variable)
import json, anthropic
import pandas as pd

client = anthropic.Anthropic()

JUDGE_PROMPT = """You are grading an AI assistant's answer.
Question: {question}
Reference answer: {reference}
Assistant's answer: {answer}

Score the answer from 1 to 5 for correctness against the reference
(5 = fully correct and complete, 1 = wrong or missing). Explain first, then score.
Respond ONLY with JSON: {"reasoning": "...", "score": <1-5>}}"""

def judge(question, reference, answer):
    msg = client.messages.create(
        model="claude-sonnet-5", max_tokens=400,
        messages=[{"role": "user", "content": JUDGE_PROMPT.format(
            question=question, reference=reference, answer=answer)}])
    return json.loads(msg.content[0].text)

eval_set = [
    {"question": "What does 'overfitting' mean?",
     "reference": "The model memorizes training data, including noise, and performs poorly on new data."},
    {"question": "Which metric matters most for cancer screening, precision or recall?",
     "reference": "Recall, because missing a real case is the costliest error."},
]

def run_eval(generate):
```

generate: any function question -> answer

```

rows = []
for case in eval_set:
    answer = generate(case["question"])
    rows.append(**case, "answer": answer, **judge(case["question"],
case["reference"], answer))
df = pd.DataFrame(rows)
print("mean score:", df["score"].mean())
return df

```

Known judge biases and how to reduce them:

- **Position bias** in pairwise comparisons (preferring the first answer): run each pair twice with the order swapped.
- **Verbosity bias** (preferring longer answers): tell the rubric that length is not quality.
- **Self-preference** (a model favouring its own outputs): use a judge from a different model family when possible.
- **Vague rubrics**: use specific, observable criteria and a small scale (pass/fail or 1–5), and ask for reasoning before the score.
- **Validate**: have humans grade 50–100 of the same answers and measure agreement (percentage agreement, or Cohen's kappa). Improve the judge prompt until agreement is close to human–human agreement.

17.7 Evaluating RAG systems

A RAG system (Chapter 18) can fail at retrieval or at generation, so measure both:

Part	Metric	Question
Retrieval	Recall@k	Is the right document among the top k results?
Retrieval	MRR (mean reciprocal rank)	How high does the first right document appear?
Generation	Faithfulness / groundedness	Is every claim supported by the retrieved text?
Generation	Answer relevance	Does it actually answer the question?
End to end	Correctness	Does it match the reference answer?

Libraries such as Ragas and DeepEval implement many of these with LLM judges.

17.8 Safety and robustness

- **Red-teaming**: people (and other models) actively try to make the system misbehave: harmful instructions, leaking private data, **prompt injection** hidden in documents or web pages.
- **Bias tests**: swap names, genders or dialects in otherwise identical prompts and compare the outputs.

- **Hallucination checks:** questions with known answers, and questions with no answer where "I don't know" is correct.
- **Robustness:** typos, rephrasings, other languages, very long inputs.

17.9 Online evaluation

Offline tests don't capture everything. After launch, run A/B tests on real traffic, log user feedback, track task-completion and escalation rates, sample conversations for human review, and watch latency and cost per request.

17.10 Perplexity, worked

A model reads four tokens and gave the actual next tokens probabilities 0.5, 0.25, 0.125 and 0.5.

1. Loss per token = $-\ln(\text{probability})$: 0.69, 1.39, 2.08, 0.69.
2. Average loss = $4.85 / 4 \approx 1.21$.
3. Perplexity = $e^{1.21} \approx \mathbf{3.4}$.

Interpretation: on average the model was as uncertain as if it were choosing between about 3.4 equally likely tokens. A perfect model scores 1; a model guessing uniformly over a 50,000-token vocabulary scores 50,000.

$$\text{perplexity} = \exp \left(-\frac{1}{T} \sum_{t=1}^T \log P(x_t | x_{<t}) \right)$$

17.11 Cohen's kappa, by hand

Two graders (a person and your LLM judge, say) mark the same 100 answers pass or fail:

	Judge: pass	Judge: fail
Human: pass	60	10
Human: fail	5	25

1. **Observed agreement** $p_o = (60 + 25) / 100 = 0.85$.
2. **Agreement expected by chance** p_e : the human passes 70% and the judge 65%, so they'd agree on "pass" $0.70 \times 0.65 = 0.455$ of the time and on "fail" $0.30 \times 0.35 = 0.105$ of the time by luck alone: $p_e = 0.56$.
3. **Kappa** = $(p_o - p_e) / (1 - p_e) = 0.29 / 0.44 \approx \mathbf{0.66}$.

An 85% raw agreement sounds high, but much of it is chance; kappa corrects for that. A common rough scale: below 0.2 slight, 0.2–0.4 fair, 0.4–0.6 moderate, 0.6–0.8 substantial, above 0.8 almost perfect. Compute it with `sklearn.metrics.cohen_kappa_score(human, judge)`. Compare the judge–human kappa with the kappa between two humans: if humans only reach 0.7 with each other, a judge at 0.66 is about as good as an extra person.

17.12 Pairwise comparisons and Elo ratings

People (and LLM judges) are more consistent at saying which of two answers is better than at giving absolute scores. Run many A-vs-B comparisons and report the **win rate**. To rank many models, convert the results into **Elo ratings** (as in chess) or the closely related Bradley–Terry scores; public leaderboards built from crowdsourced votes work this way. In Elo, the expected chance that A beats B is:

$$E_A = \frac{1}{1 + 10^{(R_B - R_A)/400}}$$

Equal ratings give 50%; a 100-point lead about 64%; 400 points about 91%.

17.13 Is the difference real?

Model A scores 78% on a 200-question eval and model B scores 74%. Is A better? A score on 200 questions has a standard error of about $\sqrt{(0.78 \times 0.22 / 200)} \approx 2.9$ points, so each score's 95% interval is roughly ± 5.7 points. A 4-point gap might be noise.

Because both models answered the *same* questions, a **paired** comparison is more sensitive. The bootstrap is simple and works for any metric:

```
import numpy as np

a = np.array(correct_a)  # 1 or 0 per question, model A
b = np.array(correct_b)  # same questions, model B
rng = np.random.default_rng(0)
diffs = []
for _ in range(10_000):
    idx = rng.integers(0, len(a), len(a))  # resample questions with replacement
    diffs.append(a[idx].mean() - b[idx].mean())
lo, hi = np.percentile(diffs, [2.5, 97.5])
print(f"A - B = {a.mean() - b.mean():.3f}, 95% interval [{lo:.3f}, {hi:.3f}]")
# if the interval includes 0, you can't be confident A is better
```

LLM outputs also vary from run to run, even at temperature 0. For important comparisons, run each question several times and average.

17.14 Writing a rubric that works

Vague rubrics ("rate helpfulness from 1 to 10") produce noisy, inconsistent scores from people and models alike. Better: a short checklist of specific, observable criteria, each pass/fail. For a customer-support answer about refunds:

1. States the correct refund window from the policy.
2. Invents no policy details that aren't in the source.
3. Asks for the order number if the customer didn't give one.
4. Is under 120 words.
5. Is polite and doesn't blame the customer.

The score is the number of criteria met, and each failed criterion tells you exactly what to fix. Include a worked example of a passing and a failing answer in the judge prompt, and test the judge on answers whose correct grades you already know before trusting it.

17.15 Structuring an evaluation harness

- **Dataset:** a versioned file (often JSON Lines: one test case per line) with inputs, reference answers or rubrics, and tags (topic, difficulty, "should refuse").
- **Runner:** calls your system with fixed settings (model, prompt version, temperature) and saves every output.
- **Graders**, cheapest first: code checks (exact match, regular expressions, JSON schema validity, unit tests for code), then an LLM judge for open-ended quality, then human spot checks on a sample.
- **Results store:** every run recorded with its model and prompt version, so you can compare any two runs.
- **Report:** overall scores, scores per tag, and a list of cases that changed between runs. The changed cases are where the insight is.
- **Gate:** in CI (Chapter 21.12), block a release if the score drops below a threshold or any critical case fails.

Try it

1. Write 20 test questions from your own work, with reference answers, and run two different models through `run_eval`.
2. Grade the same 20 answers yourself. How often do you agree with the judge?
3. Rerun the judge with a pairwise prompt ("Which answer is better, A or B?") and swap the order. Does the verdict change?

Chapter 18: Using LLMs — prompting, APIs, RAG and agents

Most AI products today aren't new models; they're systems built around an existing LLM. The three core techniques are **prompting** (asking well), **RAG** (giving the model your information) and **tool use** (letting it take actions).

18.1 Calling an LLM from code

```
# pip install anthropic (set ANTHROPIC_API_KEY; get a key at console.anthropic.com)
import anthropic
```

```
client = anthropic.Anthropic()
msg = client.messages.create(
    model="claude-sonnet-5",
    max_tokens=500,
    system="You are a patient tutor who explains AI to beginners.",
    messages=[{"role": "user", "content": "What is a transformer, in three sentences?"}],
)
print(msg.content[0].text)
```

Key ideas: the **system prompt** sets the role and rules; **messages** hold the conversation (the model has no memory — you resend the history each call); **max_tokens** caps the answer length; you pay per input and output token. Other providers (OpenAI, Google, open models via Hugging Face or Ollama) work almost identically.

18.2 Prompting that works

Technique	What it means	Example
Be clear and specific	State the task, audience, format and length	"Summarize for a busy manager in 3 bullets, max 15 words each"
Give context	Explain why and for whom	"This goes to new drivers who have never used the app"
Show examples (few-shot)	2–5 input/output examples of what good looks like	Three labelled reviews, then the new one
Structure with tags	Separate instructions from data	<document>...</document> then the question
Let it think	Ask for step-by-step reasoning before the final answer	"Work through it step by step, then give the answer in <answer> tags"
Specify the output format	JSON with named fields, for code to parse	"Return only JSON: {\"sentiment\": ..., \"confidence\": ...}"
Allow "I don't know"	Reduces invented answers	"If the document doesn't say, reply 'not stated'"
Iterate with evals	Change one thing, rerun your eval set (Chapter 17)	Prompt v1 vs v2 on 50 cases

18.3 Why RAG?

An LLM only knows its training data: nothing after its cutoff date and nothing private to you. It may also **hallucinate** — produce fluent, confident, wrong answers. **Retrieval-augmented generation (RAG)** fixes both: find the relevant passages from your documents, paste them into the prompt, and instruct the model to answer only from them, citing sources.

The RAG pipeline:

1. **Ingest**: load documents and split them into **chunks** (e.g. 200–500 words, with some overlap).
2. **Embed**: turn each chunk into a vector (Chapter 14) and store it in a vector database.
3. **Retrieve**: embed the user's question and fetch the most similar chunks.
4. **Generate**: give the LLM the question plus the retrieved chunks.

18.4 A minimal RAG system

```
# pip install sentence-transformers anthropic
import numpy as np, anthropic
from sentence_transformers import SentenceTransformer

docs = [
```

```

    "Employees get 15 vacation days per year, rising to 20 after five years.",
    "Expense reports must be submitted within 30 days with receipts attached.",
    "Remote work is allowed up to three days per week with manager approval.",
    "The office is closed on statutory holidays.",
]

embedder = SentenceTransformer("all-MiniLM-L6-v2")
doc_vecs = embedder.encode(docs, normalize_embeddings=True)

def retrieve(question, k=2):
    q = embedder.encode([question], normalize_embeddings=True)[0]
    scores = doc_vecs @ q # cosine similarity (vectors are
normalized)
    return [docs[i] for i in np.argsort(-scores)[:k]]

def answer(question):
    context = "\n".join(f"<doc>{d}</doc>" for d in retrieve(question))
    prompt = (f"Answer using ONLY the documents below. If they don't contain the
answer, "
              f"say you don't know.\n\n{context}\n\nQuestion: {question}")
    msg = anthropic.Anthropic().messages.create(
        model="claude-sonnet-5", max_tokens=300,
        messages=[{"role": "user", "content": prompt}])
    return msg.content[0].text

print(answer("How many vacation days do I get after six years?")) # 20
print(answer("What is the dress code?")) # should say it
doesn't know

```

Improving a RAG system, roughly in order of impact: better chunking (split on headings, keep tables whole), **hybrid search** (combine embeddings with keyword search such as BM25), a **reranker** model that re-scores the top 20–50 results, adding metadata filters (date, department), and evaluating retrieval and answers separately (Chapter 17.7).

18.5 Tool use and function calling

You can describe functions to the model — "get_weather(city)", "search_orders(customer_id)". Instead of answering directly, the model replies with a structured request to call one. Your code runs the function, sends the result back, and the model uses it in its answer. This lets LLMs get live data, do exact calculations and take actions. The **Model Context Protocol (MCP)** is an open standard for connecting models to tools and data sources in a reusable way.

18.6 Agents

An **agent** is an LLM running in a loop: decide the next step, call a tool, read the result, decide again, until the goal is done. Coding assistants that edit files and run tests, and research assistants that search and read many pages, are agents.

Design tips: give agents a small set of well-described tools, limit the steps they may take, ask for human confirmation before irreversible actions (sending email, spending money, deleting data), log every step, and evaluate on complete tasks, not single answers.

18.7 Choosing a model

Consideration	Trade-off
Quality	Larger, newer models are usually more capable, but slower and pricier
Latency and cost	Small models are fast and cheap; often enough for classification or extraction
Context window	How much text fits in one request; long documents need large windows or RAG
Privacy and control	Open-weight models can run on your own servers (Llama, Mistral, Qwen)
The real test	Your own eval set (Chapter 17), not a leaderboard

18.8 Anatomy of a strong prompt

A production prompt usually has these parts, in roughly this order:

ROLE	You are a support assistant for an online bookshop.
CONTEXT	Customers write in about orders, refunds and delivery. Replies are sent by email, so keep them friendly and complete.
TASK	Read the customer's message and the order record, then draft a reply.
RULES	- Use only facts from <order> and <policy>. Never invent dates or amounts. - If the order number is missing, ask for it and stop. - Escalate to a human (reply only ESCALATE) for legal threats or safety issues.
FORMAT	Plain text, under 150 words, signed "The Support Team".
EXAMPLES	<example> ...a message and an ideal reply... </example>
INPUT	<policy>{policy_text}</policy> <order>{order_record}</order> <message>{customer_message}</message>

Why each part helps: the **role and context** set the right vocabulary and tone; the **rules** state the hard limits, including what to do in edge cases; the **format** makes outputs predictable for the code that uses them; **examples** show what "good" looks like more precisely than any description; and **tags** keep untrusted input clearly separate from your instructions. Explaining *why* a rule exists ("replies are sent by email") helps the model handle cases the rules didn't anticipate.

18.9 Chunking strategies

Strategy	How	Good for
Fixed size with overlap	Every ~300 words, overlapping ~50 so sentences aren't cut off from their context	Quick start, uniform text

Strategy	How	Good for
Structure-aware	Split on headings, paragraphs, list items; keep tables whole	Manuals, policies, wikis
Semantic	Start a new chunk where the topic changes (embedding similarity drops)	Long, flowing documents
Small-to-big	Retrieve small chunks for precision, but give the LLM the surrounding section	Precise search with enough context

Smaller chunks match questions more precisely but may lack context; larger chunks carry more context but more noise. Always store **metadata** with each chunk (document title, section heading, date, source link) so answers can cite sources and searches can be filtered.

```
def chunk_words(text, size=300, overlap=50):
    words = text.split()
    step = size - overlap
    return [" ".join(words[i:i + size]) for i in range(0, max(len(words) - overlap, 1), step)]
```

18.10 Hybrid search and reranking

Embeddings find paraphrases ("time off" matches "vacation") but can miss exact terms such as product codes, error numbers and names. Keyword search (**BM25**) is the opposite. **Hybrid search** runs both and merges the rankings, commonly with **reciprocal rank fusion**: each document scores the sum of $1 / (60 + \text{its rank})$ across the rankings.

```
# pip install rank_bm25 (uses docs, embedder, doc_vecs from 18.4)
import numpy as np
from rank_bm25 import BM25Okapi

bm25 = BM25Okapi([d.lower().split() for d in docs])

def hybrid_retrieve(question, k=2):
    keyword_rank = np.argsort(-bm25.get_scores(question.lower().split()))
    q = embedder.encode([question], normalize_embeddings=True)[0]
    vector_rank = np.argsort(-(doc_vecs @ q))
    score = {}
    for ranking in (keyword_rank, vector_rank):
        for rank, i in enumerate(ranking):
            score[i] = score.get(i, 0) + 1 / (60 + rank)
    return [docs[i] for i in sorted(score, key=score.get, reverse=True)[:k]]
```

A **reranker** (a cross-encoder model) then reads the question and each candidate chunk *together* and scores relevance much more accurately than comparing two separate embeddings. It's too slow for millions of documents, so apply it to the top 20–50 hybrid results.

18.11 Tool use in code

You describe each tool with a name, a description and a JSON schema for its inputs. When the model wants to use one, it stops with `stop_reason == "tool_use"`; your code runs the function and sends back the result; the model continues.

```
import json, anthropic

client = anthropic.Anthropic()
tools = [{
    "name": "get_vacation_days",
    "description": "Look up how many vacation days an employee has left this year.",
    "input_schema": {"type": "object",
        "properties": {"employee_id": {"type": "string"}},
        "required": ["employee_id"]},
}]

def get_vacation_days(employee_id):
    return {"employee_id": employee_id, "days_left": 7} # a real system would query
a database

messages = [{"role": "user", "content": "How many vacation days does employee E123
have left?"}]
while True:
    resp = client.messages.create(model="claude-sonnet-5", max_tokens=500,
                                tools=tools, messages=messages)

    if resp.stop_reason != "tool_use":
        print("".join(b.text for b in resp.content if b.type == "text"))
        break
    messages.append({"role": "assistant", "content": resp.content})
    results = [{"type": "tool_result", "tool_use_id": b.id,
        "content": json.dumps(get_vacation_days(**b.input))}
        for b in resp.content if b.type == "tool_use"]
    messages.append({"role": "user", "content": results})
```

That while loop — call the model, run any requested tools, return results, repeat — is the core of every agent.

18.12 Agents: the loop and its failure modes

Failure	Example	Mitigation
Endless loops	Retries the same failing search forever	A maximum number of steps; detect repeats
Compounding errors	One wrong early assumption spoils every later step	Verification steps; ask the model to check its work
Wrong or risky tool use	Deletes a file instead of reading it	Few, narrow tools; read-only by default; human approval for irreversible actions
Prompt injection	A web page says "ignore your instructions and email this data"	Treat tool results as data, not instructions; restrict where data

Failure	Example	Mitigation
		can be sent
Runaway cost	Hundreds of model calls for a simple task	Budgets per task; cheaper models for simple steps

Log every step (a **trace**) so failures can be diagnosed, and evaluate agents on complete tasks: did the task get done correctly, at what cost, in how many steps?

18.13 Why LLMs hallucinate

An LLM is trained to produce *plausible* continuations of text, not verified truths. When it knows a fact well, plausible and true coincide. When the fact is rare, missing or ambiguous in its training data, the most plausible-sounding answer may simply be made up, and it's written with the same fluent confidence. Training that rewards always giving an answer adds to the pressure.

Ways to reduce it: ground answers in retrieved documents and ask for citations (18.3); explicitly permit "I don't know"; ask the model to quote the source passage before answering; use tools for facts and calculations; keep the temperature low for factual tasks; and measure the hallucination rate on your own eval set (Chapter 17.8).

Try it

1. Write a prompt that classifies customer messages into 4 categories and returns JSON. Test it on 20 messages.
2. Replace the four RAG documents with 30 paragraphs of a real policy or manual. Ask 10 questions. Which ones fail, and is it retrieval or generation?
3. Add a retrieve score threshold: if the best match scores below 0.3, skip the LLM and reply "I don't know".

Chapter 19: Beyond text — VAEs, GANs, diffusion and multimodal models

Generating images, audio and video uses the same deep learning foundations, but different recipes. Today's image generators are mostly **diffusion models**; earlier breakthroughs came from **autoencoders** and **GANs**.

19.1 Autoencoders and VAEs

An **autoencoder** is two networks: an **encoder** squeezes an image into a small vector (the **latent** code), and a **decoder** rebuilds the image from it. Trained to reproduce its input, it learns a compressed description of the data.

A **variational autoencoder (VAE)** makes the latent space smooth and organized, so you can pick a random point in it and decode a brand-new, plausible image. VAEs produce slightly blurry images alone, but they're a key part of modern systems: Stable Diffusion runs its diffusion process inside a VAE's compact latent space.

19.2 GANs: forger versus detective

A **generative adversarial network (GAN)** (2014) trains two networks against each other:

- The **generator** turns random noise into fake images.
- The **discriminator** tries to tell real images from fakes.

As the detective improves, the forger must improve too. GANs produced the first photorealistic fake faces (StyleGAN), but training is notoriously unstable, and they can get stuck making only a few kinds of output (**mode collapse**).

19.3 Diffusion models

A **diffusion model** learns to reverse a gradual noising process:

1. **Forward process** (fixed, no learning): take a real image and add a little random noise, step by step, over e.g. 1,000 steps, until it's pure static.
2. **Training**: show the network a noisy image and the step number, and train it to predict the noise that was added. The loss is just mean squared error (Chapter 3) between the predicted and the real noise.
3. **Generation**: start from pure random noise and repeatedly subtract the predicted noise. An image gradually appears.

$$L = \|\epsilon - \epsilon_{\theta}(x_t, t)\|^2$$

ϵ is the real noise, ϵ_{θ} is the network's guess, x_t is the noisy image at step t . The network is usually a U-Net or, in newer models, a transformer (**diffusion transformer**).

19.4 Guiding with text

To make "a cat astronaut in watercolour", the prompt is turned into embeddings by a text encoder (often **CLIP**, which learned a shared space for images and captions from hundreds of millions of pairs). The denoising network attends to those embeddings through **cross-attention** at every step. **Classifier-free guidance** strengthens how closely the image follows the text.

```
# pip install diffusers transformers accelerate (needs a GPU runtime)
import torch
from diffusers import AutoPipelineForText2Image

pipe = AutoPipelineForText2Image.from_pretrained(
    "stabilityai/sd-turbo", torch_dtype=torch.float16).to("cuda")
image = pipe("a watercolour painting of a lighthouse at dawn",
             num_inference_steps=1, guidance_scale=0.0).images[0] # sd-turbo is
distilled to need 1 step
image.save("lighthouse.png")
image
```

19.5 Other modalities

Modality	Common approach	Examples of tasks
Speech to text	Transformer encoder-decoder	Transcription (Whisper)

Modality	Common approach	Examples of tasks
	on audio spectrograms	
Text to speech	Transformers or diffusion over audio tokens	Voice assistants, narration
Music	Transformers over audio tokens, diffusion	Song generation
Video	Diffusion transformers over space-time patches	Text-to-video
Vision-language	An image encoder feeding a text LLM	Describing photos, reading charts, answering questions about documents

Multimodal LLMs read images (and sometimes audio) alongside text, by turning the image into a sequence of "visual tokens" that the transformer handles like words. That's how an assistant can read a screenshot or a scanned receipt.

19.6 Comparing generative model families

Family	Strengths	Weaknesses
Autoregressive transformers	Excellent for text and code; flexible	Generate one token at a time, so long outputs are slow
VAEs	Stable training, useful latent space	Blurry outputs on their own
GANs	Very fast generation, sharp images	Unstable training, mode collapse
Diffusion	Highest image quality, diverse, controllable	Many denoising steps (distillation reduces this)

19.7 The VAE loss and the reparameterization trick

A VAE's encoder doesn't output a single latent point. It outputs a *distribution*: a mean μ and a spread σ for each latent number. We then sample z from it and decode.

Sampling is random, and you can't backpropagate through a coin toss. The **reparameterization trick** rewrites the sample so the randomness sits outside the network:

$$z = \mu + \sigma \odot \epsilon, \quad \epsilon \sim \mathcal{N}(0, I)$$

Now μ and σ are ordinary outputs that gradients can flow through. The loss has two parts:

$$L = \underbrace{\|x - \hat{x}\|^2}_{\text{reconstruction}} + \underbrace{\text{KL}(\mathcal{N}(\mu, \sigma^2) \parallel \mathcal{N}(0, I))}_{\text{keep the latent space tidy}}$$

The reconstruction term makes outputs resemble inputs. The KL term pulls every encoded distribution towards a standard bell curve, so the latent space has no empty gaps: any random point you sample there decodes into something plausible. The tension between the two terms is also why plain VAE images are a little blurry.

19.8 The GAN's minimax game

The discriminator D tries to output 1 for real images and 0 for fakes; the generator G tries to make D output 1 for its fakes:

$$\min_G \max_D \mathbb{E}_{x \sim \text{data}} [\log D(x)] + \mathbb{E}_{z \sim \text{noise}} [\log (1 - D(G(z)))]$$

Training alternates: one step improving D on a batch of real and fake images, one step improving G. In practice the generator maximizes $\log D(G(z))$ instead, which gives stronger gradients early on when its fakes are easy to spot. GANs are hard to train because the two networks must stay balanced: if D wins too easily, G gets no useful signal; if G finds one output that fools D, it may produce only that (**mode collapse**). Variants such as Wasserstein GANs and spectral normalization made training more stable.

19.9 Noise schedules and the one-step shortcut

The forward process doesn't need to be run step by step. A noisy image at any step t can be produced in one jump from the clean image x_0 :

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

$\bar{\alpha}_t$ ("alpha bar") falls from almost 1 (barely any noise) at $t = 1$ to almost 0 (pure noise) at the last step. The **noise schedule** is how fast it falls; a cosine-shaped schedule works better than a straight line.

Thanks to the shortcut, training is simple:

```
# one diffusion training step (sketch)
t = torch.randint(0, T, (batch_size,))          # a random step for each
image                                           # the noise we add
eps = torch.randn_like(x0)
ab = alpha_bar[t].view(-1, 1, 1, 1)
xt = ab.sqrt() * x0 + (1 - ab).sqrt() * eps      # jump straight to step t
loss = F.mse_loss(model(xt, t), eps)            # predict the noise
```

Generation originally took about 1,000 denoising steps. Better samplers (such as DDIM) cut that to 20–50, and **distilled** models like the sd-turbo in 19.4 need only 1–4.

19.10 Latent diffusion

A 512×512 colour image has 786,432 numbers, and running a large network on all of them at every step is slow. **Latent diffusion** first uses a VAE (19.1) to compress the image to a $64 \times 64 \times 4$ latent of 16,384 numbers — 48 times smaller — runs the whole diffusion process in that compact space, and decodes back to pixels only at the end. This is the "Stable" in Stable Diffusion: it made high-quality generation affordable on ordinary GPUs.

19.11 Classifier-free guidance

During training, the text prompt is randomly replaced by an empty prompt some of the time (for example 10%), so one network learns both to denoise *with* the prompt and *without* it. At generation time, it predicts the noise both ways and pushes the result further in the direction the prompt points:

$$\hat{\epsilon} = \epsilon_{\text{uncond}} + w (\epsilon_{\text{cond}} - \epsilon_{\text{uncond}})$$

The **guidance scale** w (often 5–8 for standard models) trades variety for faithfulness: higher values follow the prompt more strictly but can look over-saturated and less varied. A **negative prompt** ("blurry, text, watermark") replaces the empty prompt, steering away from those features.

19.12 Evaluating image generators

- **FID** (Fréchet Inception Distance): compares the statistics of features extracted from real and generated images. Lower is better; it measures realism and variety together.
- **CLIP score**: how well each image matches its prompt, using CLIP's shared text–image space.
- **Human preference studies**: people choose the better of two images; still the most trusted measure.

Try it

1. Generate the same prompt with three different random seeds (`generator=torch.manual_seed(n)` in the pipe call). How different are the results?
2. Change one word in the prompt at a time and observe which words have the most influence.
3. Explain in your own words why predicting noise is an easier training task than generating an image in one step.

Chapter 20: Reinforcement learning basics

Reinforcement learning (RL) learns by trial and error. Nobody gives the right answer; an **agent** acts in an **environment** and receives **rewards**, and it learns which actions lead to the most reward over time. RL taught computers to beat world champions at Go, and it's how LLMs are tuned with human preferences and trained to reason (Chapter 16).

20.1 The vocabulary

Term	Meaning	Example: a robot vacuum
Agent	The learner and decision-maker	The vacuum's controller
Environment	Everything it interacts with	The room
State	The current situation	Position, battery, dirt map
Action	A choice the agent can make	Move forward, turn, return to dock
Reward	A number saying how good the last step was	+1 per dirty patch cleaned, −10 for hitting stairs

Term	Meaning	Example: a robot vacuum
Policy	The agent's strategy: state → action	"If battery low, return to dock"
Value	Expected total future reward from a state	A room corner with lots of dirt has high value
Episode	One run from start to finish	One cleaning session

The loop: observe the state → choose an action → get a reward and a new state → learn → repeat.

20.2 Key ideas

- **Delayed reward:** the move that wins a chess game may come 30 moves before the win. The agent must learn which earlier actions deserve credit.
- **Discount factor (γ):** future rewards count a bit less than immediate ones ($\gamma = 0.99$ means a reward 100 steps away is worth about 37%).
- **Exploration vs exploitation:** try new actions to discover better options, or use the best one known so far? The common **ϵ -greedy** rule explores randomly with probability ϵ , which shrinks as the agent learns.
- **Markov decision process (MDP):** the maths framework behind RL — states, actions, transition probabilities, rewards.

20.3 Q-learning

Q-learning keeps a table $Q(\text{state}, \text{action})$ estimating the total future reward of taking an action in a state. After each step it nudges the estimate towards "the reward I just got + the best value I can reach from the new state":

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

α is the learning rate, r the reward, s' the new state.

20.4 Code: learning to cross a frozen lake

In FrozenLake, an agent walks a 4×4 grid from start to goal without falling into holes. The only reward is +1 for reaching the goal.

```
# pip install gymnasium
import numpy as np
import gymnasium as gym

env = gym.make("FrozenLake-v1", is_slippery=False)
Q = np.zeros((env.observation_space.n, env.action_space.n)) # 16 states x 4 actions
alpha, gamma, epsilon = 0.1, 0.99, 1.0
rng = np.random.default_rng(0)

for episode in range(2000):
    state, _ = env.reset(seed=episode)
    done = False
    while not done:
```

```

    if rng.random() < epsilon:
        action = env.action_space.sample()          # explore
    else:
        action = int(np.argmax(Q[state]))           # exploit
    new_state, reward, terminated, truncated, _ = env.step(action)
    target = reward + gamma * np.max(Q[new_state]) * (not terminated)
    Q[state, action] += alpha * (target - Q[state, action])
    state, done = new_state, terminated or truncated
    epsilon = max(0.05, epsilon * 0.995)           # explore less over time

# test the learned policy
wins = 0
for _ in range(100):
    state, _ = env.reset(); done = False
    while not done:
        state, reward, terminated, truncated, _ = env.step(int(np.argmax(Q[state])))
        done = terminated or truncated
    wins += reward
print("wins out of 100:", wins)    # 100

```

20.5 Deep RL and policy methods

A table works for 16 states, not for the pixels of a video game. **Deep Q-Networks (DQN)** replace the table with a neural network that estimates Q from the raw state — this is how DeepMind's agents learned Atari games from pixels in 2013–2015.

Policy-gradient methods learn the policy directly: increase the probability of actions that led to high reward. **PPO** (proximal policy optimization) is a stable, popular version, and it's the classic algorithm behind RLHF: the LLM is the policy, generating a response is the action, and the reward model's score is the reward. **AlphaGo** combined deep networks with tree search and self-play.

20.6 When to use RL

RL fits problems of sequential decisions with a clear reward: games, robotics, recommendation strategies, resource scheduling, and fine-tuning LLMs. It's data-hungry and hard to debug, so if you have labelled examples of the right answer, supervised learning is usually simpler.

20.7 The Bellman equation

The value of a state equals the reward you get now plus the discounted value of where you end up. For the best possible behaviour:

$$Q^*(s, a) = \mathbb{E} \left[r + \gamma \max_{a'} Q^*(s', a') \right]$$

This is the **Bellman equation**, and the Q-learning update in 20.3 simply nudges Q towards its right-hand side. A worked example: in a corridor where only the goal pays reward 1 and $\gamma = 0.9$, the square next to the goal is worth 1, the one before it 0.9, then 0.81, then 0.73. Value "flows backwards" from the reward, one step at a time, and the best policy is just "move to the neighbour with the highest value".

20.8 Families of RL methods

Distinction	Option A	Option B
What is learned	Value-based: learn Q, act greedily (Q-learning, DQN)	Policy-based: learn the policy directly (REINFORCE, PPO)
Which experience is used	Off-policy: learn the best policy while behaving differently, can reuse old experience (Q-learning)	On-policy: learn only from the current policy's own fresh experience (SARSA, PPO)
Is there a model of the world?	Model-free: learn from trial and error only	Model-based: learn or use a model of how the environment responds, and plan with it (AlphaZero uses the known rules; MuZero learns its own model)

20.9 Policy gradients: REINFORCE and actor–critic

A policy-gradient method makes the policy a neural network $\pi_\theta(a \mid s)$ that outputs action probabilities. **REINFORCE**, the simplest, runs a full episode and then raises the log-probability of each action in proportion to the total reward G that followed it:

$$\nabla_\theta J = \mathbb{E} \left[\sum_t G_t \nabla_\theta \log \pi_\theta(a_t \mid s_t) \right]$$

Actions in good episodes become more likely; actions in bad ones less likely. It works, but the updates are very noisy: a good action in an unlucky episode gets punished. The fix is to compare against a **baseline** — how good was this outcome *relative to what we expected?* — giving the **advantage** $A = G - V(s)$.

Actor–critic methods learn both at once: the **actor** is the policy, and the **critic** is a network estimating $V(s)$ to supply the baseline. **PPO** adds a clip that stops any single update from changing the policy too much, which makes training stable; that's why it was the default for RLHF. Some recent LLM reasoning training uses **GRPO**, which samples several answers to the same prompt and uses their average reward as the baseline, removing the need for a separate critic network.

20.10 Reward hacking

An RL agent maximizes the reward you *wrote*, not the goal you *meant*. Famous examples: an agent in a boat-racing game discovered it scored more by circling forever to collect bonus items than by finishing the race; a language model rewarded by a reward model that slightly preferred longer answers learned to become verbose; a coding agent rewarded for passing tests may find ways to make tests pass without solving the problem.

This is **Goodhart's law**: when a measure becomes a target, it stops being a good measure. Defences: design rewards carefully and test them adversarially; combine several reward signals; keep a penalty for drifting far from sensible behaviour (the KL term in 16.9); and keep humans reviewing samples of what the agent actually does.

Try it

1. Set `is_slippery=True` (the agent sometimes slides the wrong way). How many wins now? Try more episodes.
2. Print `Q.argmax(axis=1).reshape(4, 4)` to see the learned direction in each square (0 = left, 1 = down, 2 = right, 3 = up).
3. Set `epsilon = 0` from the start. Why does the agent fail to learn?

Part 5 — Engineering and MLOps

Chapter 21: Git and GitHub for ML projects

Git records every version of your project's files, so you can go back in time, try ideas safely on a side branch, and work with others without overwriting each other. **GitHub** (or GitLab, Bitbucket) hosts Git projects online. Every professional AI team uses them.

21.1 Core concepts

Concept	Meaning
Repository (repo)	A project folder whose history Git tracks
Commit	A saved snapshot of changes, with a message explaining why
Staging area	The changes you've chosen to include in the next commit
Branch	An independent line of work; <code>main</code> is the default
Merge	Combine one branch's changes into another
Remote	A copy of the repo on a server, usually called <code>origin</code>
Push / pull	Send your commits to the remote / fetch and merge others'
Clone	Download a full copy of a remote repo
Pull request (PR)	A request on GitHub to merge a branch, where teammates review the code

21.2 Setup

Install Git from git-scm.com (it's pre-installed on macOS and most Linux; Colab has it too). Then tell it who you are, once:

```
git config --global user.name "Your Name"
git config --global user.email "you@example.com"
```

21.3 The everyday workflow

```
mkdir titanic-model && cd titanic-model
git init                    # start tracking this folder

# ... create train.py, requirements.txt ...
git status                  # what changed?
git add train.py requirements.txt # stage the files
git commit -m "Add baseline random forest training script"
git log --oneline           # see history
```

Connect to GitHub: create an empty repository on github.com, then:

```
git remote add origin https://github.com/<your-username>/titanic-model.git
git branch -M main
git push -u origin main
```

21.4 Branches and pull requests

Never experiment directly on `main`. Create a branch for each idea:

```
git switch -c try-gradient-boosting    # create and move to a new branch
# ... edit, run, evaluate ...
git add train.py
git commit -m "Try HistGradientBoosting: CV accuracy 0.83 vs 0.80"
git push -u origin try-gradient-boosting
```

On GitHub, open a **pull request** from that branch. A teammate reviews, automated tests run (21.7), and you merge. To update your local copy afterwards: `git switch main && git pull`.

If two people changed the same lines, Git reports a **merge conflict** and marks both versions in the file between `<<<<<<`, `=====` and `>>>>>>`. Edit the file to keep the right version, delete the markers, then `git add` and `git commit`.

21.5 Useful undo commands

Situation	Command
Discard unsaved changes to a file	<code>git restore train.py</code>
Unstage a file	<code>git restore --staged train.py</code>
See what changed	<code>git diff (unstaged) / git diff --staged</code>
Undo a pushed commit safely	<code>git revert <commit-id></code> (adds a new commit that reverses it)
Look at an old version	<code>git switch --detach <commit-id></code> , then <code>git</code>

Situation	Command
	<code>switch main to come back</code>

21.6 What not to commit

Git is for code and small text files. Keep these out with a `.gitignore` file in the repo's root:

```
# .gitignore
__pycache__/
.venv/
.ipynb_checkpoints/
data/          # large datasets
models/        # trained weights (*.pt, *.joblib)
.env           # API keys and passwords: NEVER commit secrets
wandb/
```

If you ever commit a secret (an API key), consider it leaked: revoke it and create a new one, even after deleting the file.

For data and models, use **DVC** (Data Version Control) or **Git LFS**, which store the big files elsewhere and keep a small pointer in Git; or a model registry such as MLflow or the Hugging Face Hub.

21.7 A typical ML repository

```
titanic-model/
├── README.md          # what it does, how to run it
├── requirements.txt    # exact library versions
├── .gitignore
├── notebooks/         # exploration
├── src/
│   ├── train.py       # trains and saves the model
│   ├── evaluate.py    # runs the evaluation set
│   └── serve.py       # the prediction API (Chapter 22)
├── tests/             # automated checks
├── Dockerfile         # Chapter 22
└── .github/workflows/ # CI: run tests on every pull request
```

Continuous integration (CI) with GitHub Actions runs your tests — and your model evaluations (Chapter 17) — automatically on every push, so a change that breaks something is caught before it's merged.

Write good commit messages: a short summary line in the imperative ("Add...", "Fix..."), and for experiments, the result ("CV accuracy 0.83").

21.8 Git's three areas, and what a commit really is

Your files live in three places:

1. **Working directory:** the files you edit.
2. **Staging area** (the *index*): changes you've marked with `git add` for the next commit.
3. **Repository:** the permanent history, written by `git commit`.

The staging area exists so you can commit related changes together and leave unrelated ones out, even when they're in the same folder.

A **commit** is a snapshot of the whole project, plus the author, date, message and a pointer to the previous commit (its **parent**). It's identified by a hash such as `a1b2c3d`. A **branch** is just a movable label pointing at a commit; committing moves it forward. **HEAD** is the label for "where you are now". Once this clicks, most Git commands become easy to reason about: switching branches moves HEAD; merging creates a commit with two parents.

21.9 Merge vs rebase

Your branch and `main` have both moved on. Two ways to combine them:

```
git switch try-gradient-boosting
git merge main      # adds a merge commit that joins the two histories
# or
git rebase main     # replays your commits on top of the latest main, as if you'd
started there
```

- **Merge** keeps the true history, including when branches split and joined. Safe at any time.
- **Rebase** produces a tidy, straight-line history, but it rewrites your commits (new hashes).
- **The golden rule**: never rebase commits that others have already pulled. Rebase only your own local work.

On GitHub, **squash and merge** combines all of a pull request's commits into one tidy commit on `main` — a popular choice for small teams.

21.10 Tags and releases

A **tag** gives a permanent name to one commit, usually a release:

```
git tag -a v1.0.0 -m "First production model: random forest, CV accuracy 0.82"
git push origin v1.0.0
```

Use **semantic versioning**, MAJOR.MINOR.PATCH: increase PATCH for fixes, MINOR for compatible improvements (a retrained model with the same inputs), MAJOR for breaking changes (the API's inputs change). Tag the exact code that produced each deployed model, and record that tag in your model registry and Docker image tag (`titanic-api:1.0.0`), so you can always trace a prediction back to the code and data behind it.

21.11 Branching strategies

Strategy	How it works	Good for
GitHub flow	<code>main</code> is always deployable; every change is a short-lived branch merged by pull request	Most teams; start here
Trunk-based development	Everyone commits to <code>main</code> very often, hiding unfinished work behind feature flags	Mature teams with strong automated tests
Git flow	Long-lived <code>develop</code> and release branches alongside <code>main</code>	Software with scheduled, versioned releases

For ML work, give each experiment its own branch and put the results (metrics, plots, the data version used) in the pull request description. The PR history then doubles as an experiment log.

21.12 A GitHub Actions workflow

GitHub Actions runs a workflow file on every push or pull request. This one installs the project, trains the model, runs the tests, and fails the pull request if quality drops:

```
# .github/workflows/ci.yml
name: CI
on: [push, pull_request]

jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-python@v5
        with:
          python-version: "3.12"
      - run: pip install -r requirements.txt pytest
      - run: python src/train.py
      - run: pytest tests/
      - run: python src/evaluate.py --min-accuracy 0.78 # exits with an error below
the bar
    env:
      ANTHROPIC_API_KEY: ${ secrets.ANTHROPIC_API_KEY } # only if evaluation
calls an LLM judge
```

A simple test file:

```
# tests/test_model.py
import joblib, pandas as pd

def test_model_returns_a_probability():
    model = joblib.load("model.joblib")
    passenger = pd.DataFrame([{"pclass": 1, "sex": "female", "age": 29,
                              "sibsp": 0, "fare": 100, "embarked": "S"}])
    p = model.predict_proba(passenger)[0, 1]
    assert 0.0 <= p <= 1.0

def test_women_in_first_class_more_likely_to_survive_than_men_in_third():
    model = joblib.load("model.joblib")
    rows = pd.DataFrame([
        {"pclass": 1, "sex": "female", "age": 30, "sibsp": 0, "fare": 80, "embarked":
"S"},
        {"pclass": 3, "sex": "male", "age": 30, "sibsp": 0, "fare": 8, "embarked":
"S"},
    ])
    p = model.predict_proba(rows)[:, 1]
    assert p[0] > p[1] # a "sanity" test: the model must respect a known pattern
```

Store secrets such as API keys under the repository's *Settings* → *Secrets and variables* → *Actions*, never in the workflow file.

Try it

1. Put your Chapter 8 code into a repo, commit it, and push it to GitHub.
2. Create a branch, change the model, open a pull request to yourself and merge it.
3. Deliberately create a merge conflict (edit the same line on two branches) and resolve it.

Chapter 22: Docker — packaging a model

"It works on my machine" is the classic deployment failure: a different Python version, a missing library.

Docker packages your code, its libraries and a minimal operating system into one **image** that runs identically on any computer or cloud.

22.1 Core concepts

Concept	Meaning	Analogy
Image	A read-only package: OS + libraries + your code + model	A recipe, or a class in Python
Container	A running instance of an image, isolated from the rest of the machine	A dish cooked from the recipe; an object
Dockerfile	A text file of instructions for building an image	The recipe's steps
Layer	Each instruction adds a cached layer; unchanged layers are reused	Pre-prepared ingredients
Registry	A place to store and share images (Docker Hub, GitHub, AWS ECR, Google Artifact Registry)	A recipe library
Port mapping	Connects a port inside the container to one on your machine	A door to the kitchen
Volume	Storage that survives when a container is deleted	A pantry outside the kitchen

Containers are much lighter than virtual machines: they share the host's operating system kernel and start in seconds.

Install **Docker Desktop** from [docker.com](https://docs.docker.com/docker-for-mac/install/) (Mac, Windows, Linux). Docker doesn't run inside Colab, so this chapter runs on your own computer.

22.2 Step 1: train and save the model

```
# train.py
import joblib, seaborn as sns
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier

raw = sns.load_dataset("titanic")[["survived", "pclass", "sex", "age", "sibsp",
"fare", "embarked"]]
X, y = raw.drop(columns="survived", raw["survived"])
prep = ColumnTransformer([
    ("num", make_pipeline(SimpleImputer(strategy="median"), StandardScaler()),
["pclass", "age", "sibsp", "fare"]),
    ("cat", make_pipeline(SimpleImputer(strategy="most_frequent"),
OneHotEncoder(handle_unknown="ignore")), ["sex", "embarked"]),
])
model = Pipeline([("prep", prep), ("model", RandomForestClassifier(n_estimators=300,
max_depth=5, random_state=42))])
model.fit(X, y) # final model: trained on all data after
evaluation
joblib.dump(model, "model.joblib")
```

22.3 Step 2: a prediction API with FastAPI

An **API** lets other programs send data and get predictions over HTTP. **FastAPI** is the standard Python choice.

```
# serve.py
import joblib
import pandas as pd
from fastapi import FastAPI
from pydantic import BaseModel

app = FastAPI(title="Titanic survival API")
model = joblib.load("model.joblib") # load once, at start-up

class Passenger(BaseModel): # validates incoming data
    pclass: int
    sex: str
    age: float | None = None
    sibsp: int = 0
    fare: float
    embarked: str = "S"

@app.get("/health")
def health():
    return {"status": "ok"}

@app.post("/predict")
def predict(p: Passenger):
    X = pd.DataFrame([p.model_dump()])
    return {"survival_probability": round(float(model.predict_proba(X)[0, 1]), 3)}
```

With requirements.txt pinning the versions:

```
fastapi==0.115.*
uvicorn==0.32.*
scikit-learn==1.5.*
pandas==2.2.*
joblib==1.4.*
```

Pin the scikit-learn version to the one you trained with (python -c "import sklearn; print(sklearn.__version__)"); a saved model may not load in a different version.

22.4 Step 3: the Dockerfile

```
# Dockerfile
FROM python:3.12-slim

WORKDIR /app

# Install libraries first: this layer is cached until requirements.txt changes
COPY requirements.txt .
RUN pip install --no-cache-dir -r requirements.txt

# Then copy the code and model
COPY serve.py model.joblib ./

EXPOSE 8000
CMD ["uvicorn", "serve:app", "--host", "0.0.0.0", "--port", "8000"]
```

22.5 Step 4: build, run, test

```
python train.py                                # creates model.joblib
docker build -t titanic-api:1.0 .                # build the image
docker run -d -p 8000:8000 --name titanic titanic-api:1.0 # run it in the
background

curl http://localhost:8000/health
curl -X POST http://localhost:8000/predict \
  -H "Content-Type: application/json" \
  -d '{"pclass": 1, "sex": "female", "age": 29, "fare": 100}'
# {"survival_probability": 0.97}
```

Open <http://localhost:8000/docs> in a browser: FastAPI generates an interactive page where you can try the API.

22.6 Everyday Docker commands

Command	Does
<code>docker ps / docker ps -a</code>	List running / all containers
<code>docker logs -f titanic</code>	Follow a container's output
<code>docker exec -it titanic bash</code>	Open a shell inside a running container
<code>docker stop titanic / docker rm titanic</code>	Stop / delete a container

Command	Does
<code>docker images / docker rmi <image></code>	List / delete images
<code>docker tag titanic-api:1.0 <user>/titanic-api:1.0 then docker push <user>/titanic-api:1.0</code>	Publish an image to a registry

22.7 Good practice

- **Small images:** use `-slim` base images; add a `.dockerignore` (like `.gitignore`) to keep data, `.git` and notebooks out.
- **Version tags:** tag images `1.0`, `1.1`... not just `latest`, so you can roll back.
- **No secrets in images:** pass API keys at run time, `docker run -e ANTHROPIC_API_KEY=...`, or through a secrets manager.
- **Big models:** rather than baking multi-gigabyte LLM weights into the image, download them at start-up or mount them as a volume (`-v /models:/models`).
- **GPUs:** base the image on an NVIDIA CUDA or PyTorch image and run with `docker run --gpus all ...`. For LLMs, ready-made serving images such as vLLM or Hugging Face TGI give fast, batched inference.
- **Several services** (API + database + vector store): describe them together in a `docker-compose.yml` and start them all with `docker compose up`.

22.8 Dockerfile instructions

Instruction	What it does
FROM	The base image to start from, e.g. <code>python:3.12-slim</code>
WORKDIR	Sets (and creates) the folder that later instructions run in
COPY	Copies files from your project into the image
RUN	Runs a command while building, e.g. installing packages
ENV	Sets an environment variable inside the image
ARG	A variable available only while building (<code>docker build --build-arg</code>)
EXPOSE	Documents which port the app listens on
USER	Which user the app runs as (avoid root)
HEALTHCHECK	A command Docker runs periodically to check the app is healthy

Instruction	What it does
CMD	The default command when a container starts
ENTRYPOINT	A fixed command; CMD then supplies default arguments to it

22.9 How layer caching works

Each instruction creates a **layer**, and Docker caches it. When you rebuild, Docker reuses every layer up to the first one whose inputs changed, then rebuilds everything after it. So order instructions from *rarely changes* to *often changes*:

```
# Slow: any code change invalidates the COPY, so packages reinstall every build
COPY . .
RUN pip install -r requirements.txt
```

```
# Fast: packages reinstall only when requirements.txt changes
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY . .
```

With large ML libraries such as PyTorch, this ordering can turn a 10-minute rebuild into a few seconds.

22.10 Multi-stage builds and running as a non-root user

A **multi-stage build** uses one stage to build or install things and a second, clean stage that copies in only the results. Build tools and caches never reach the final image, so it's smaller and has fewer security vulnerabilities.

```
# Stage 1: install dependencies
FROM python:3.12 AS builder
WORKDIR /app
COPY requirements.txt .
RUN pip install --no-cache-dir --prefix=/install -r requirements.txt
```

```
# Stage 2: the slim image that actually runs
FROM python:3.12-slim
WORKDIR /app
COPY --from=builder /install /usr/local
COPY serve.py model.joblib ./
RUN useradd --create-home appuser
USER appuser # don't run as root
EXPOSE 8000
CMD ["uvicorn", "serve:app", "--host", "0.0.0.0", "--port", "8000"]
```

Running as a non-root user means that if someone found a flaw in your app, they would have far fewer permissions inside the container.

22.11 Networking, environment variables and storage

- **Ports:** `-p 8080:8000` maps port 8080 on your computer to port 8000 inside the container.
- **Networks:** containers on the same Docker network reach each other by name. With Compose (22.12), a service called `db` is simply reachable at the hostname `db`.

- **Environment variables** configure the same image differently for development and production: `docker run -e LOG_LEVEL=debug ...`, or many at once from a file with `--env-file .env`.
- **Storage**: a container's own files vanish when it's removed. Use a **named volume** (`-v dbdata:/var/lib/postgresql/data`) for data Docker manages, or a **bind mount** (`-v $(pwd)/data:/app/data`) to share a folder from your computer, handy during development.

22.12 Several services with Docker Compose

This runs the prediction API together with a PostgreSQL database for logging predictions:

```
# docker-compose.yml
services:
  api:
    build: .
    ports:
      - "8000:8000"
    environment:
      DATABASE_URL: postgresql://app:${DB_PASSWORD}@db:5432/predictions
    depends_on:
      - db
  db:
    image: postgres:16
    environment:
      POSTGRES_USER: app
      POSTGRES_PASSWORD: ${DB_PASSWORD}
      POSTGRES_DB: predictions
    volumes:
      - dbdata:/var/lib/postgresql/data
volumes:
  dbdata:
```

```
echo "DB_PASSWORD=change-me" > .env      # Compose reads .env automatically; keep it
out of Git
docker compose up -d --build              # build and start both services
docker compose logs -f api                # follow the API's logs
docker compose down                       # stop (add -v to also delete the database
volume)
```

The API connects to the hostname `db`, the service's name. The same pattern runs a RAG app with a vector database, or an LLM server next to your application.

Try it

1. Build and run the image, then call `/predict` for three different passengers.
2. Change `serve.py` to also return the predicted class, rebuild as `1.1`, and run both versions on different ports (`-p 8001:8000`).
3. Write a `.dockerignore` and compare image sizes with `docker images`.

Chapter 23: Kubernetes and MLOps

Docker runs one container on one machine. **Kubernetes** (often written **K8s**) runs many containers across many machines: it restarts crashed ones, spreads traffic between copies, adds copies when load rises,

and rolls out new versions without downtime. **MLOps** is the wider practice of keeping models reliable once they're in production.

23.1 Kubernetes concepts

Object	What it is
Cluster	The whole system: a control plane plus worker machines
Node	One machine (virtual or physical) in the cluster
Pod	The smallest unit: one or more containers running together
Deployment	Declares "keep N copies of this pod running, with this image"; handles updates and rollbacks
Service	A stable name and address that load-balances traffic to a Deployment's pods
Ingress	Routes outside HTTP traffic (a domain name) to Services
ConfigMap / Secret	Configuration and passwords/API keys, kept out of the image
HorizontalPodAutoscaler (HPA)	Adds or removes pods based on CPU or other metrics
Namespace	A folder-like partition, e.g. staging and production

The key idea is **declarative**: you describe the desired state in YAML files, and Kubernetes continually works to make reality match.

23.2 Try it locally

Install `kubectl` (the command-line tool) and **minikube** (a one-node cluster on your laptop) from kubernetes.io. In the cloud, managed services (Amazon EKS, Google GKE, Azure AKS) run the cluster for you.

```
minikube start
minikube addons enable metrics-server      # needed for autoscaling
minikube image load titanic-api:1.0       # make the Chapter 22 image available to
the cluster
```

23.3 Deploying the model API

```
# k8s.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
```

```

  name: titanic-api
spec:
  replicas: 2                                # two copies for reliability
  selector:
    matchLabels: {app: titanic-api}
  template:
    metadata:
      labels: {app: titanic-api}
    spec:
      containers:
      - name: api
        image: titanic-api:1.0
        imagePullPolicy: IfNotPresent
        ports:
        - containerPort: 8000
        resources:
          requests: {cpu: "250m", memory: "256Mi"} # guaranteed share
          limits: {cpu: "1", memory: "512Mi"}      # maximum allowed
        readinessProbe:                            # only send traffic when ready
          httpGet: {path: /health, port: 8000}
          initialDelaySeconds: 5
        livenessProbe:                              # restart if it stops responding
          httpGet: {path: /health, port: 8000}
          initialDelaySeconds: 10
---
apiVersion: v1
kind: Service
metadata:
  name: titanic-api
spec:
  selector: {app: titanic-api}
  ports:
  - port: 80
    targetPort: 8000
---
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: titanic-api
spec:
  scaleTargetRef: {apiVersion: apps/v1, kind: Deployment, name: titanic-api}
  minReplicas: 2
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target: {type: Utilization, averageUtilization: 70} # add pods above 70% CPU

```

```

kubectl apply -f k8s.yaml                # create everything
kubectl get pods                          # two pods, STATUS Running
kubectl port-forward service/titanic-api 8080:80
curl -X POST http://localhost:8080/predict -H "Content-Type: application/json" \
  -d '{"pclass": 3, "sex": "male", "fare": 8}'

```

```
kubectl delete pod <one-pod-name>          # watch Kubernetes replace it
automatically
kubectl set image deployment/titanic-api api=titanic-api:1.1 # rolling update
kubectl rollout undo deployment/titanic-api # roll back if something breaks
kubectl logs -l app=titanic-api             # logs from all pods
```

For GPU models, nodes with GPUs are added to the cluster and a pod requests one with resources: `{limits: {nvidia.com/gpu: 1}}`. Tools such as **KServe**, **Ray Serve** and **vLLM** add model-specific serving on top of Kubernetes.

23.4 The MLOps lifecycle

1. **Data**: versioned datasets and a documented pipeline for collecting and cleaning them.
2. **Experiment tracking**: log every run's code version, data version, hyperparameters and metrics (MLflow, Weights & Biases).
3. **Model registry**: store approved models with versions and stages (staging, production).
4. **CI/CD**: on every change, automatically run tests and the evaluation set; build the Docker image; deploy to staging, then production.
5. **Serving**: online (API, milliseconds), batch (score millions of rows overnight), or streaming (events as they arrive).
6. **Monitoring**: watch the system and the model (23.6).
7. **Retraining**: on a schedule or when monitoring shows decay (**continuous training**), then back to step 2.

23.5 Safe rollout strategies

Strategy	How it works	Why
Rolling update	Replace pods a few at a time	Kubernetes' default; no downtime
Canary	Send ~5% of traffic to the new model, compare, then increase	Limits damage from a bad model
Shadow	New model gets copies of real requests; its answers are logged, not used	Test on real data with zero user risk
A/B test	Split users between models and measure the business outcome	Proves the new model is actually better
Blue-green	Run old and new side by side, switch traffic at once	Instant rollback

23.6 Monitoring models in production

A model can be "up" and still be silently wrong. Monitor three levels:

- **System**: latency, error rate, throughput, CPU/GPU and memory use, cost.

- **Data: data drift** — incoming inputs no longer look like training data (new customer types, a changed sensor). Compare feature distributions with statistical tests or the population stability index (PSI).
- **Model: concept drift** — the relationship between inputs and outcome changes (fraud tactics evolve). Track accuracy once true outcomes arrive, and prediction distributions before they do.

For LLM applications (**LLMOps**), also version prompts like code, gate releases on your eval set (Chapter 17), trace each step of chains and agents, track token usage and cost, sample conversations for review, and apply guardrails (input/output filters) against prompt injection and unsafe content.

23.7 How Kubernetes works inside

A cluster has a **control plane** (the brain) and **worker nodes** (the muscle).

Component	Where	Job
API server	Control plane	The front door: kubectl and every other component talk only to it
etcd	Control plane	A database storing the cluster's desired and current state
Scheduler	Control plane	Chooses a node for each new pod, based on requested CPU, memory and GPUs
Controller manager	Control plane	Runs the loops that make reality match the desired state, e.g. replacing missing pods
kubelet	Every node	Starts the pod's containers on its node and reports their health
Container runtime	Every node	Actually runs containers (usually containerd)
kube-proxy	Every node	Routes Service traffic to the right pods

What happens on `kubectl apply -f k8s.yaml`: the API server validates the file and stores it in etcd; the Deployment controller notices that 2 pods are wanted and 0 exist, and creates 2 pod records; the scheduler assigns each to a node with enough free resources; each node's kubelet pulls the image and starts the container; once the readiness probe passes, the Service starts sending it traffic. If a node dies, the controller sees pods missing and the whole cycle repeats elsewhere. That **reconciliation loop** is the heart of Kubernetes.

23.8 ConfigMaps and Secrets

Keep configuration out of the image so one image runs in every environment:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: titanic-config
data:
  LOG_LEVEL: "info"
  MODEL_VERSION: "1.0.0"
```

```
# create a Secret from a local file of KEY=value lines (keeps the key out of your
shell history)
kubectl create secret generic titanic-secrets --from-env-file=secrets.env
Then load both into the container as environment variables, inside the Deployment's container section:
```

```
envFrom:
- configMapRef: {name: titanic-config}
- secretRef: {name: titanic-secrets}
```

By default, Secrets are only base64-encoded, not encrypted, so restrict who can read them; production clusters usually connect to a dedicated secrets manager (AWS Secrets Manager, HashiCorp Vault, and similar).

23.9 Ingress: a public address

A Service is reachable only inside the cluster. An **Ingress** routes outside HTTP traffic for a domain to it. It needs an **ingress controller** running in the cluster (in minikube: `minikube addons enable ingress`).

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: titanic-api
spec:
  ingressClassName: nginx
  rules:
  - host: titanic.example.com
    http:
      paths:
      - path: /
        pathType: Prefix
        backend:
          service:
            name: titanic-api
            port: {number: 80}
```

HTTPS certificates are usually added automatically with the **cert-manager** tool.

23.10 Helm: packages for Kubernetes

Real applications need many YAML files that differ slightly between staging and production. **Helm** packages them as a **chart**: templates with placeholders, plus a `values.yaml` file that fills them in.

```
helm install titanic ./titanic-chart -f values-production.yaml # install your own
chart
helm upgrade titanic ./titanic-chart --set image.tag=1.1.0      # roll out a new
version
helm rollback titanic 1                                          # go back to
```

```
revision 1
```

```
# install ready-made software, e.g. Prometheus + Grafana monitoring
helm repo add prometheus-community https://prometheus-community.github.io/helm-charts
helm install monitoring prometheus-community/kube-prometheus-stack
```

Prometheus collects metrics (request counts, latencies, CPU) and **Grafana** turns them into dashboards and alerts — the usual monitoring pair on Kubernetes.

23.11 Detecting data drift with PSI

The **population stability index (PSI)** compares a feature's distribution in production with its distribution in training. Split the training values into bins (often 10, by percentile), then compare the share of values in each bin:

$$\text{PSI} = \sum_{\text{bins}} (\text{actual\%} - \text{expected\%}) \times \ln \frac{\text{actual\%}}{\text{expected\%}}$$

A common rule of thumb: below 0.1, stable; 0.1–0.25, a moderate shift worth watching; above 0.25, a major shift that calls for investigation and probably retraining.

```
import numpy as np

def psi(expected, actual, bins=10):
    edges = np.quantile(expected, np.linspace(0, 1, bins + 1))
    edges[0], edges[-1] = -np.inf, np.inf
    e = np.histogram(expected, edges)[0] / len(expected)
    a = np.histogram(actual, edges)[0] / len(actual)
    e, a = np.clip(e, 1e-6, None), np.clip(a, 1e-6, None) # avoid log(0)
    return np.sum((a - e) * np.log(a / e))

rng = np.random.default_rng(0)
training_fares = rng.lognormal(3.0, 1.0, 5000)
this_week_same = rng.lognormal(3.0, 1.0, 1000)
this_week_shifted = rng.lognormal(3.5, 1.0, 1000) # fares have risen
print(round(psi(training_fares, this_week_same), 2)) # about 0.01: stable
print(round(psi(training_fares, this_week_shifted), 2)) # about 0.25: major shift
Run a check like this on a schedule for every important feature and for the model's output scores, and
alert when PSI crosses the threshold. Libraries such as Evidently produce full drift reports.
```

Try it

1. Deploy the API to minikube and call it through port-forward.
2. Scale manually with `kubectl scale deployment/titanic-api --replicas=4` and watch `kubectl get pods -w`.
3. Write down which drift you'd expect for a driver-risk model if a company switched from city delivery vans to long-haul trucks, and what you'd monitor.

Chapter 24: Responsible AI

A model that's accurate on average can still be unfair to some people, impossible to explain, leak private data, or be manipulated. Responsible AI is the set of practices that catch these problems before they cause harm. It's part of doing the job well, not an optional extra.

24.1 Where bias comes from

Source	Example
Historical bias	Past hiring data reflects past discrimination, and the model learns it
Representation bias	A face dataset with few darker-skinned faces performs worse on them
Measurement bias	Using "arrests" as a stand-in for "crime" when policing differs by neighbourhood
Proxy features	Postcode quietly standing in for ethnicity or income, even if those columns are removed
Feedback loops	A model's decisions shape the future data it's retrained on

24.2 Measuring fairness

Compute your metrics separately for each group (Chapter 9.7), and compare:

- **Demographic parity:** does each group receive positive predictions at a similar rate?
- **Equal opportunity:** among people who truly qualify, is the recall the same for each group?
- **Equalized odds:** are both the true-positive and false-positive rates similar across groups?
- **Calibration within groups:** does a 70% score mean 70% for everyone?

These definitions can conflict mathematically; you can't always satisfy all at once. Choosing which matters is a decision for people who understand the context, not only for engineers.

```
import pandas as pd
from sklearn.metrics import recall_score

# using best, X_test, y_test from Chapter 8
results = pd.DataFrame({"group": X_test["sex"], "y": y_test, "pred":
best.predict(X_test)})
for g, part in results.groupby("group"):
    print(g, "positive rate:", round(part["pred"].mean(), 2),
          "recall:", round(recall_score(part["y"], part["pred"]), 2))
```

The library **Fairlearn** automates these comparisons and offers mitigation methods.

24.3 Explainability

- **Global:** which features matter overall? Use feature importance (Chapter 6) or **permutation importance** — shuffle one column and see how much the score drops.
- **Local:** why *this* prediction? **SHAP** assigns each feature a contribution for each individual prediction; **LIME** fits a simple model around one example.
- **Inherently interpretable models:** for high-stakes decisions (credit, health, employment), a slightly less accurate but explainable model is often the right choice.

- **LLMs**: explaining large neural networks is an open research field (**mechanistic interpretability**). In practice, ask for cited sources (RAG) and log inputs and outputs for review.

```
from sklearn.inspection import permutation_importance

r = permutation_importance(best, X_test, y_test, n_repeats=10, random_state=42)
print(pd.Series(r.importances_mean,
index=X_test.columns).sort_values(ascending=False))
# pip install shap -> shap.Explainer for per-prediction explanations
```

24.4 Privacy

- Collect only the data you need, and remove or mask **personally identifiable information (PII)**: names, emails, IDs, precise locations.
- Models can **memorize** training data; LLMs have been shown to repeat rare training text. Don't train on secrets.
- **Differential privacy** adds carefully measured noise so no single person's data can be recovered from the model.
- **Federated learning** trains across devices or organizations without moving raw data to one place.
- Follow the law where you operate (e.g. GDPR in Europe, PIPEDA in Canada) and your organization's data policies.

24.5 Security

- **Adversarial examples**: tiny, invisible changes to an input that flip a model's prediction.
- **Data poisoning**: corrupting training data to plant a hidden behaviour.
- **Prompt injection**: text hidden in a web page, email or document telling an LLM to ignore its instructions. Treat all retrieved content as untrusted data, limit what tools an agent can use, and require human approval for sensitive actions.
- **Model and data theft**: protect API keys, rate-limit endpoints, log access.

24.6 Transparency and oversight

- **Model cards** document what a model is for, how it was trained and evaluated, its known limitations and groups it performs worse on. **Datasheets** do the same for datasets.
- **Human in the loop**: for consequential decisions, the model advises and a person decides, with a way for affected people to appeal.
- **Regulation** is growing: the EU AI Act, for example, sets requirements by risk level. Check the rules that apply to your country and industry.
- **Environmental cost**: large training runs use significant energy; prefer smaller models when they're good enough.

24.7 A fairness check, worked by hand

A loan model scores 200 applicants, 100 from group A and 100 from group B. In each group, 60 people would truly repay the loan.

	Group A	Group B
--	---------	---------

	Group A	Group B
Approved	55 (50 who would repay, 5 who wouldn't)	35 (30 who would repay, 5 who wouldn't)
Approval rate (demographic parity)	55%	35%
Recall among people who would repay (equal opportunity)	$50 / 60 = 0.83$	$30 / 60 = 0.50$
False-positive rate	$5 / 40 = 0.125$	$5 / 40 = 0.125$
Precision	$50 / 55 = 0.91$	$30 / 35 = 0.86$

The false-positive rates are equal, and precision looks similar, so a quick glance might say "fair enough". But a qualified applicant in group B is approved only half the time, compared with 83% in group A. That's an equal-opportunity failure that a single overall accuracy number would never show. The ratio of approval rates, $0.35 / 0.55 \approx 0.64$, is also below the **four-fifths (0.8) guideline** that US employment regulators use as a red flag for adverse impact.

Next steps would be to look for proxy features, check whether group B is under-represented or differently labelled in the training data, and decide with legal and domain experts which fairness criterion applies and how to correct it.

24.8 How SHAP works

SHAP is based on **Shapley values** from game theory: a fair way to split a team's payout among its players, by averaging each player's extra contribution over every possible order in which they could have joined the team. In SHAP the "players" are features and the "payout" is the difference between this prediction and the average prediction.

The contributions always add up exactly. For example (illustrative numbers): the average survival prediction is 0.38; for one passenger, sex = female adds +0.30, class = 1st adds +0.15, fare adds +0.05 and age adds +0.02, giving $0.38 + 0.52 = \mathbf{0.90}$.

```
# pip install shap    (forest and X_test from Chapter 6)
import shap

explainer = shap.TreeExplainer(forest)
sv = explainer(X_test)
shap.plots.waterfall(sv[0, :, 1])    # why passenger 0 got their survival score
(class 1 = survived)
shap.plots.beeswarm(sv[:, :, 1])    # which features matter most, across all
passengers
```

Two cautions. SHAP explains what the *model* does, not what causes outcomes in the real world. And when features are strongly correlated (fare and class), credit can be split between them in surprising ways.

24.9 A red-teaming checklist for LLM applications

Test each category deliberately, and turn every failure you find into a permanent evaluation case (Chapter 17.5):

- ☐ **Harmful requests:** does it refuse what it should?
- ☐ **Over-refusal:** does it wrongly refuse ordinary, legitimate requests?
- ☐ **Prompt injection:** instructions hidden in uploaded documents, web pages, emails or tool results.
- ☐ **Data leakage:** revealing its system prompt, other users' data, or personal information from documents.
- ☐ **Hallucination:** unanswerable questions, made-up citations, fake policy details.
- ☐ **Bias:** identical requests with different names, genders, dialects or nationalities.
- ☐ **Jailbreaks:** role-play, "hypothetically", encoded text, many-turn persuasion.
- ☐ **Tool misuse:** can it be talked into an action it shouldn't take?
- ☐ **Robustness:** typos, other languages, very long or empty inputs.

24.10 Designing human oversight

Level	What the person does	Suitable for
Human in the loop	Approves every decision before it takes effect	High stakes: medical, credit, hiring, safety
Human on the loop	Monitors and can step in or reverse decisions	Medium stakes, high volume
Fully automated	Reviews samples and metrics after the fact	Low stakes: spam filtering, recommendations

Oversight only works if it's real. Watch for **automation bias**, where reviewers approve whatever the system suggests. Show the model's confidence and reasons; make overriding it quick; track how often people override (near zero may mean rubber-stamping); give affected people a way to appeal; and name an owner who is accountable for the system's outcomes.

Try it

1. Run the fairness code above. Is the model's recall similar for men and women? What might explain a gap?
2. Write a one-page model card for your Titanic model: purpose, data, metrics, limitations.
3. Think of an AI system at your workplace. Who could it treat unfairly, and how would you test for it?

Appendices

Appendix A: Math essentials

You can build useful AI with very little math, but four areas help you understand what's happening: linear algebra (how data is represented), calculus (how models learn), probability and statistics (how uncertainty is handled), and a little information theory (the losses). Each idea below links to where it appears in the book.

A.1 Linear algebra: data as vectors and matrices

- A **vector** is a list of numbers: one passenger's features, or one token's embedding (Chapter 14).
- A **matrix** is a grid of numbers: a whole dataset (rows × features), or a layer's weights (Chapter 10).
- The **dot product** multiplies two vectors element by element and adds up: $a \cdot b = a_1b_1 + a_2b_2 + \dots$. It measures how aligned they are; it's the heart of every neuron and of attention ($Q \cdot K$, Chapter 15).
- **Matrix multiplication** applies many dot products at once. A neural network layer is one matrix multiply plus an activation.
- The **norm** (length) of a vector, and **cosine similarity** = dot product divided by both lengths (Chapter 14).

```
import numpy as np
a, b = np.array([1, 2, 3]), np.array([4, 5, 6])
print(a @ b) # dot product: 32
print(np.linalg.norm(a)) # length: 3.74
print(a @ b / (np.linalg.norm(a) * np.linalg.norm(b))) # cosine similarity: 0.97
W = np.random.randn(3, 2) # a layer: 3 inputs -> 2
outputs
print(a @ W) # the layer's output
```

A.2 Calculus: how learning works

- A **derivative** is the slope: how fast the output changes when you nudge the input.
- A **partial derivative** is the slope with respect to one input while the others stay fixed.
- The **gradient** is the vector of all partial derivatives: it points uphill. Gradient descent steps the opposite way (Chapter 3).
- The **chain rule** says the slope of a chain of functions is the product of their slopes. Backpropagation is the chain rule applied layer by layer (Chapter 10).

$$\frac{d}{dx}f(g(x)) = f'(g(x)) \cdot g'(x)$$

A.3 Probability

- A **probability distribution** assigns likelihoods to outcomes. An LLM's output is a distribution over the next token (Chapter 15).
- **Conditional probability** $P(A | B)$: the chance of A given that B happened.
- **Bayes' theorem** updates beliefs with evidence; it's behind naive Bayes (Chapter 5):

$$P(A | B) = \frac{P(B | A) P(A)}{P(B)}$$

- **Expected value** is the probability-weighted average outcome; **variance** measures spread.
- The **normal (Gaussian) distribution** is the bell curve; noise in Chapters 3 and 19 is drawn from it.

A.4 Statistics

- **Mean vs median**: the median resists outliers (Chapter 4).
- **Standard deviation**: typical distance from the mean; used in standardization.
- **Correlation**: how two variables move together, from -1 to 1 . Correlation is not causation.
- **Sampling and confidence intervals**: a test score from 178 examples is an estimate with uncertainty (Chapter 9.7).
- **Hypothesis testing and p-values**: is a difference (for example, in an A/B test) likely to be real or chance?

A.5 Information theory

- **Entropy**: how uncertain a distribution is. A fair coin has high entropy; a two-headed coin has zero. Used to choose tree splits (Chapter 6).
- **Cross-entropy**: how surprised your model's distribution is by the true outcome. It's the loss for classification and for LLMs (Chapters 5, 15, 16).
- **KL divergence**: how different two distributions are. RLHF uses it to keep the tuned model close to the original (Chapter 16).

$$H(p, q) = - \sum_x p(x) \log q(x)$$

A.6 Worked examples, by hand

Matrix multiplication. Each entry of the result is the dot product of a row of the first matrix with a column of the second:

$$\begin{bmatrix} 1 & 2 \\ 3 & 4 \end{bmatrix} \begin{bmatrix} 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 1 \cdot 5 + 2 \cdot 7 & 1 \cdot 6 + 2 \cdot 8 \\ 3 \cdot 5 + 4 \cdot 7 & 3 \cdot 6 + 4 \cdot 8 \end{bmatrix} = \begin{bmatrix} 19 & 22 \\ 43 & 50 \end{bmatrix}$$

The shape rule: $(m \times n)$ times $(n \times p)$ gives $(m \times p)$, and the inner numbers must match. A neural network layer taking 3 features to 2 neurons is (1×3) times $(3 \times 2) = (1 \times 2)$. A "shapes cannot be multiplied" error means the inner numbers differ.

A derivative, two ways. For $f(x) = x^2$, the rule says the derivative is $2x$, so at $x = 3$ the slope is 6. Check numerically by nudging x a tiny amount: $(f(3.001) - f(3)) / 0.001 = (9.006001 - 9) / 0.001 \approx 6.001$.

The chain rule. For $f(x) = (2x + 1)^2$, treat it as an outer function (something squared) and an inner one $(2x + 1)$. Outer derivative: $2 \times (2x + 1)$. Inner derivative: 2. Multiply: $4(2x + 1)$. At $x = 1$ that's 12. Numerical check: $((3.002)^2 - 9) / 0.001 \approx 12.004$. Backpropagation (Chapter 10) is this, repeated through every layer.

Bayes and the medical test. A disease affects 1% of people. A test catches 90% of real cases and wrongly flags 5% of healthy people. You test positive. How likely are you to have it?

$$P(\text{disease} | +) = \frac{0.90 \times 0.01}{0.90 \times 0.01 + 0.05 \times 0.99} = \frac{0.009}{0.0585} \approx 0.15$$

Only about 15%. In counts: out of 10,000 people, 100 are sick and 90 of them test positive; 9,900 are healthy and 495 of them test positive. So 90 of the 585 positives are real. This is the same effect as a classifier's low precision on a rare class (Chapter 9): when the condition is rare, most alarms are false even with a good test.

Mean, variance and standard deviation. Data: 2, 4, 4, 4, 5, 5, 7, 9. Mean = $40 / 8 = 5$. Squared distances from the mean: 9, 1, 1, 1, 0, 0, 4, 16, totalling 32. Variance = $32 / 8 = 4$. Standard deviation = $\sqrt{4} = 2$. Standardizing the value 9 gives $(9 - 5) / 2 = 2$: it's two standard deviations above the mean.

Entropy and cross-entropy of coins (in bits, using log base 2):

- A fair coin: $-(0.5 \log_2 0.5 + 0.5 \log_2 0.5) = \mathbf{1 \text{ bit}}$. Maximum uncertainty.
- A 90/10 coin: $-(0.9 \log_2 0.9 + 0.1 \log_2 0.1) \approx 0.137 + 0.332 = \mathbf{0.47 \text{ bits}}$. Fairly predictable.
- A two-headed coin: **0 bits**. No uncertainty at all.
- If the coin is really fair but your model believes 90/10, the cross-entropy is $-(0.5 \log_2 0.9 + 0.5 \log_2 0.1) \approx \mathbf{1.74 \text{ bits}}$. The extra 0.74 bits over the true entropy is the KL divergence: the cost of wrong beliefs. Training a classifier or an LLM with cross-entropy loss is exactly the process of shrinking that extra cost.

```
import numpy as np
print(np.array([[1, 2], [3, 4]]) @ np.array([[5, 6], [7, 8]])) # [[19 22] [43 50]]
print(np.std([2, 4, 4, 4, 5, 5, 7, 9])) # 2.0
print(-(0.5 * np.log2(0.9) + 0.5 * np.log2(0.1))) # 1.74
```

A.7 Where to learn more

- [3Blue1Brown](#): visual video series on linear algebra, calculus and neural networks.
- [Khan Academy](#): free courses on every topic above, from the basics.
- *Mathematics for Machine Learning* by Deisenroth, Faisal and Ong: a free textbook written exactly for this purpose.

Appendix B: The concept checklist and where to go next

This is the full list of concepts an AI/ML student is expected to know, mapped to the chapter that teaches each one. Tick them off as you can explain them in your own words and write the code without looking.

B.1 Foundations (Chapters 1–3, Appendix A)

- ☐ AI vs ML vs deep learning vs generative AI
- ☐ Model, parameters, training, inference
- ☐ Python basics, NumPy, pandas, Matplotlib
- ☐ Loss functions (MSE, cross-entropy)
- ☐ Gradient descent and learning rate
- ☐ Vectors, matrices, dot products; derivatives, gradients, chain rule; probability, Bayes' theorem; basic statistics

B.2 Classic machine learning (Chapters 4–9)

- ☐ Features, labels, data types; EDA
- ☐ Missing values, outliers, encoding, scaling, feature engineering
- ☐ Train / validation / test split; data leakage
- ☐ Supervised, unsupervised, self-supervised, reinforcement learning
- ☐ Linear and logistic regression, Ridge, Lasso
- ☐ kNN, SVM, naive Bayes
- ☐ Decision trees, random forests, gradient boosting (XGBoost, LightGBM)
- ☐ k-means, hierarchical clustering, DBSCAN; PCA, t-SNE, UMAP; anomaly detection
- ☐ Overfitting, underfitting, bias–variance trade-off, regularization
- ☐ Cross-validation, hyperparameter tuning, pipelines
- ☐ Confusion matrix, precision, recall, F1, ROC-AUC, PR-AUC; MAE, RMSE, R^2 ; baselines; thresholds; calibration

B.3 Deep learning (Chapters 10–13)

- ☐ Neurons, layers, activation functions
- ☐ Forward pass, backpropagation
- ☐ Epochs, batches, optimizers (SGD, Adam, AdamW), learning-rate schedules
- ☐ Dropout, weight decay, batch/layer norm, early stopping, data augmentation
- ☐ PyTorch tensors, autograd, the training loop, GPUs
- ☐ CNNs, pooling, transfer learning, fine-tuning
- ☐ RNNs, LSTMs, vanishing gradients, sequence-to-sequence

B.4 Generative AI and LLMs (Chapters 14–19)

- ☐ Tokenization (BPE), embeddings, semantic search, vector databases
- ☐ Self-attention (Q, K, V), multi-head attention, causal masking, positional encoding
- ☐ Transformer blocks, residual connections; encoder vs decoder models
- ☐ Sampling: temperature, top-k, top-p; scaling laws
- ☐ Pre-training, SFT, RLHF, DPO, reasoning via RL
- ☐ LoRA, QLoRA, when to fine-tune vs prompt vs RAG
- ☐ Perplexity, benchmarks, task eval sets, LLM-as-judge, RAG evaluation, red-teaming
- ☐ Prompt engineering, APIs, RAG pipelines, tool use, agents
- ☐ VAEs, GANs, diffusion models, CLIP, multimodal models

B.5 Reinforcement learning (Chapter 20)

- ☐ Agent, environment, state, action, reward, policy, value
- ☐ Exploration vs exploitation; Q-learning; DQN; policy gradients and PPO

B.6 Engineering, MLOps and responsibility (Chapters 21–24)

- ☐ Git: commits, branches, pull requests, merge conflicts, .gitignore

- ☐ Docker: images, containers, Dockerfiles, registries
- ☐ Serving a model as an API (FastAPI)
- ☐ Kubernetes: pods, deployments, services, autoscaling, rollouts
- ☐ Experiment tracking, model registry, CI/CD, monitoring, data and concept drift
- ☐ Fairness metrics, explainability (SHAP), privacy, security, prompt injection, model cards

B.7 Where to go next

Books, in the order suggested at the start of our conversation:

- *Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow*, Aurélien Géron — more depth on Parts 2–3.
- *Understanding Deep Learning*, Simon Prince — free online; deeper theory for Parts 3–4.
- *Build a Large Language Model (From Scratch)*, Sebastian Raschka — extends Chapter 15 to a real GPT.
- *Hands-On Large Language Models*, Jay Alammar and Maarten Grootendorst — visual guide to using LLMs.
- *AI Engineering*, Chip Huyen — building and evaluating LLM products; pairs with Chapters 17, 18 and 23.

Courses and videos: Andrew Ng's *Machine Learning Specialization* and *Deep Learning Specialization*; fast.ai's *Practical Deep Learning for Coders*; Andrej Karpathy's *Neural Networks: Zero to Hero* (builds GPT step by step); the free Hugging Face courses on LLMs and diffusion.

Landmark papers to read once you're comfortable: *Attention Is All You Need* (2017, transformers), *BERT* (2018), *Language Models are Few-Shot Learners* (2020, GPT-3), *Training language models to follow instructions with human feedback* (2022, RLHF), *Denoising Diffusion Probabilistic Models* (2020), *LoRA* (2021).

Practice: Kaggle competitions and datasets; rebuild each chapter's project on data from your own work; publish your code on GitHub (Chapter 21). A portfolio of three finished projects — for example a tuned tabular model served in Docker, a RAG assistant with an evaluation set, and a small fine-tuned LLM — demonstrates most of this checklist.